

# DEEP NEURAL NETWORKS

---

Trends and Applications

Ekapol Chuangsuwanich

July 14, 2017

# DNNs (Deep Neural Networks)

- Why deep learning?
- Greatly improved performance in ASR and other tasks (Computer Vision, Robotics, Machine Translation, NLP, etc.)
- Surpassed human performance in many tasks

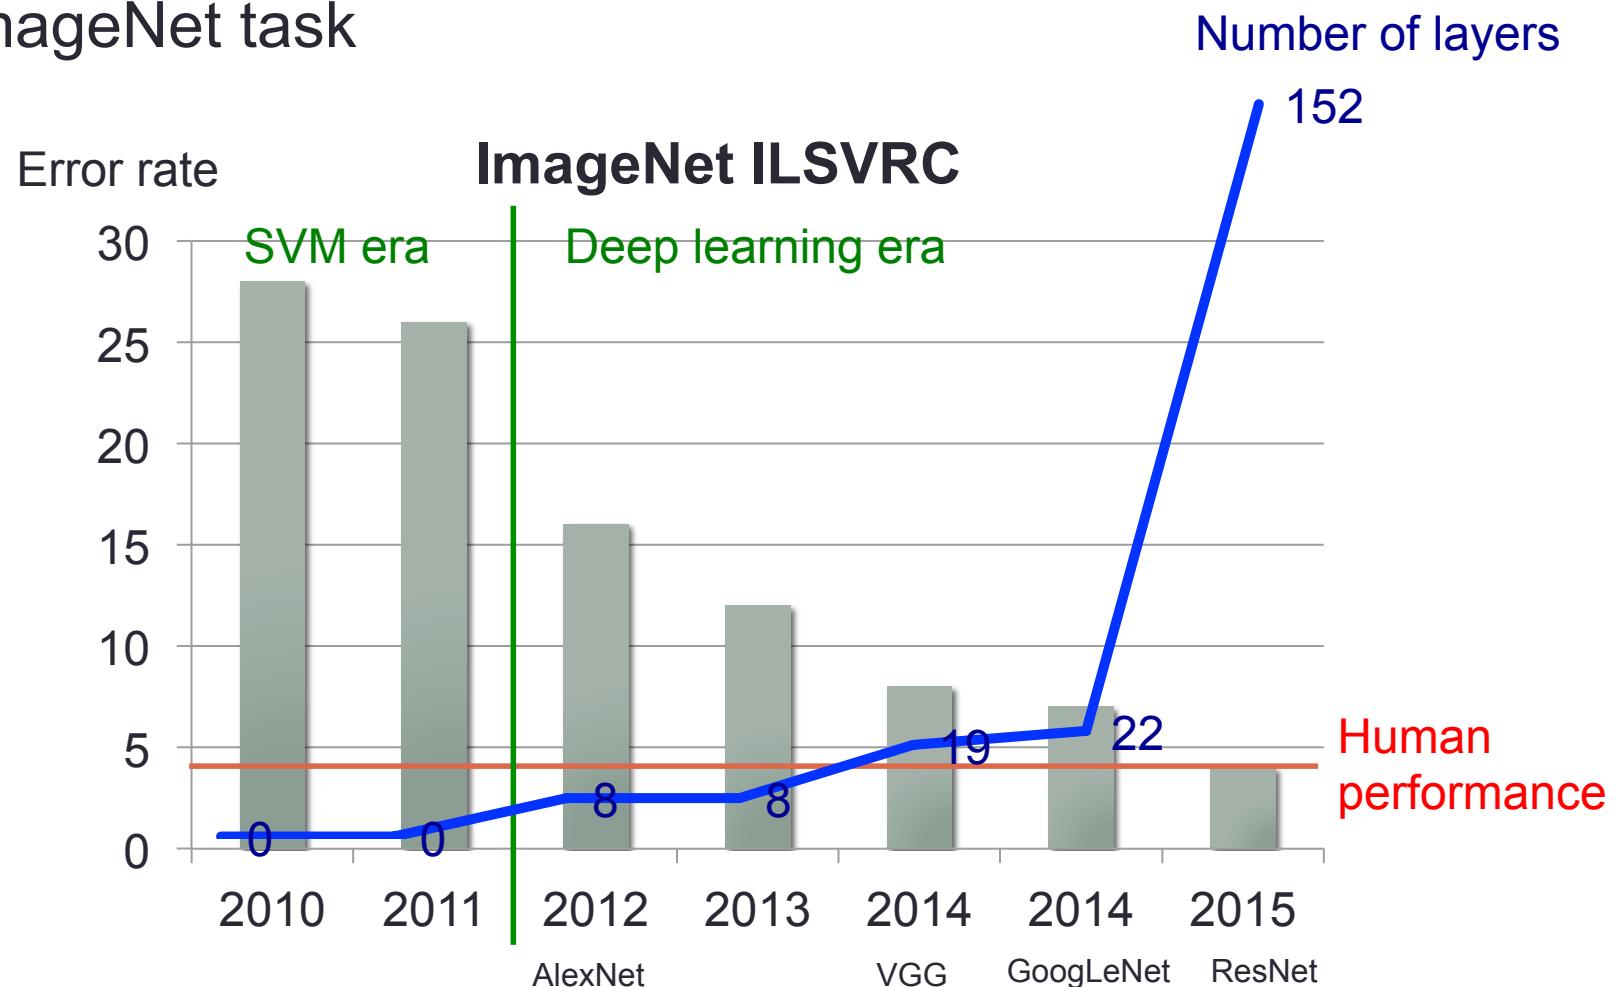
Task	Previous state-of-the-art	Deep learning (2012)	Deep learning (2017)
TIMIT	24.4%	20.0%	17.0%
Switchboard	23.6%	16.1%	5.5%
Google voice search	16.0%	12.3%	4.9%

# Why now

- Neural Networks has been around since 1990s
- Big data – DNN can take advantage of large amounts of data better than other models
- GPU – Enable training bigger models possible
- Deep – Easier to avoid local minima when the model is large

# Wider and deeper networks

- ImageNet task

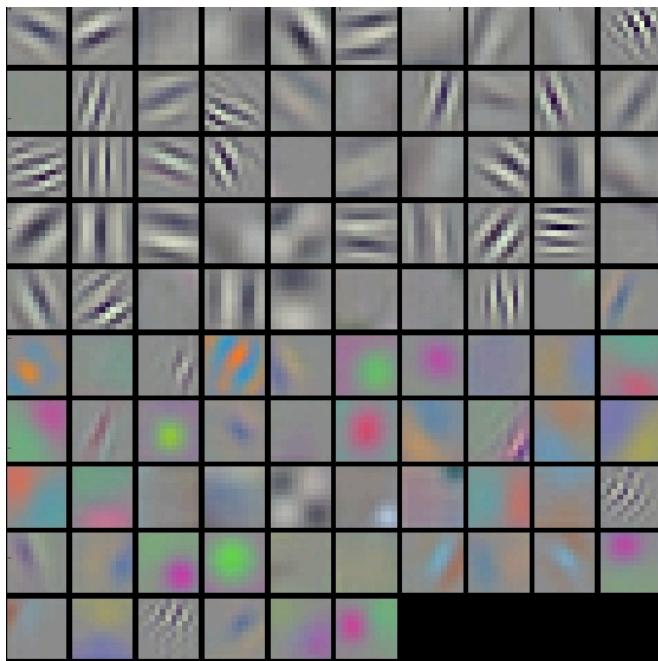


# Why is deep learning good

- Traditional machine learning approaches need feature engineering (MFCC, SIFT, HoG, etc.)
  - Features are based on human understanding of the phenomena
  - Have some simplifying assumptions
- Models also has simplifying assumptions
- Human knowledge captures **known knowns**, and **Known unknowns** NOT **unknown unknowns**
- Deep learning has enough parameters and model freedom to automatically learn the unknown unknowns (and everything else)
  - Downside: Needs LOTS of data

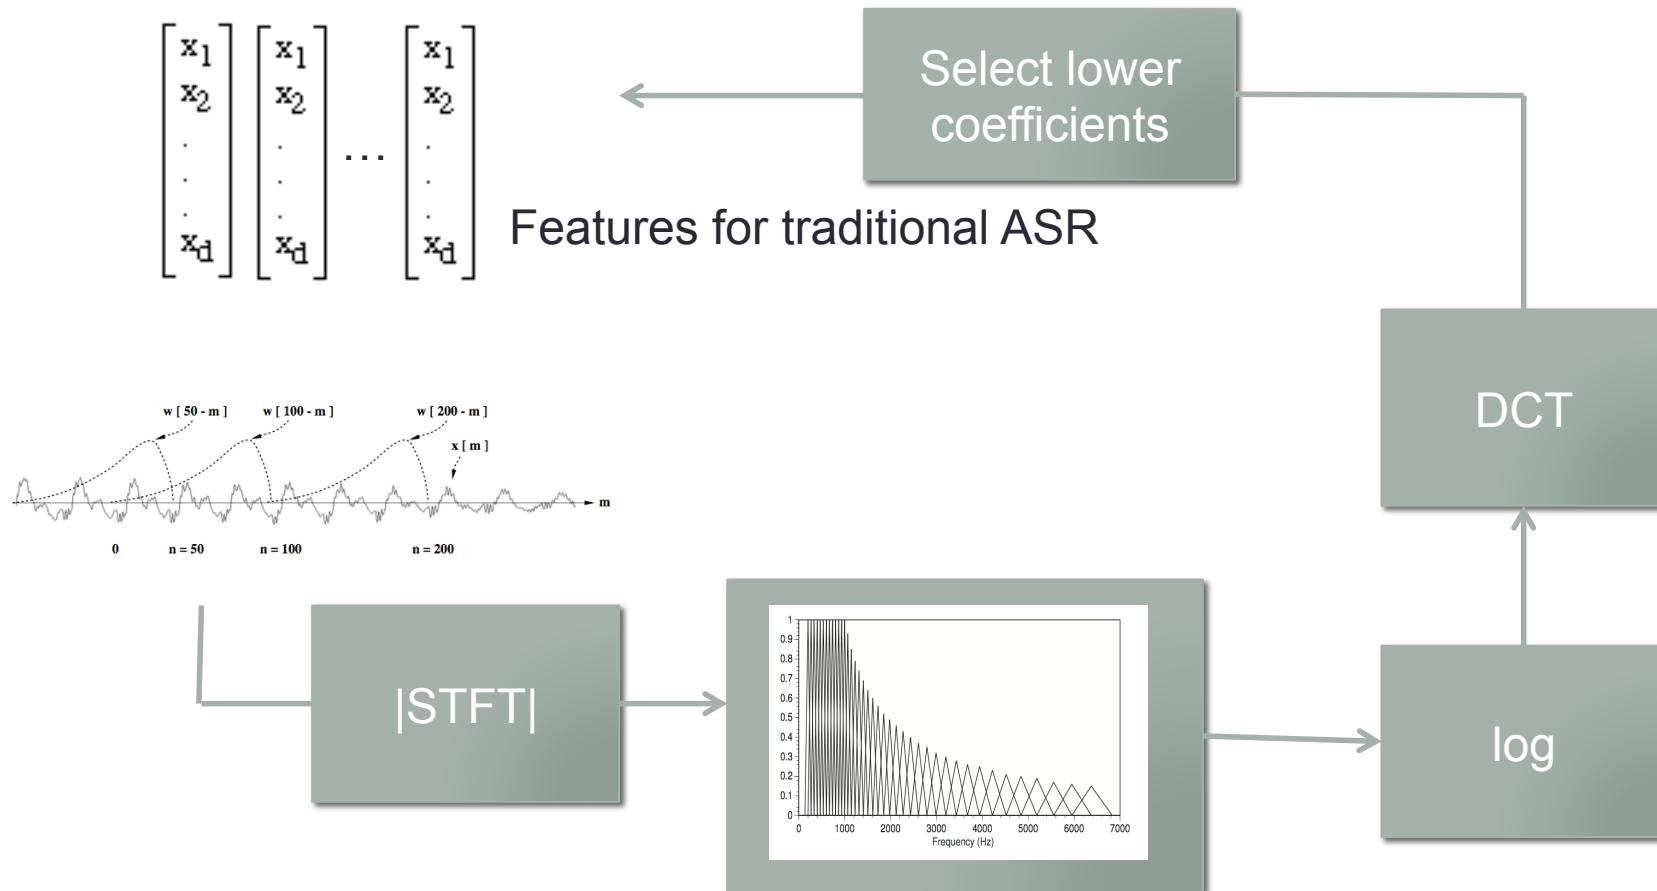
# Why is deep learning good

- DNN can handle higher dimensional input feature
- Deep learning also combines features engineering and modeling into one single optimization problem
  - Guarantees the matching between the feature and the model



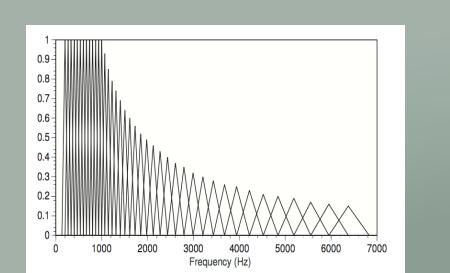
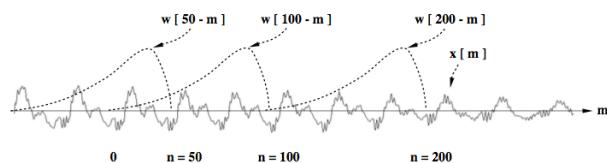
Features learned from a CNN for a vision task.  
Inputs are now raw pixels instead of descriptors.

# MFCC (Mel Frequency Cepstral Coefficient)

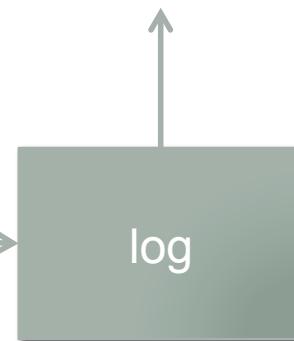


# MFCC (Mel Frequency Cepstral Coefficient)

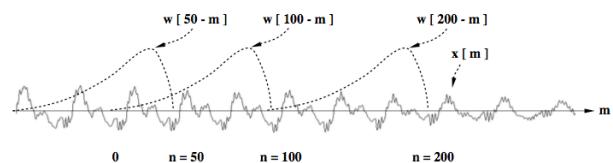
Features for ASR with DNN



$$\begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix}, \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix}, \dots, \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix}$$



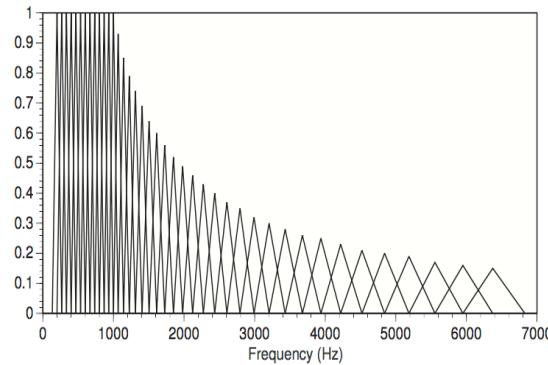
# MFCC (Mel Frequency Cepstral Coefficient)



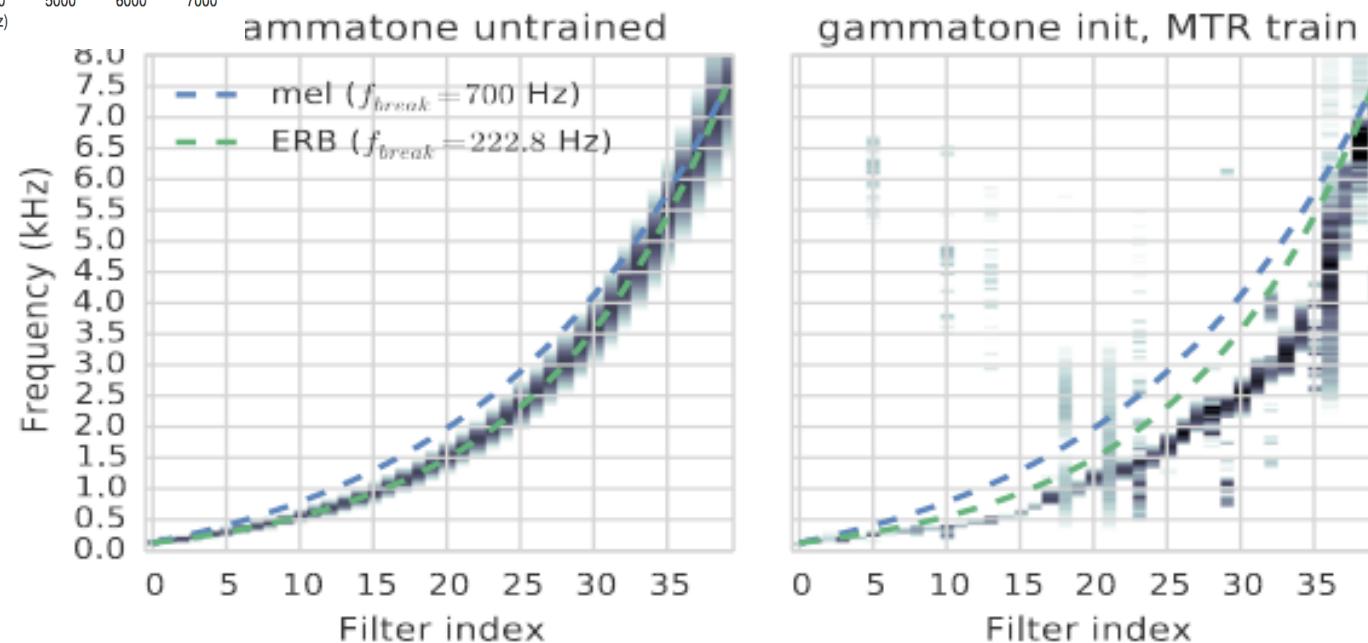
Features for ASR with DNN

$$\begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix} \dots \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix}$$

# Learning the Mel filters



Filter type	WER
DNN+Random	16.4
DNN+Gammatone (untrained)	16.4
DNN+Gammatone (trained)	16.2

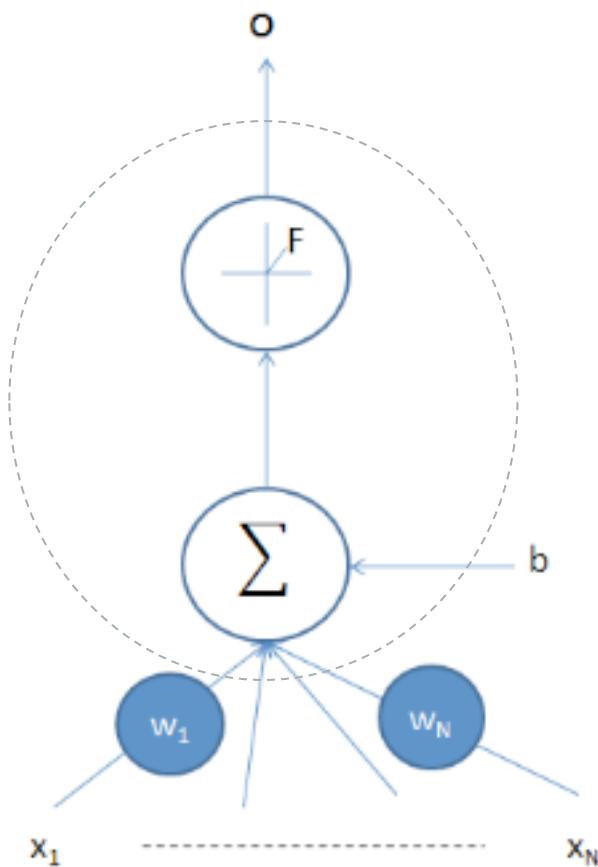


# DNNs architectures

- Fully connected networks
  - Components in neural networks
  - How is neural networks trained
- Convolutional neural networks
- Recurrent neural networks
- Gated recurrent units
- Long short-term memory networks
- Encoder-Decoder

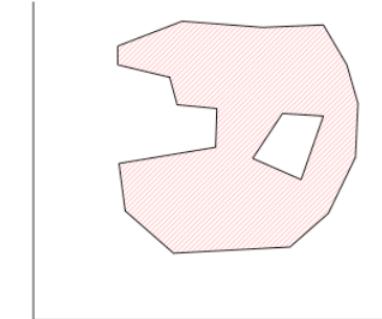
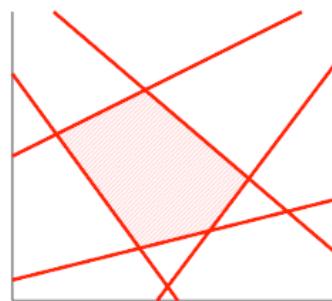
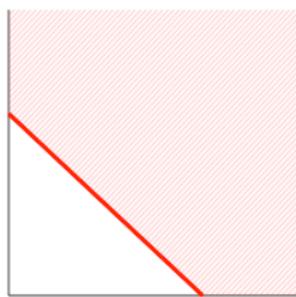
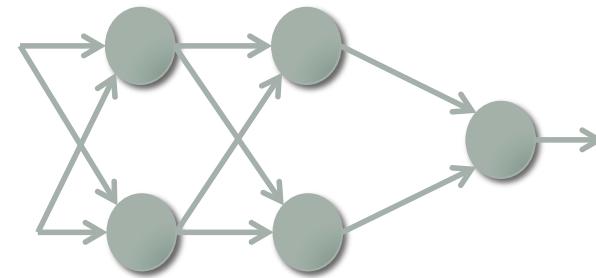
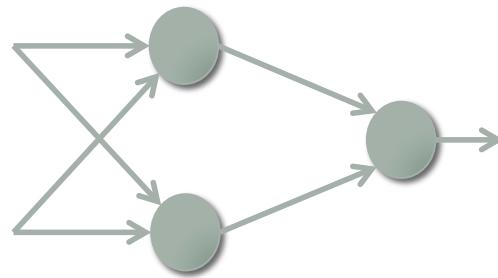
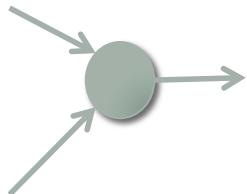
# Fully connected networks

- Many names: feed forward networks or deep neural networks or Multilayer perceptron
  - Composed of multiple neurons
- 
- $o = F(\mathbf{w}^T \mathbf{x} + b)$
  - $o$  - activation output
  - $F$  - nonlinear function
  - $\mathbf{w}$  - weight vector
  - $b$  - bias

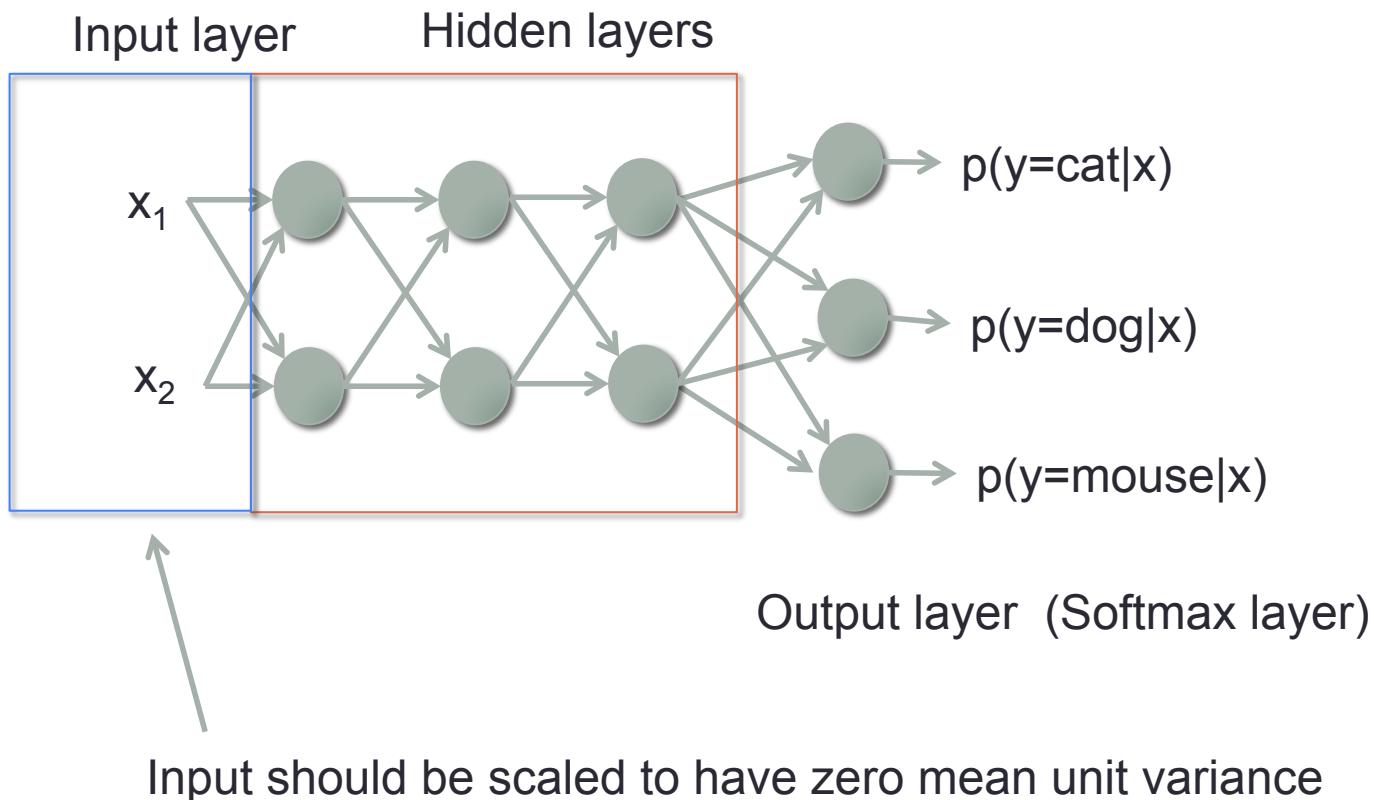


# Combining neurons

- Each neuron splits the feature space with a hyperplane
- Stacking neuron creates more complicated decision boundaries
- More powerful but prone to overfitting

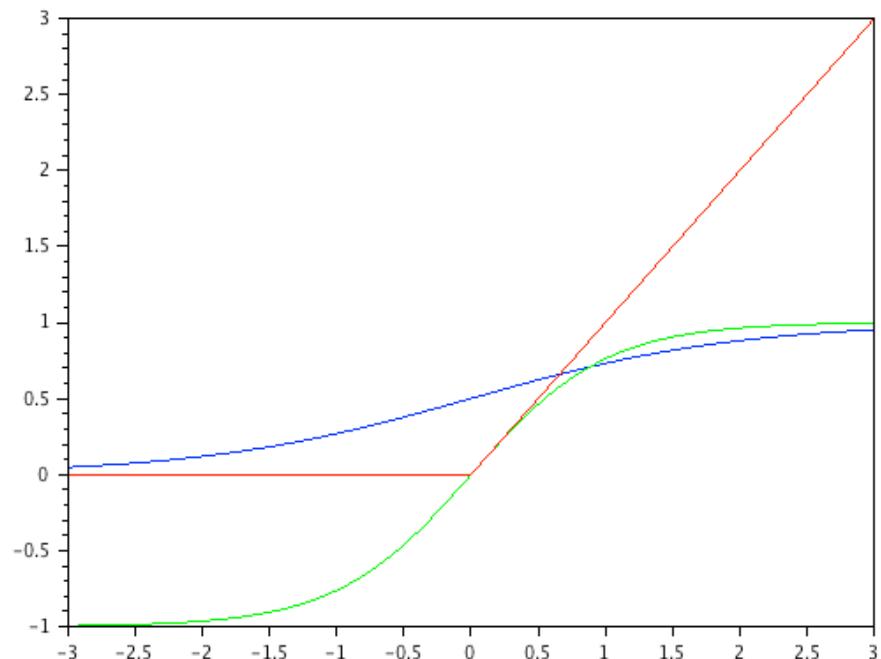


# Terminology



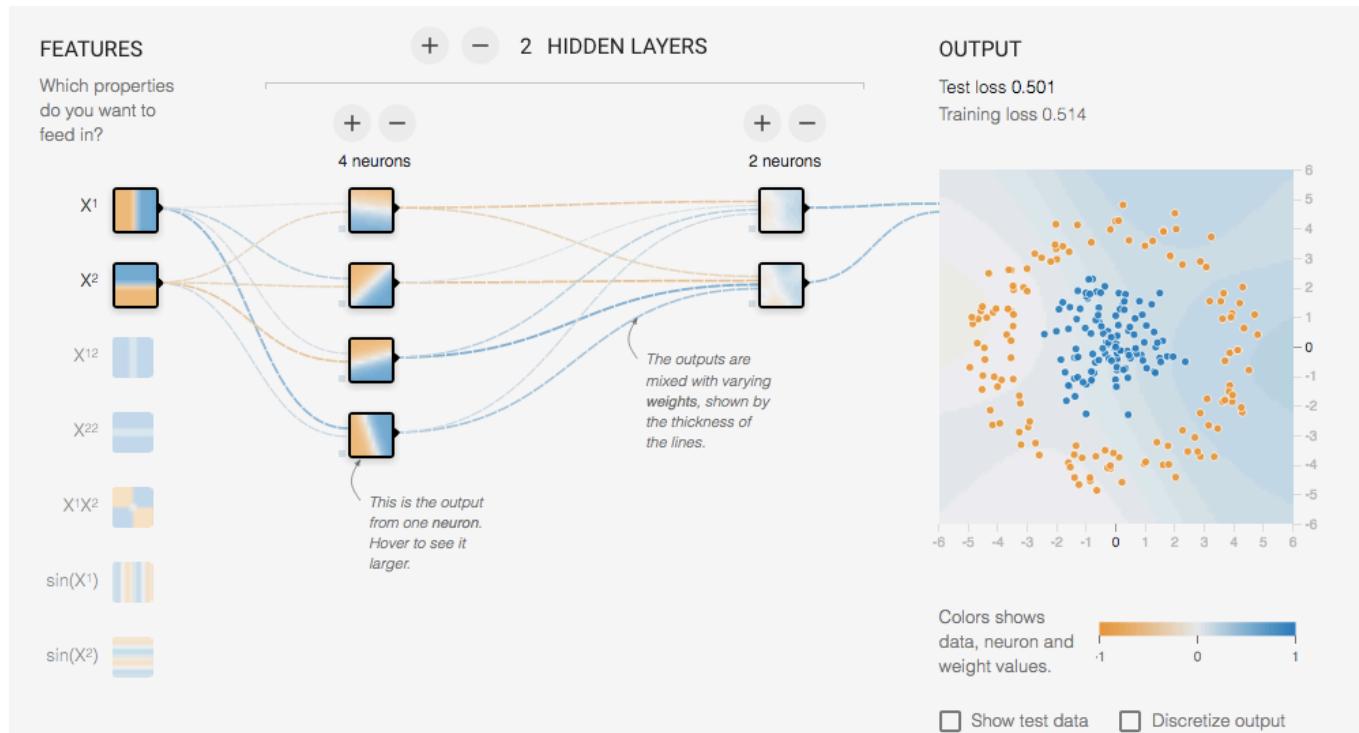
# Non-linearity

- The F Non-linearity is important in order to stack neurons
  - If F is linear, a multi layered network can be collapsed as a single layer (by just multiplying weights together)
- Sigmoid
- tanh
- Rectified Linear Unit (ReLU)
- Most popular is ReLU and its variants (Fast to train, and more stable)

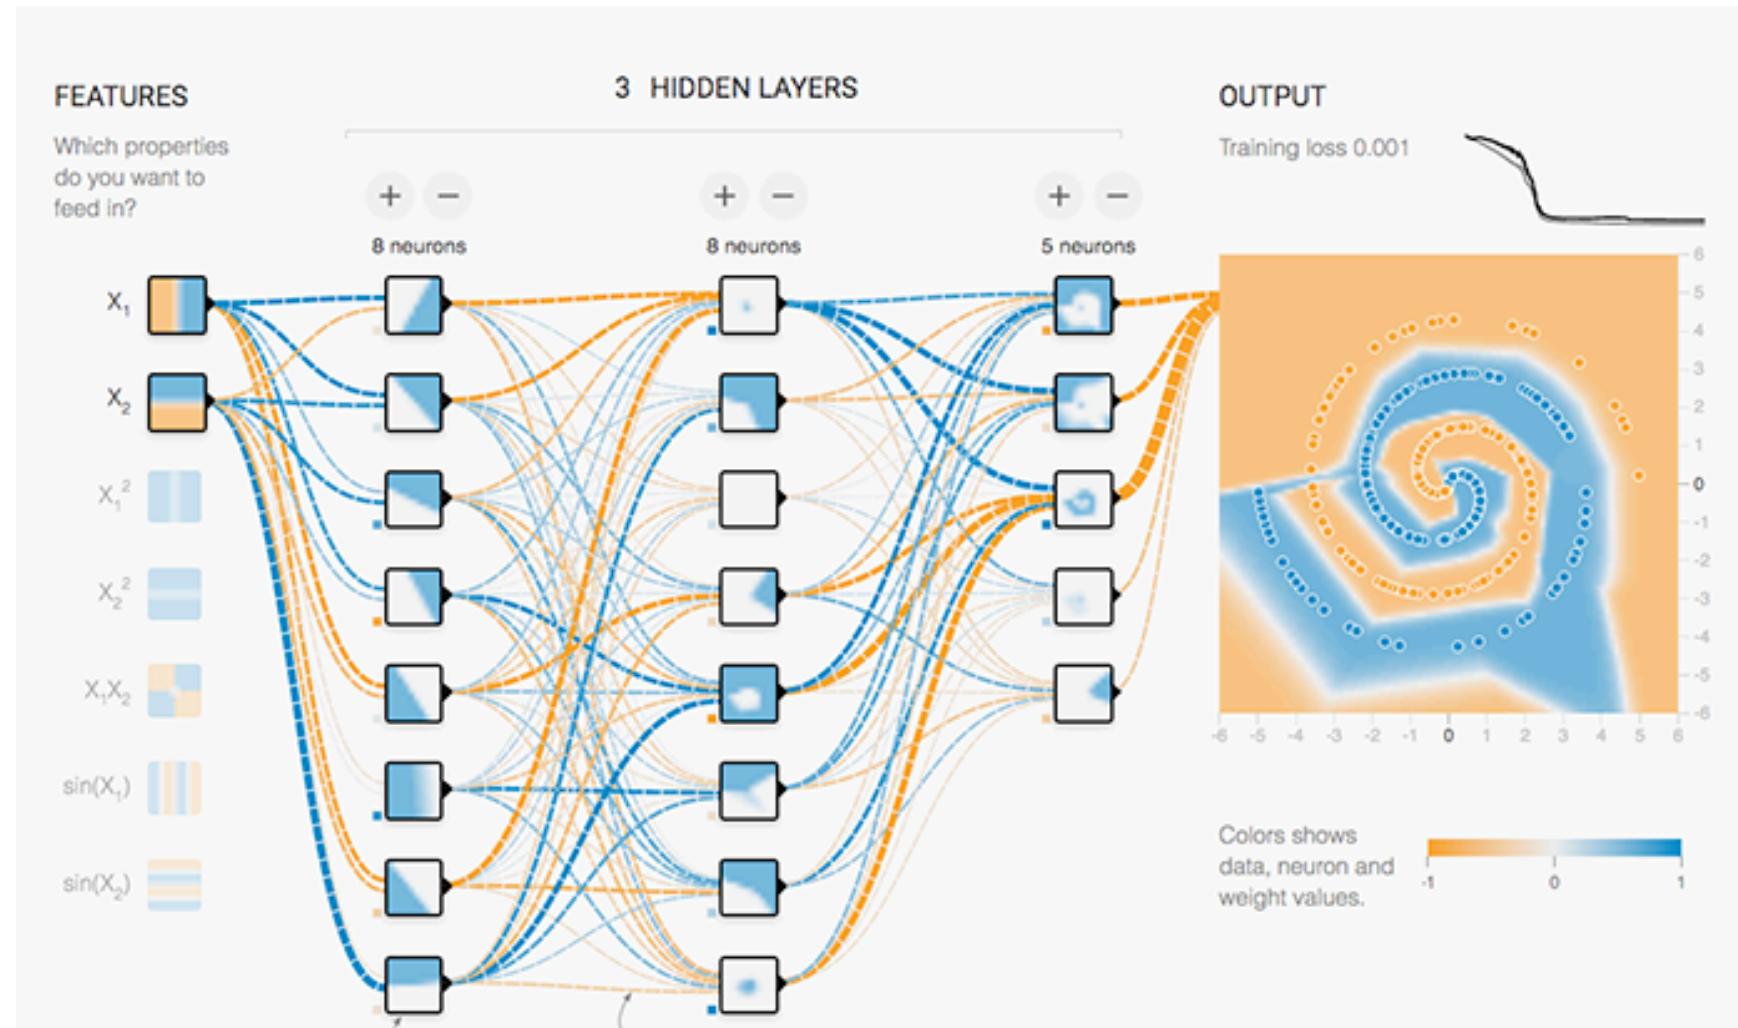


# Playground

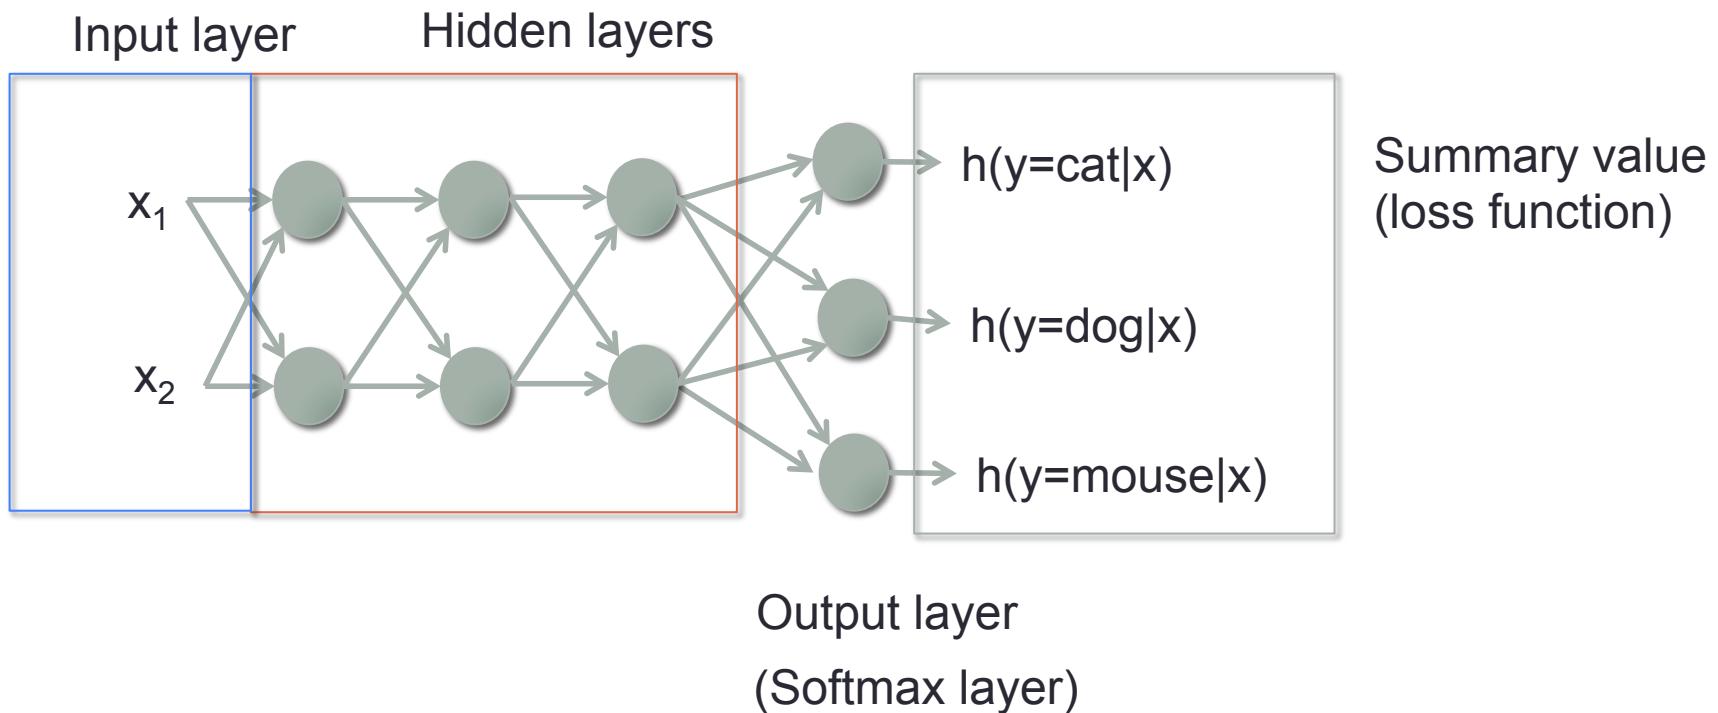
- <http://playground.tensorflow.org/>



# Playground



# Terminology



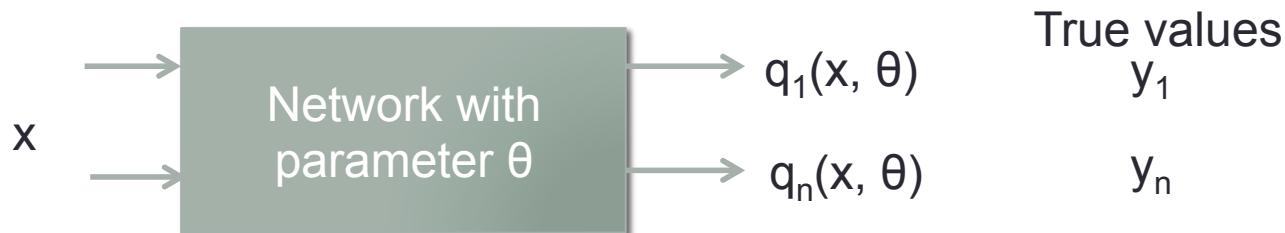
# Objective function (Loss function)

- Can be any function that summarizes the performance into a single number
- Cross entropy
  - Used for softmax outputs (probabilities)

$$L = -\sum_n y_n \log q_n(x, \theta)$$

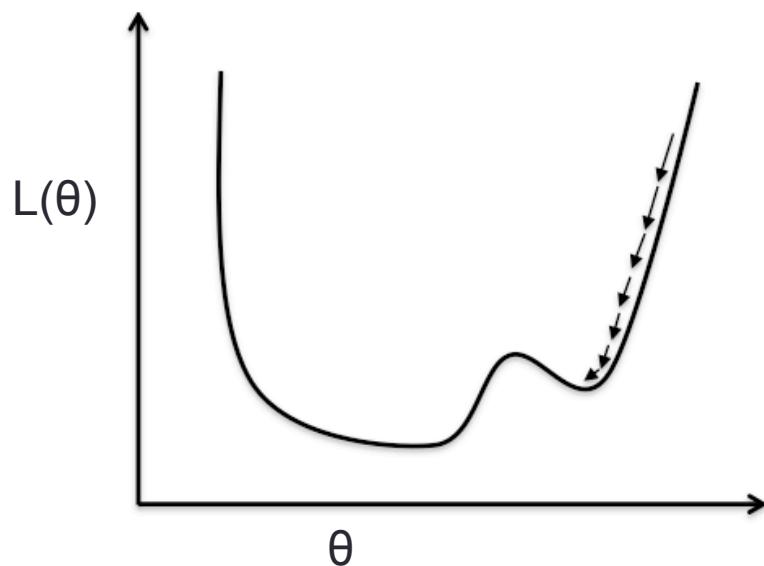
- Sum of square errors
  - Used for any real valued outputs

$$L = -\frac{1}{2} \sum_n (y_n - q_n(x, \theta))^2$$



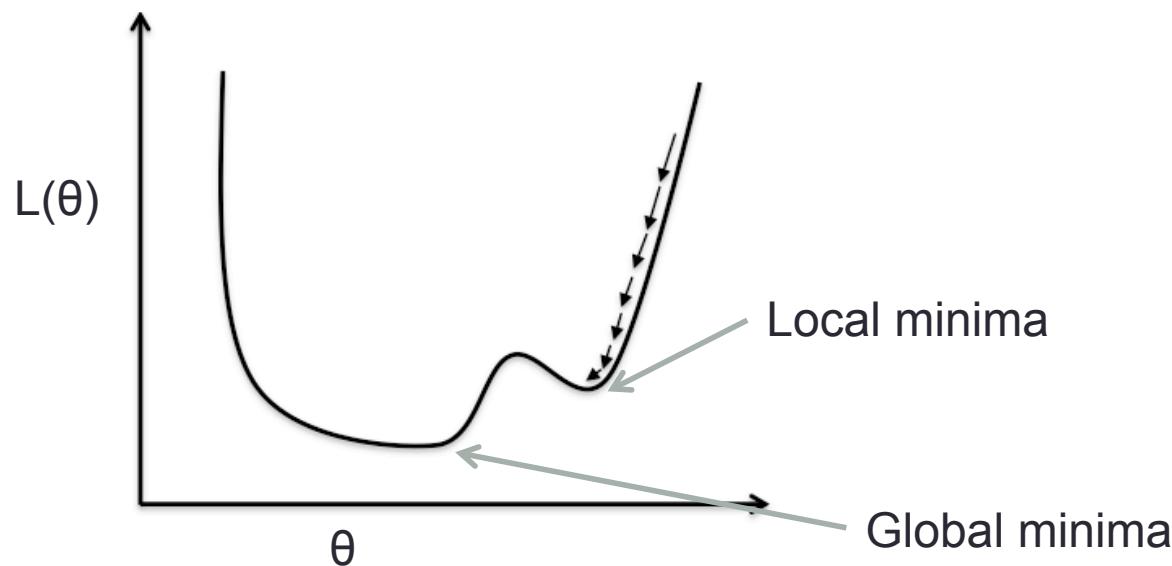
# Minimization using gradient descent

- We want to minimize  $L$  with respect to  $\theta$ 
  - Differentiate with respect to  $\theta$
  - Gradients passes through the network by Back Propagation



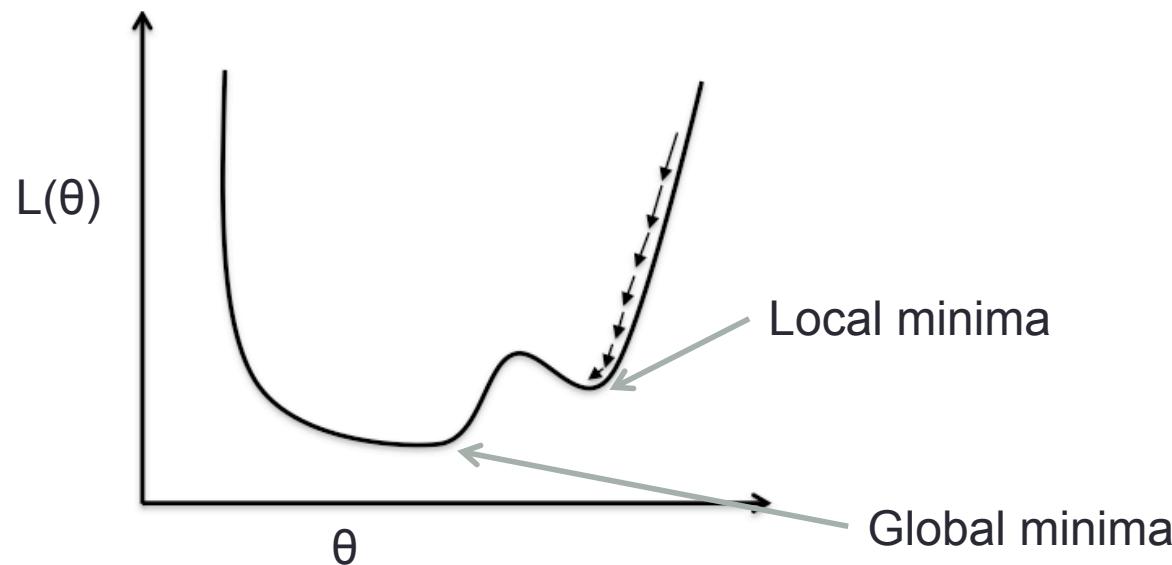
# Deep vs Shallow

- The loss function of neural network is non-convex (and non-concave)
  - Local minimas can be avoided with convexity
  - Convexity gives easier training



# Deep vs Shallow

- If deep, most local minimas are the global minima!
  - Always a way to lower the loss in the network with millions of parameters
  - Enough parameters to remember every training examples
  - Does not imply anything about generalization

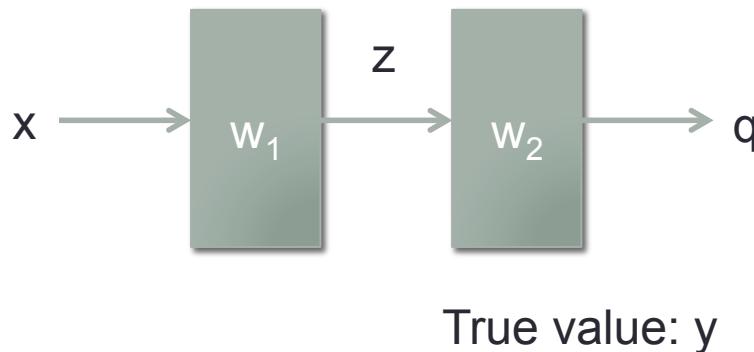


# Stochastic gradient descent (SGD)

- Main method to train a neural network
- Consider you have one million training examples
  - Gradient descent computes the objective function of **all** samples, then decide direction of descent
    - Takes too long
    - SGD computes the objective function on **subsets** of samples
      - The subset should not be biased and properly randomized to ensure no correlation between samples
  - The subset is called a mini-batch
  - Size of the mini-batch determines the training speed and accuracy
    - Usually somewhere between 32-1024 samples per mini-batch

# Back propagation

- To update the network, we need to compute the direction (gradient) that minimizes the loss function
  1. Compute the loss function of a mini-batch (**forward pass**)
  2. Compute the gradient (partial derivative) with respect to each parameter (**backward pass**)
- Key: Use chain rule



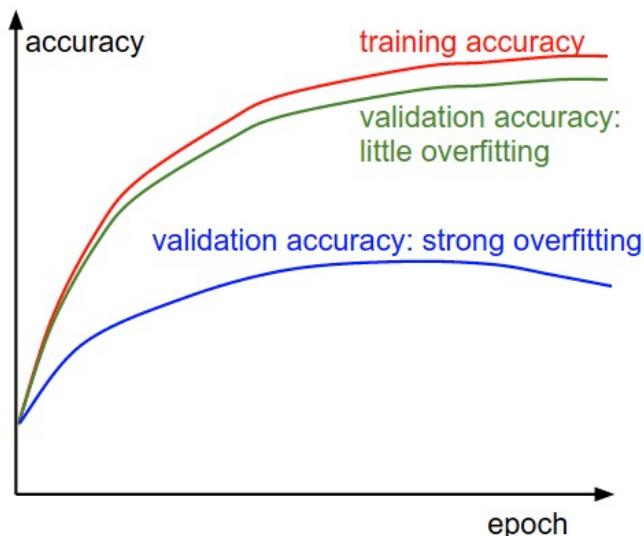
$$\begin{aligned}z &= w_1 x & q &= w_2 z \\L &= (y - q)^2\end{aligned}$$

$$\begin{aligned}\frac{dL}{dw_2} &= \frac{dL}{dq} \frac{dq}{dw_2} \\&= [2(y - q)(-1)][z]\end{aligned}$$

$$\begin{aligned}\frac{dL}{dw_1} &= \frac{dL}{dq} \frac{dq}{dz} \frac{dz}{dw_1} \\&= [2(y - q)(-1)][w_2][x]\end{aligned}$$

# Overfitting

- You can keep doing back propagation forever!
- The training loss will always go down
- But it overfits
- Need to monitor performance on a held out set
- Stop or decrease learning rate when overfit happens

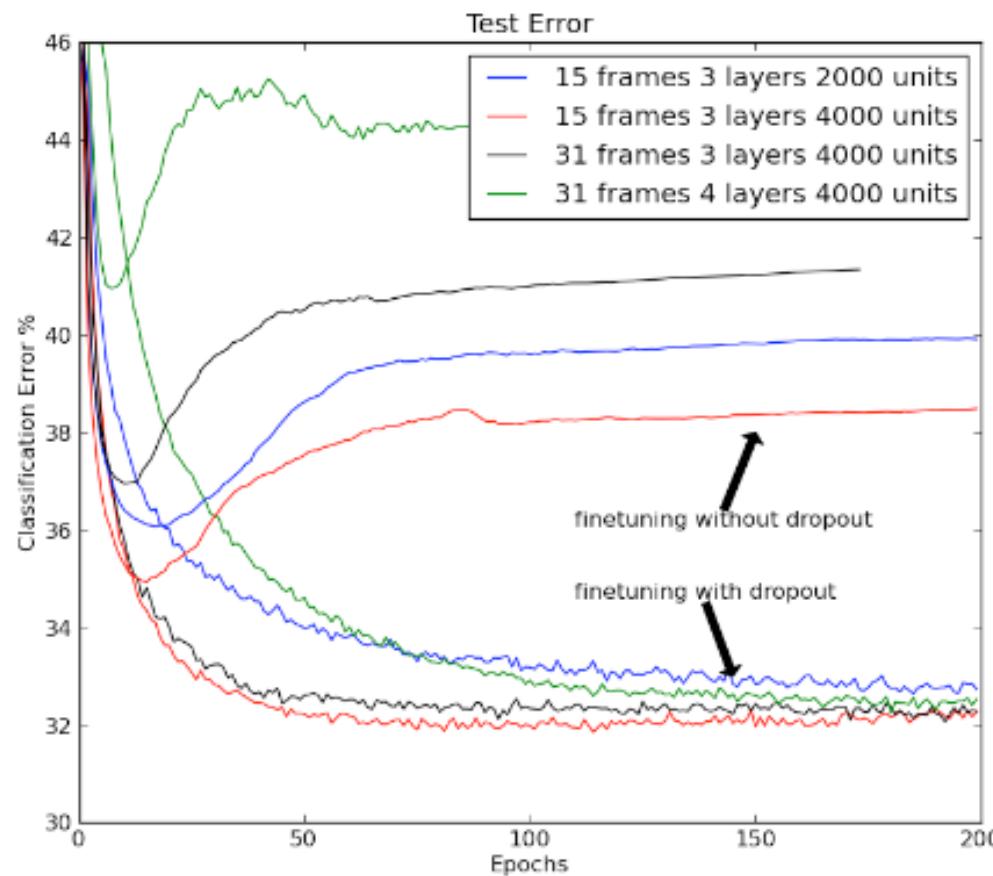


# Reducing overfitting - dropout

- A *RECENT* (2012) regularization technique for reducing overfitting
- Randomly turn off different subset of neurons during training
  - Network no longer depend on any particular neuron
  - Force the model to have redundancy – robust to any corruption in input data
  - A form of performing model averaging (assemble of experts)
- Now a standard technique

# Dropout on TIMIT

- A phoneme recognition task



# Vanishing/Exploding gradient

- Backprop introduces many multiplications down chain
- The gradient value gets smaller and smaller
  - The deeper the network the smaller the gradient in the lower layers
  - Lower layers changes too slowly (or not at all)
  - Hard to train very deep networks (>6 layers)
- The opposite can also be true. The gradient explodes from repeated multiplication
  - Put a maximum value for the gradient (Gradient clipping)

# Residual networks

- Adds a direct connection between layers (bypass non linearity)
- Gradient now has two components (direct connection and traditional direction)

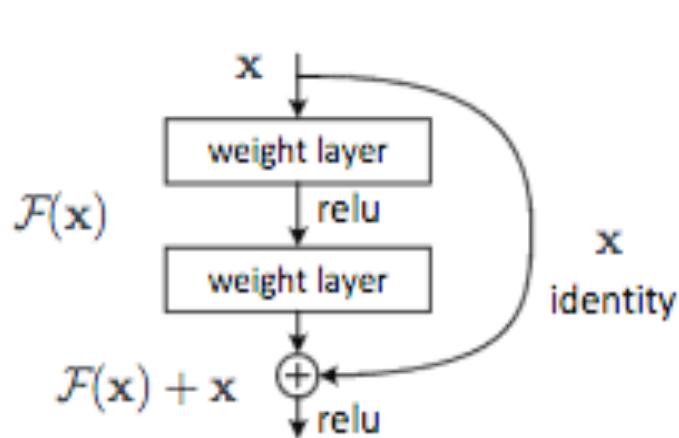


Image recognition task

	plain	ResNet
18 layers	27.94	27.88
34 layers	28.54	<b>25.03</b>

Kaiming He, Deep Residual Learning for Image Recognition

# Residual networks

ASR task

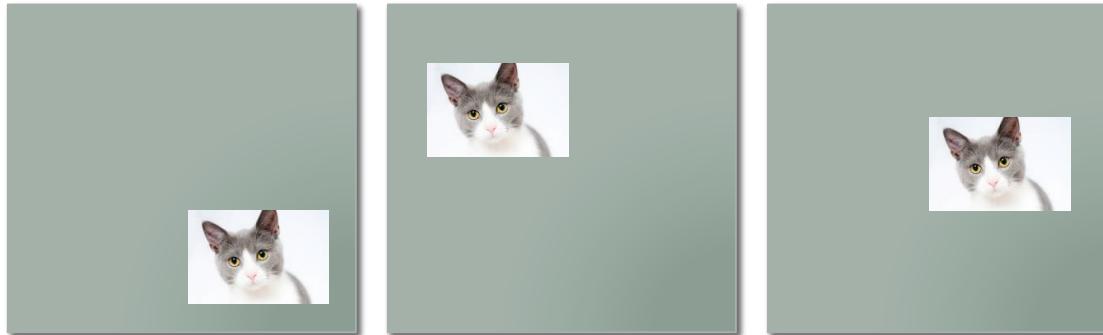
System	#layers	with overlap	no overlap
LSTMP	3	50.7	41.7
LSTMP	8	52.6	43.8
HLSTMP	3	50.4	41.2
HLSTMP	8	50.7	41.3

# DNNs architectures

- ~~Fully connected networks~~
- Convolutional neural networks
- Recurrent neural networks
- Gated recurrent units
- Long short-term memory networks
- Encoder-decoder

# Convolutional Neural Networks (CNNs)

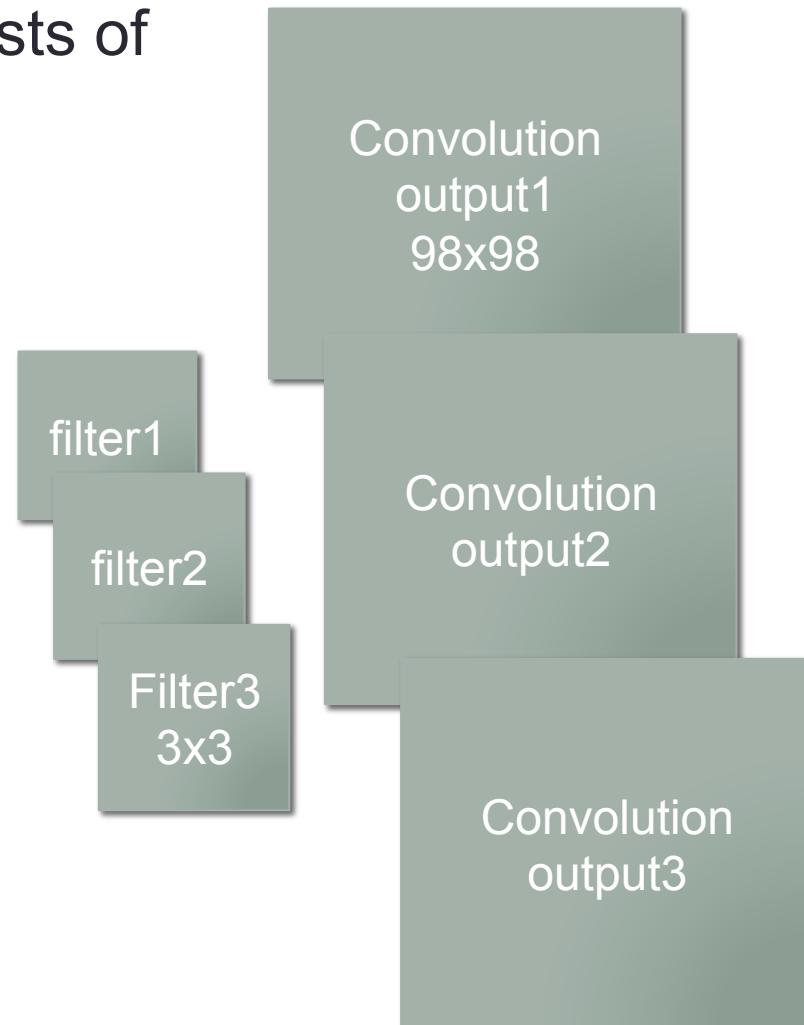
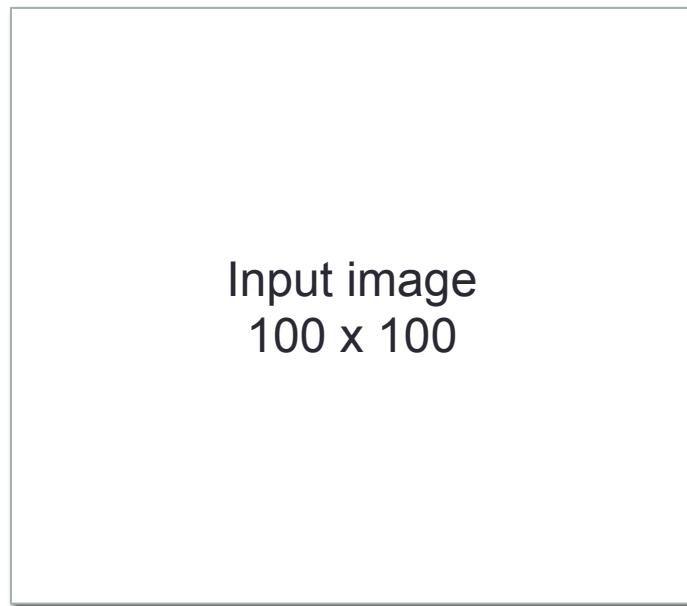
- Consider an image of a cat. DNNs need different neurons to learn every possible location a cat can be



- Can we use the same parameters to learn that a cat exists regardless of location
  - Shift-invariance
  - 2 parts: convolutional layer and pooling layer

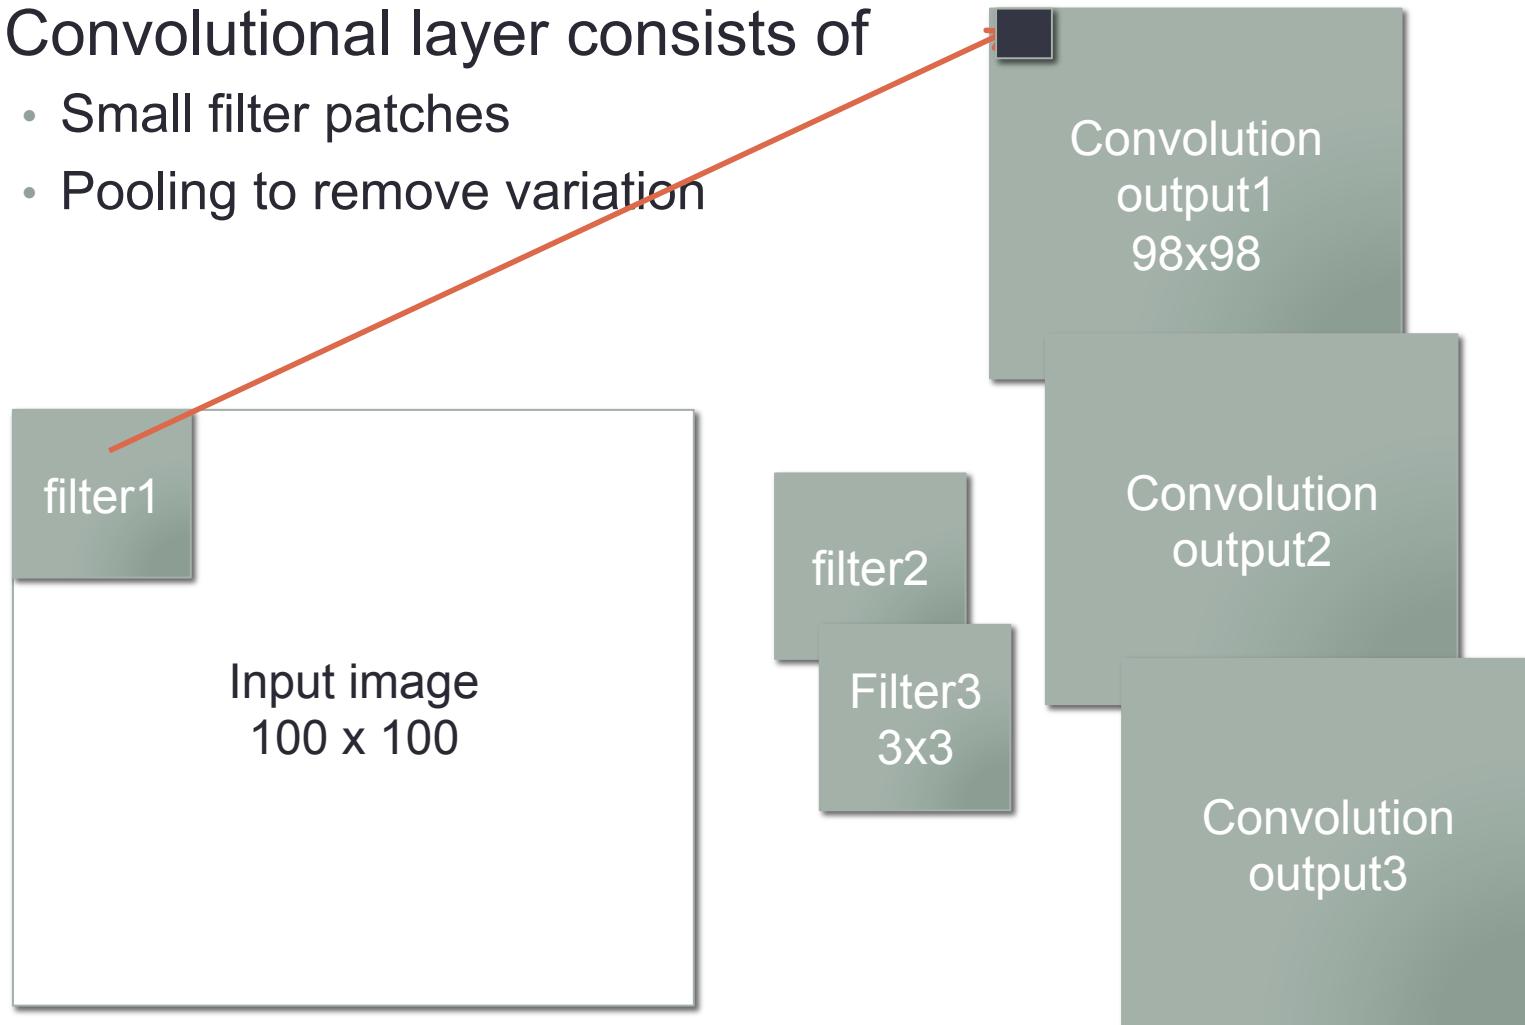
# Convolutional filters

- Convolutional layer consists of
  - Small filter patches
  - Pooling to remove variation



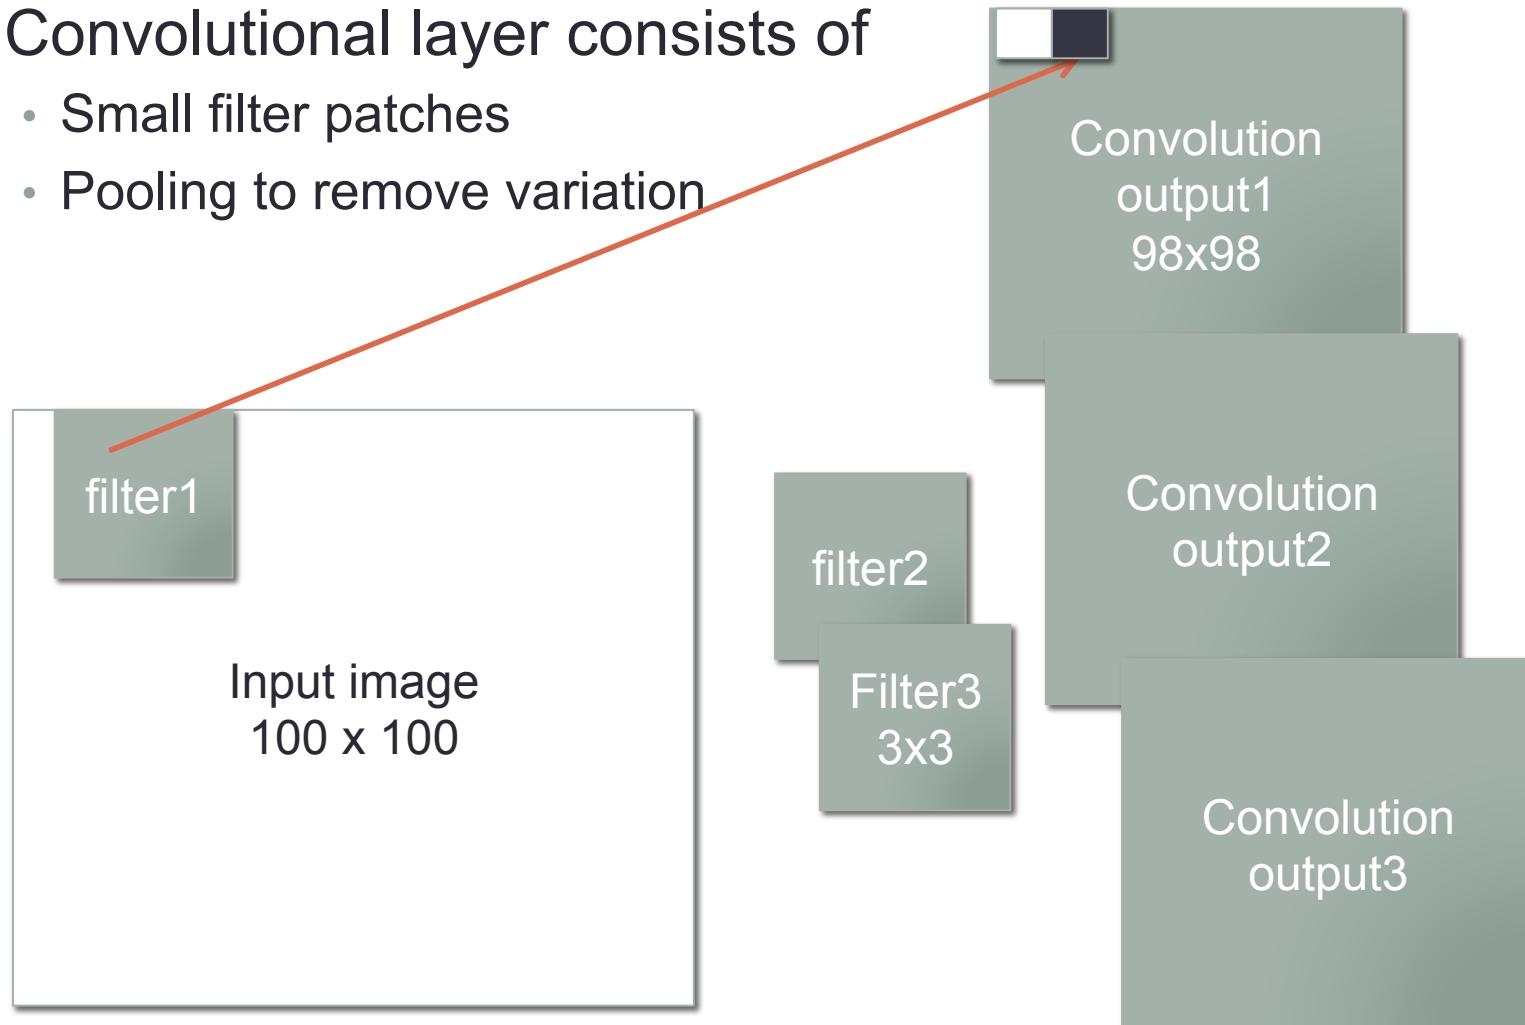
# Convolutional filters

- Convolutional layer consists of
  - Small filter patches
  - Pooling to remove variation



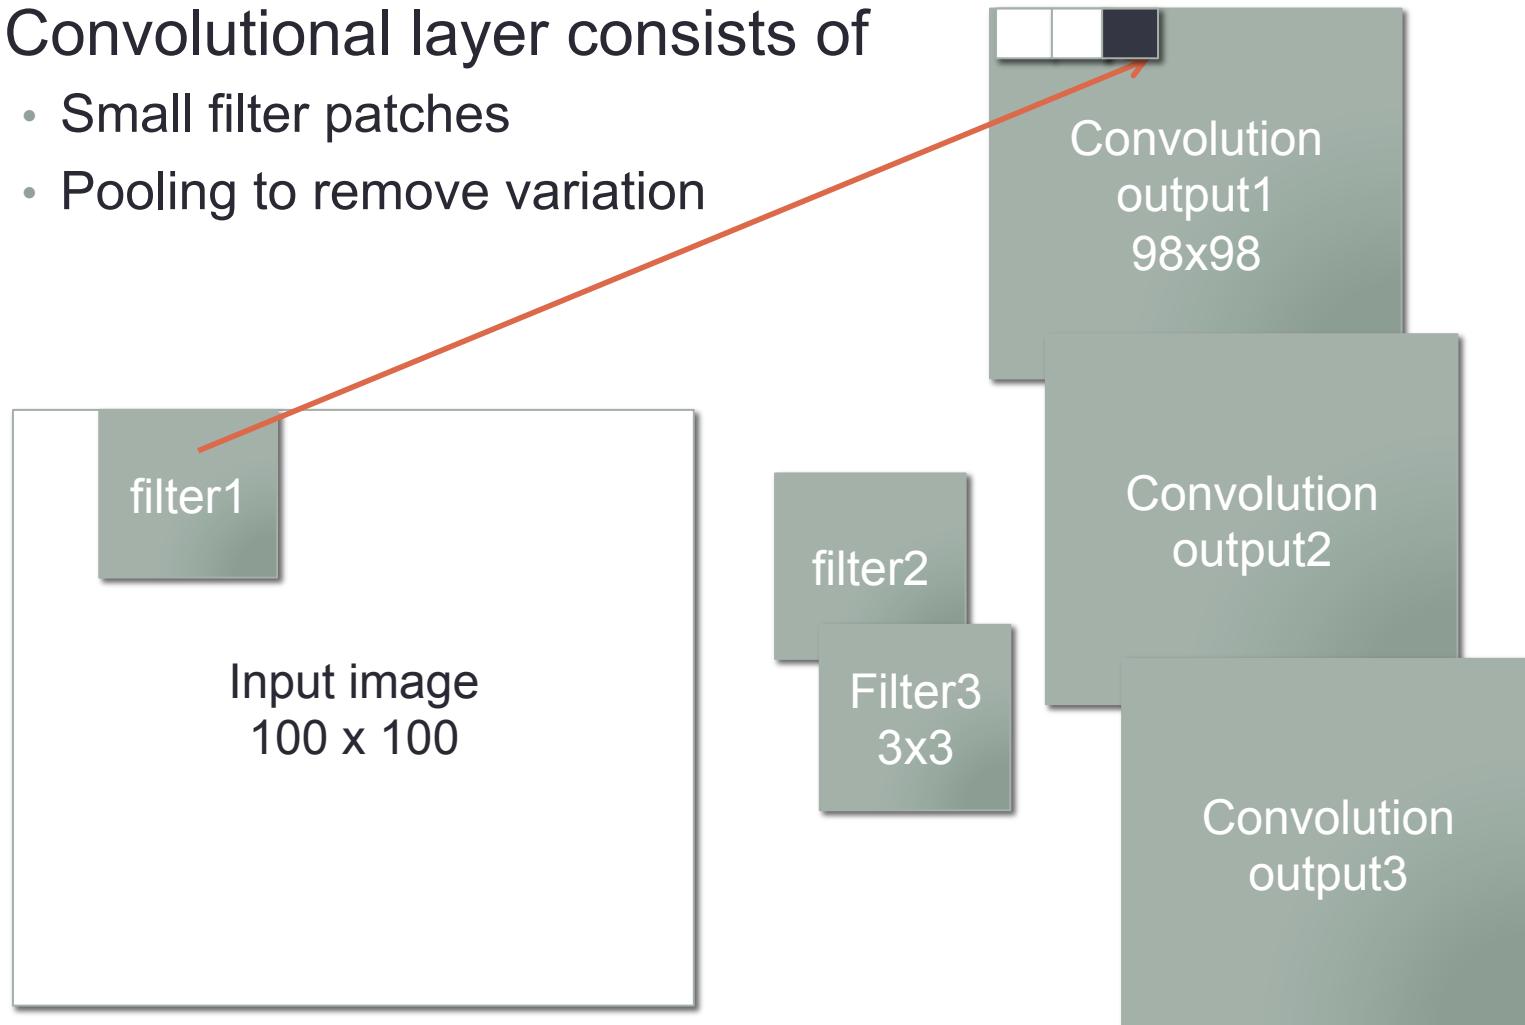
# Convolutional filters

- Convolutional layer consists of
  - Small filter patches
  - Pooling to remove variation



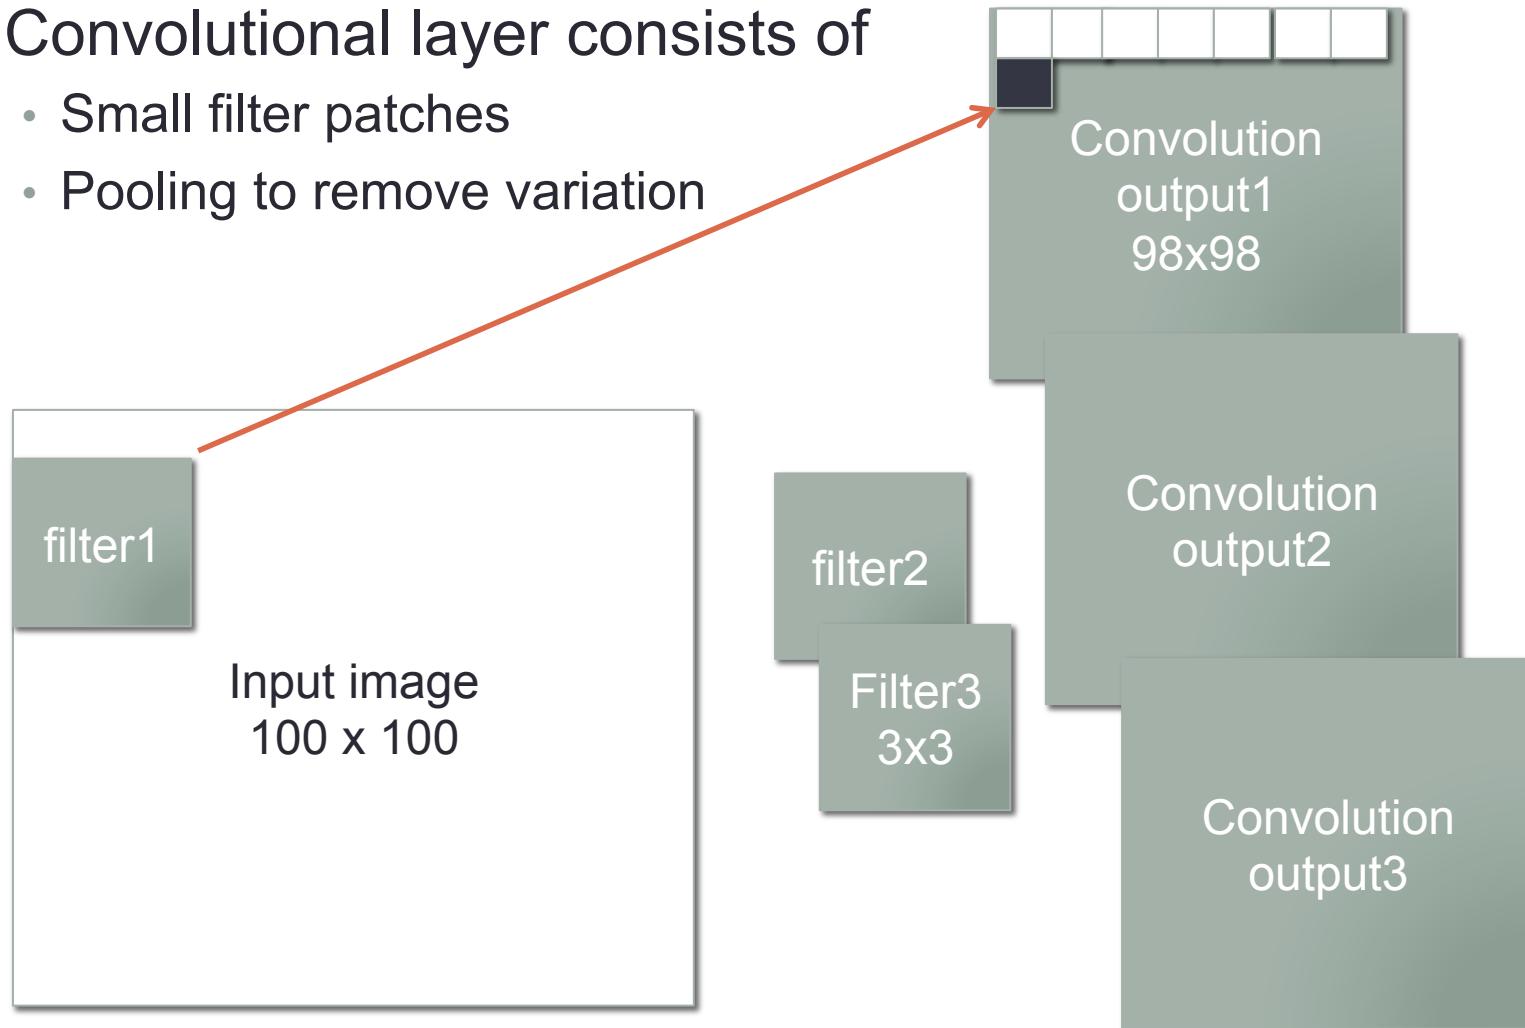
# Convolutional filters

- Convolutional layer consists of
  - Small filter patches
  - Pooling to remove variation



# Convolutional filters

- Convolutional layer consists of
  - Small filter patches
  - Pooling to remove variation



# Pooling/subsampling

- Convolutional layer consists of
  - Small filter patches
  - Pooling to remove variation

Convolution  
output1  
 $98 \times 98$

Convolution  
output2

3x3 Max filter  
with no overlap



Layer output1  
 $33 \times 33$

Layer output2

# Pooling/subsampling

- Convolutional layer consists of
  - Small filter patches
  - Pooling to remove variation



# Pooling/subsampling

- Convolutional layer consists of
  - Small filter patches
  - Pooling to remove variation



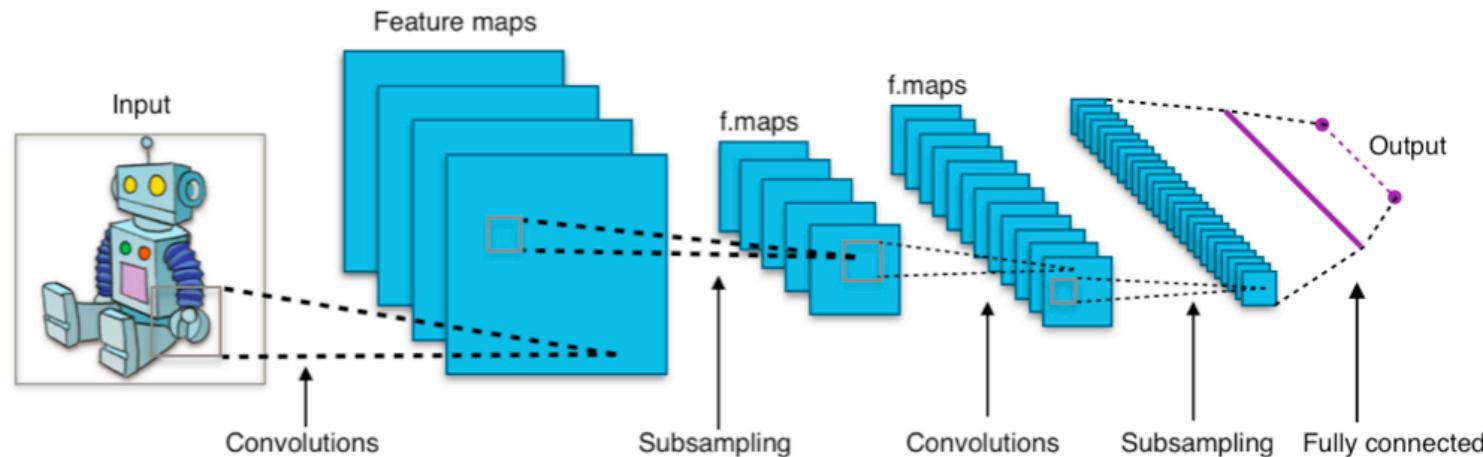
# Pooling/subsampling

- Convolutional layer consists of
  - Small filter patches
  - Pooling to remove variation



# CNN

- Filter size, number of filters, and pooling rate are all parameters
- Usually followed by a fully connected network at the end
  - CNN is good at learning low level features
  - DNN combines the features into high level features and classify

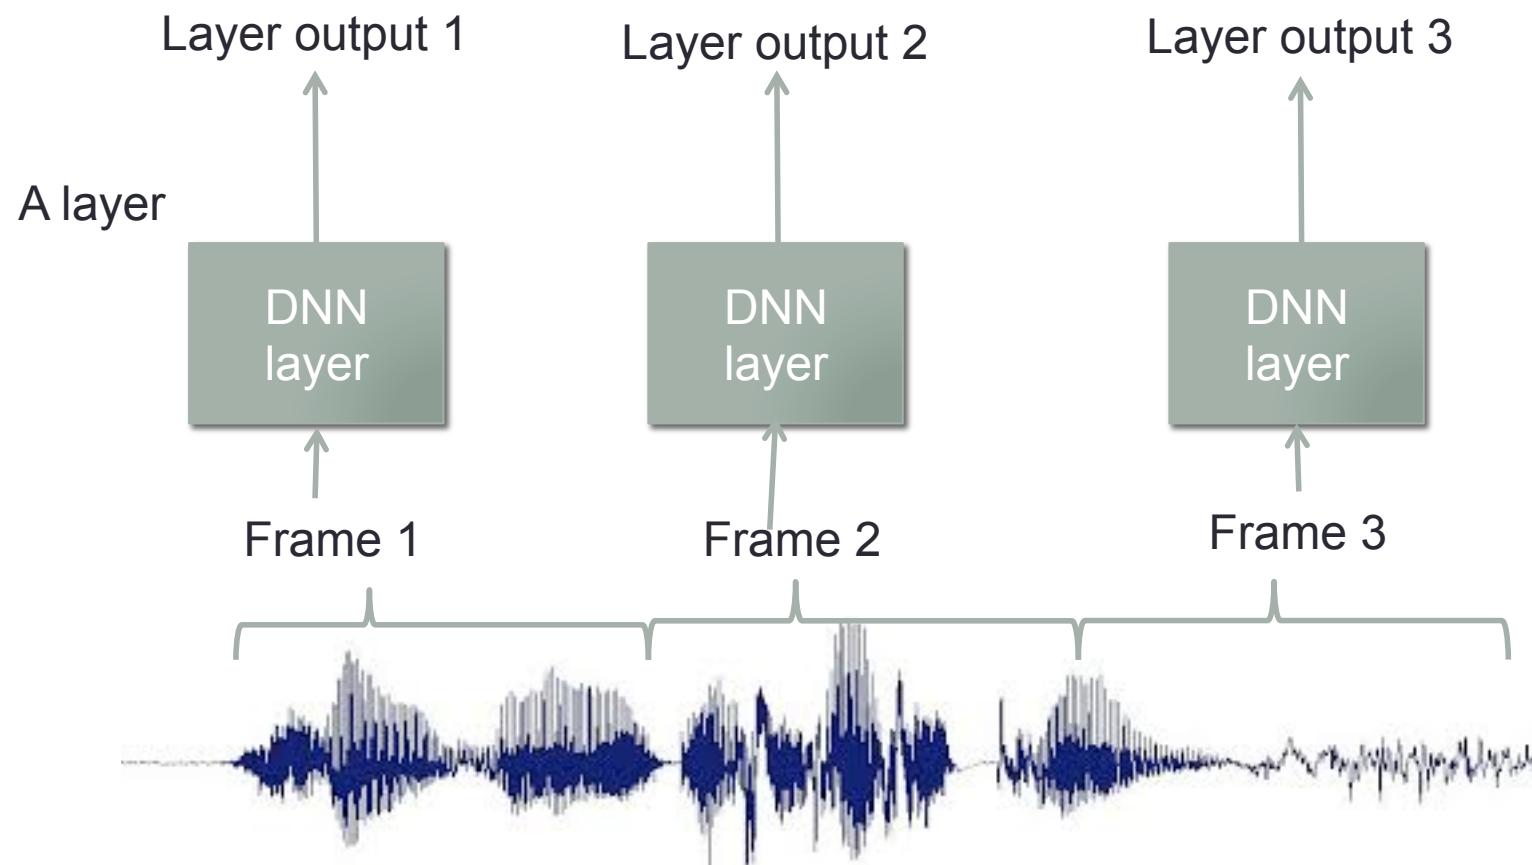


# Recurrent neural network (RNN)

- Many tasks has temporal relationship
  - Speech – CVC structure
  - NLP – open bracket should follow by an end bracket
- DNN parameters are fixed (for the forward pass). Thus, it does not have the memory to keep track of these relationships
- RNN try to force the model to take into account of previous inputs and outputs

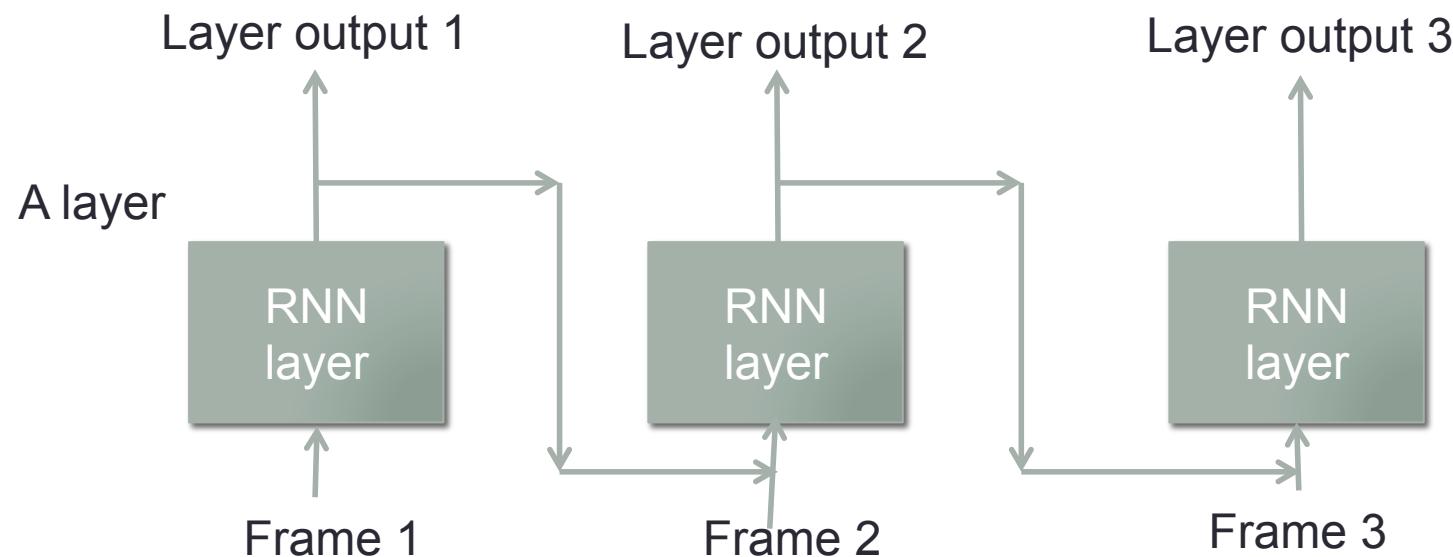
# Recurrent neural network (RNN)

- DNN framework



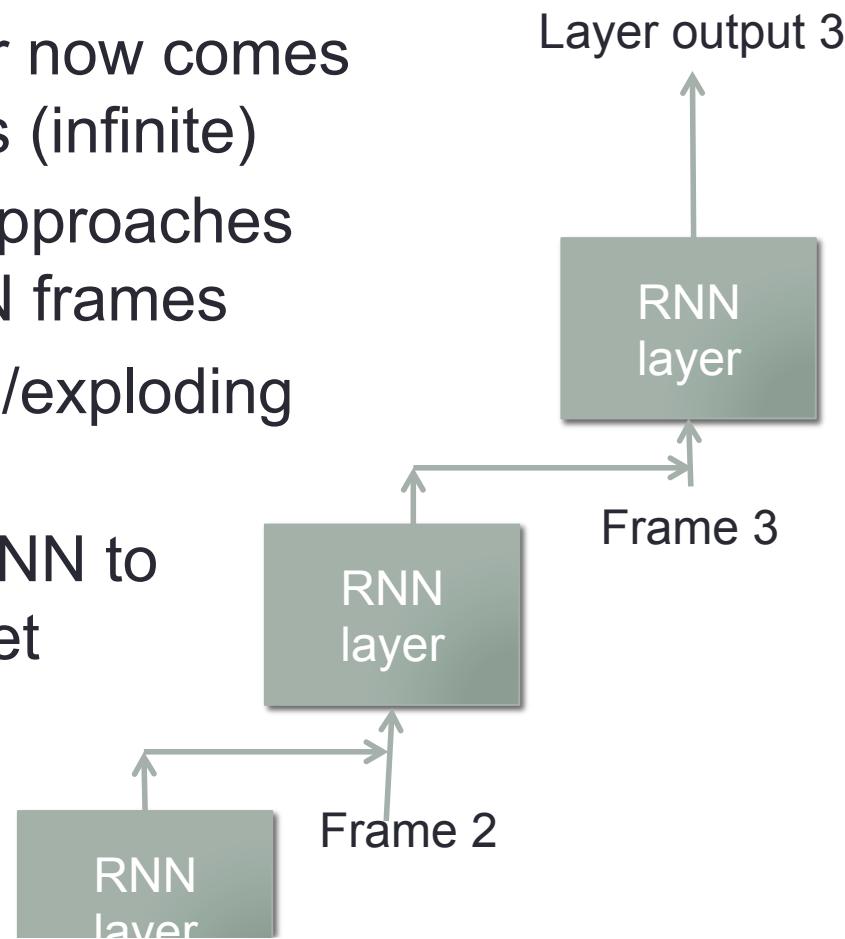
# Recurrent neural network (RNN)

- RNN framework



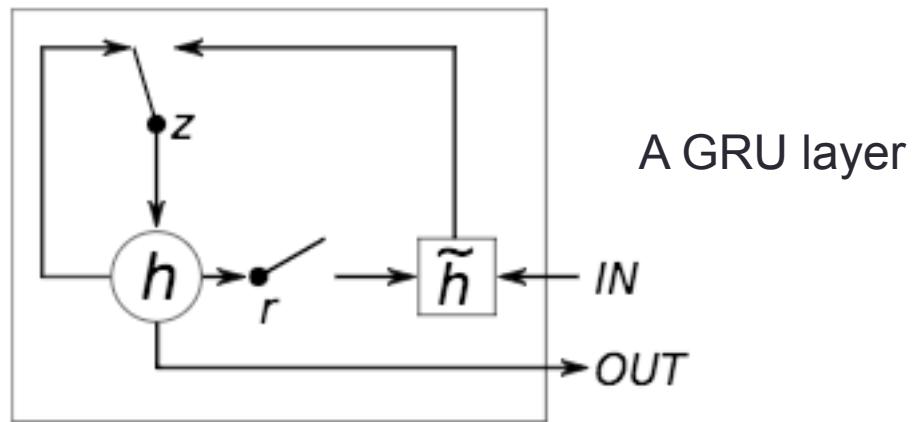
# Back propagation through time

- RNN weights are now shared across time. Updating the weights is now trickier
- Updates for a layer now comes from multiple terms (infinite)
- Practical training approaches will truncate after N frames
- Gradient vanishing/exploding problem!
- Needs a way for RNN to know when to forget



# Gated recurrent unit (GRU)

- Forms a Gated Recurrent Neural Networks (GRNN)
- Add gates that can choose to reset ( $r$ ) or update ( $z$ )



$$z_t = \sigma_g(W_z x_t + U_z h_{t-1} + b_z)$$

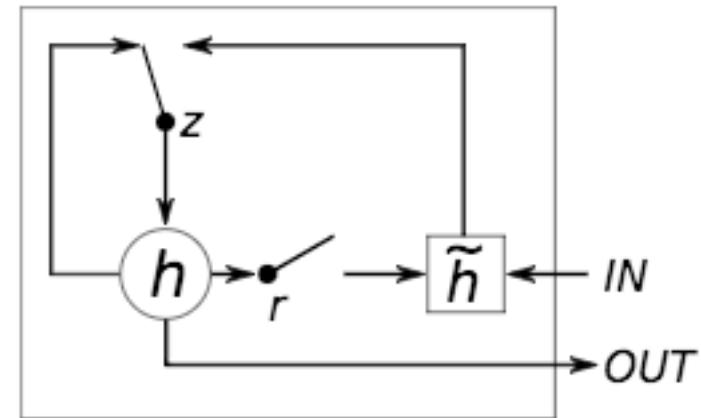
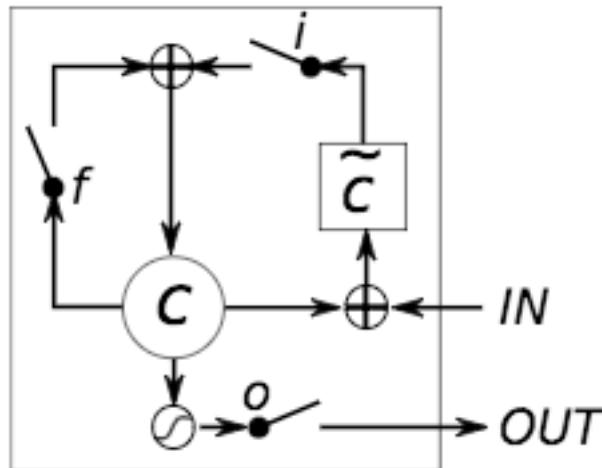
$$r_t = \sigma_g(W_r x_t + U_r h_{t-1} + b_r)$$

$$h_t = z_t \circ h_{t-1} + (1 - z_t) \circ \underbrace{\sigma_h(W_h x_t + U_h (r_t \circ h_{t-1}) + b_h)}_{\tilde{h}}$$

$\tilde{h}$

# LSTMs

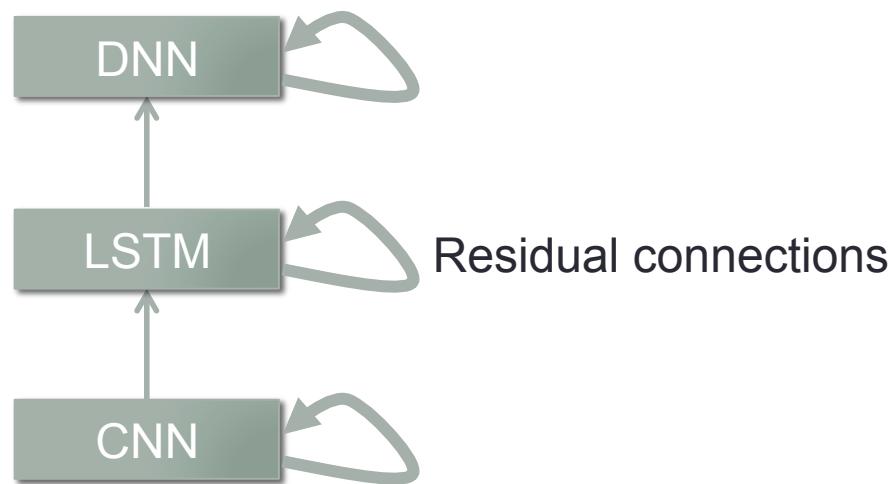
- Have 3 gates, forget ( $f$ ), input ( $i$ ), output ( $o$ )
- Has a memory cell ( $c$ ) that can be different from the output



- No clear conclusion whether LSTM or GRU is better yet
- For smaller data set, GRU seems to be better

# LSTMs

- Usually combined with CNN and DNN for real applications



# LSTM memories remember meaningful things

Cell sensitive to position in line:

```
The sole importance of the crossing of the Berezina lies in the fact  
that it plainly and indubitably proved the fallacy of all the plans for  
cutting off the enemy's retreat and the soundness of the only possible  
line of action--the one Kutuzov and the general mass of the army  
demanded--namely, simply to follow the enemy up. The French crowd fled  
at a continually increasing speed and all its energy was directed to  
reaching its goal. It fled like a wounded animal and it was impossible  
to block its path. This was shown not so much by the arrangements it  
made for crossing as by what took place at the bridges. When the bridges  
broke down, unarmed soldiers, people from Moscow and women with children  
who were with the French transport, all--carried on by vis inertiae--  
pressed forward into boats and into the ice-covered water and did not,  
surrender.
```

Cell that turns on inside quotes:

```
"You mean to imply that I have nothing to eat out of.... On the  
contrary, I can supply you with everything even if you want to give  
dinner parties," warmly replied Chichagov, who tried by every word he  
spoke to prove his own rectitude and therefore imagined Kutuzov to be  
animated by the same desire.
```

```
Kutuzov, shrugging his shoulders, replied with his subtle penetrating  
smile: "I meant merely to say what I said."
```

Cell that robustly activates inside if statements:

```
static int __dequeue_signal(struct sigpending *pending, sigset_t *mask,  
    siginfo_t *info)  
{  
    int sig = next_signal(pending, mask);  
    if (sig) {  
        if (current->notifier) {  
            if (sigismember(current->notifier_mask, sig)) {  
                if (! (current->notifier)(current->notifier_data)) {  
                    clear_thread_flag(TIF_SIGPENDING);  
                    return 0;  
                }  
            }  
            collect_signal(sig, pending, info);  
        }  
        return sig;  
    }  
}
```

A large portion of cells are not easily interpretable. Here is a typical example:

```
/* Unpack a filter field's string representation from user-space  
 * buffer. */  
char *audit_unpack_string(void **bufp, size_t *remain, size_t len)  
{  
    char *str;  
    if (!*bufp || (len == 0) || (len > *remain))  
        return ERR_PTR(-EINVAL);  
    /* Of the currently implemented string fields, PATH_MAX  
     * defines the longest valid length.
```

Cell that turns on inside comments and quotes:

```
/* Duplicate LSM field information. The lsm_rule is opaque, so
 * re-initialized. */
static inline int audit_dupe_lsm_field(struct audit_field *df,
    struct audit_field *sf)
{
    int ret = 0;
    char *lsm_str;
    /* our own copy of lsm_str */
    lsm_str = kstrdup(sf->lsm_str, GFP_KERNEL);
    if (unlikely(!lsm_str))
        return -ENOMEM;
    df->lsm_str = lsm_str;
    /* our own (refreshed) copy of lsm_rule */
    ret = security_audit_rule_init(df->type, df->op, df->lsm_str,
        [void *] &df->lsm_rule);
    /* Keep currently invalid fields around in case they
     * become valid after a policy reload. */
    if (ret == -EINVAL) {
        pr_warn("Audit rule for LSM %'ss' is invalid\n",
            df->lsm_str);
        ret = 0;
    }
    return ret;
}
```

Cell that is sensitive to the depth of an expression:

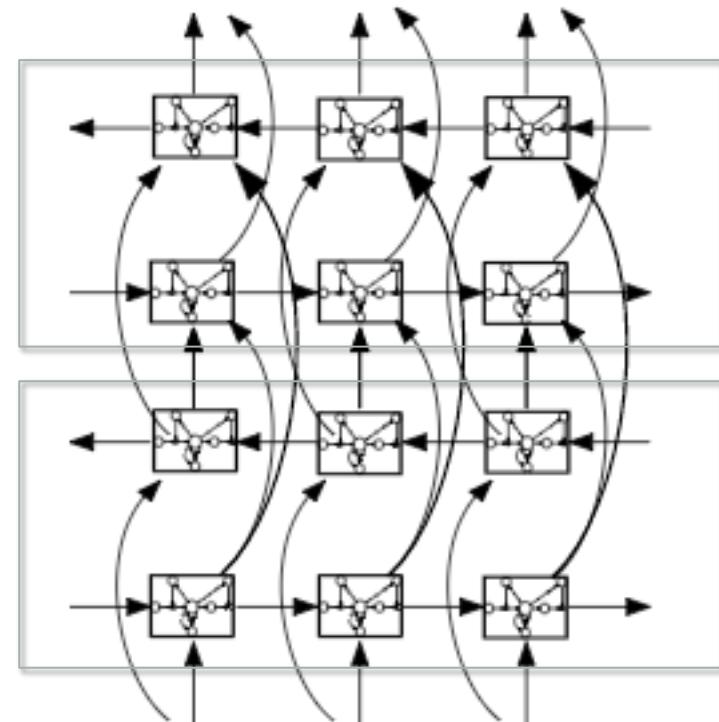
```
#ifdef CONFIG_AUDITSYSCALL
static inline int audit_match_class_bits(int class, u32 *mask)
{
    int i;
    if (classes[class]) {
        for (i = 0; i < AUDIT_BITMASK_SIZE; i++)
            if (mask[i] & classes[class][i])
                return 0;
    }
    return 1;
}
```

Cell that might be helpful in predicting a new line. Note that it only turns on for some "y":

```
char *audit_unpack_string(void **bufp, size_t *remain, size_t len)
{
    char *str;
    if (!*bufp || (len == 0) || (len > *remain))
        return ERR_PTR(-EINVAL);
    /* of the currently implemented string fields, PATH_MAX
     * defines the longest valid length.
     */
    if (len > PATH_MAX)
        return ERR_PTR(-ENAMETOOLONG);
    str = kmalloc(len + 1, GFP_KERNEL);
    if (unlikely(!str))
        return ERR_PTR(-ENOMEM);
    memcpy(str, *bufp, len);
    str[len] = 0;
    *bufp += len;
    *remain -= len;
    return str;
}
```

# Bi-directional LSTMs

- Uni-directional LSTMs only consider information from the past
- Bi-directional LSTMs consider information from the future and the past
- For each bi-directional layer, there are two LSTMs, one from the past, one from the future.

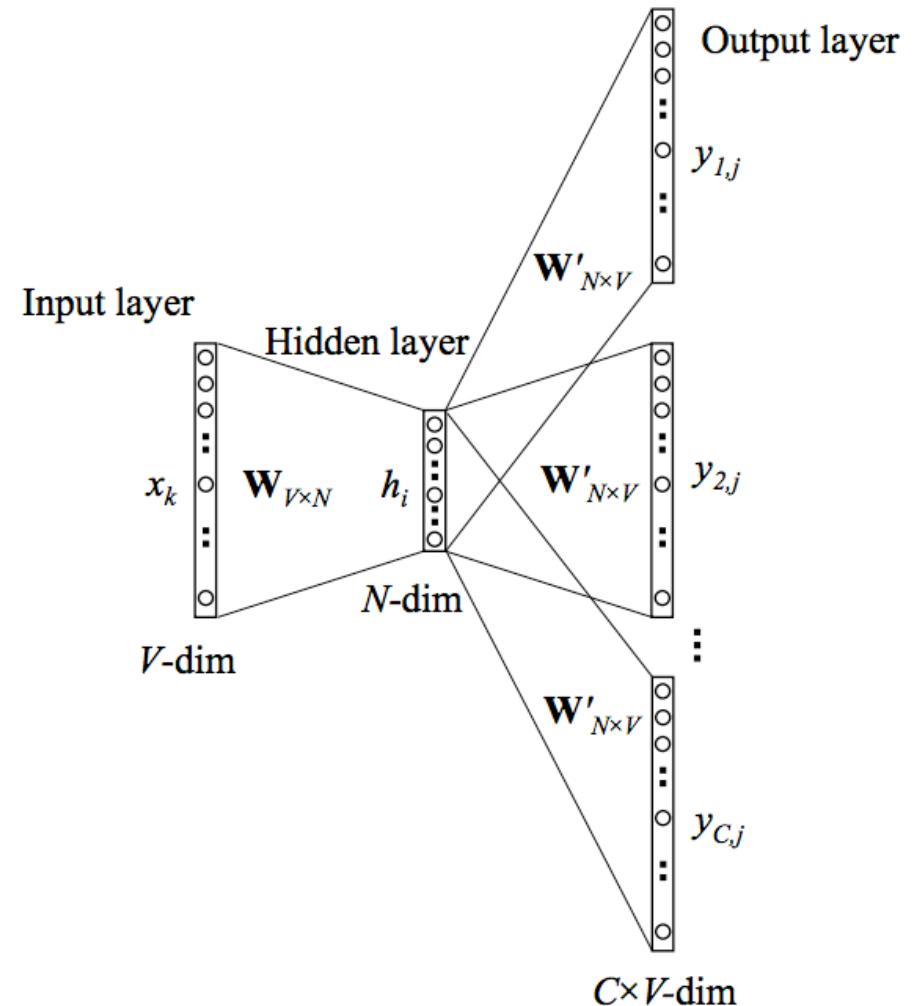


# Encoder-decoder

- We know neural networks can learn representations internally
- Can we use those internal representations?

# Word2Vec

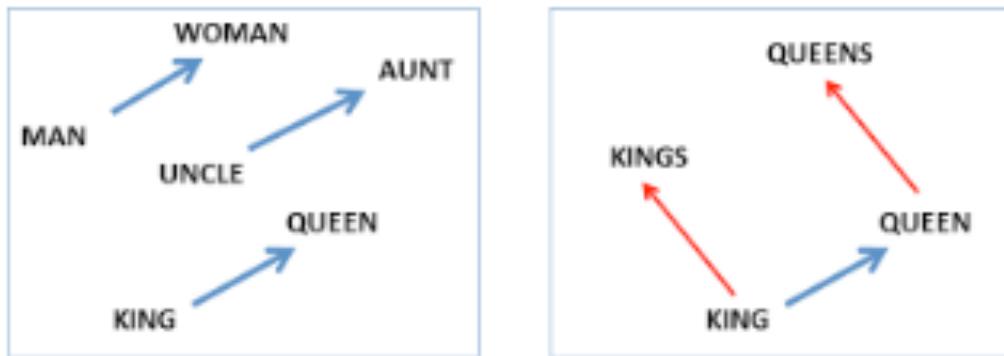
- We can understand the meaning of a word from the context
- Train a neural network that predicts what are likely words around an input word
- Use the hidden activations as features



“A quick brown XXX jumps over the fence”

# Word2Vec

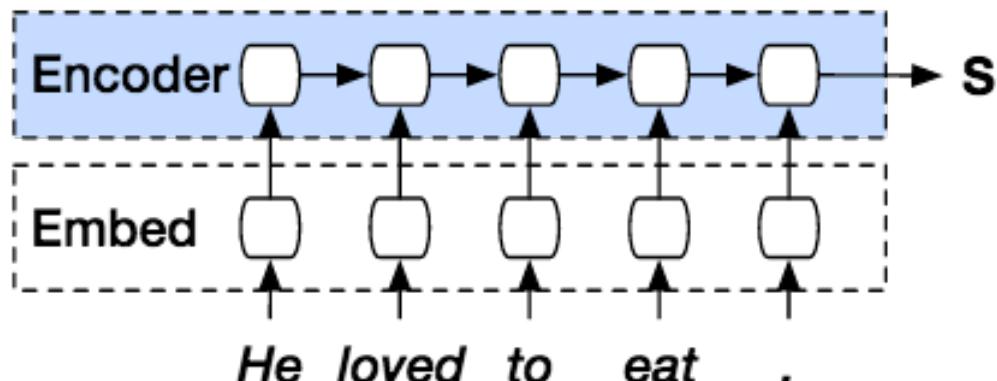
- The vector representations (**embeddings**) has interesting properties
- Similar meaning words are near each other
- You can also add and subtract meanings
- Word2vec is now used in most NLP tasks as an input feature



Mikolov, Linguistic Regularities in Continuous Space Word Representations, 2013

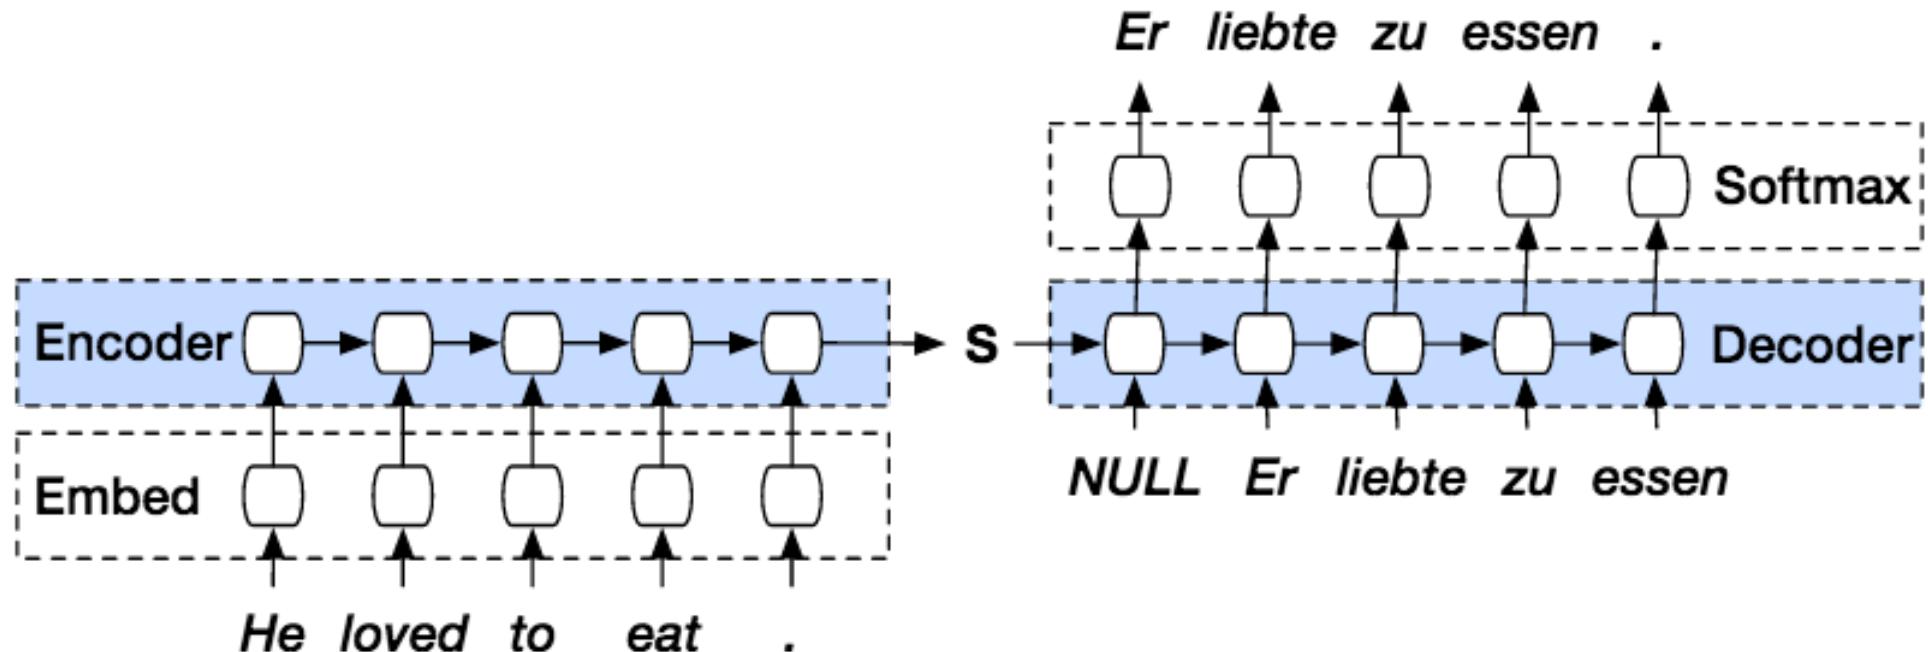
# Sentence encoder

- How can we represent sentences?
- Similarly, use the internal states as the representation
  - LSTMs/GRUs are good for sequence task
  - Use LSTMs/GRUs memory cells outputs as representation



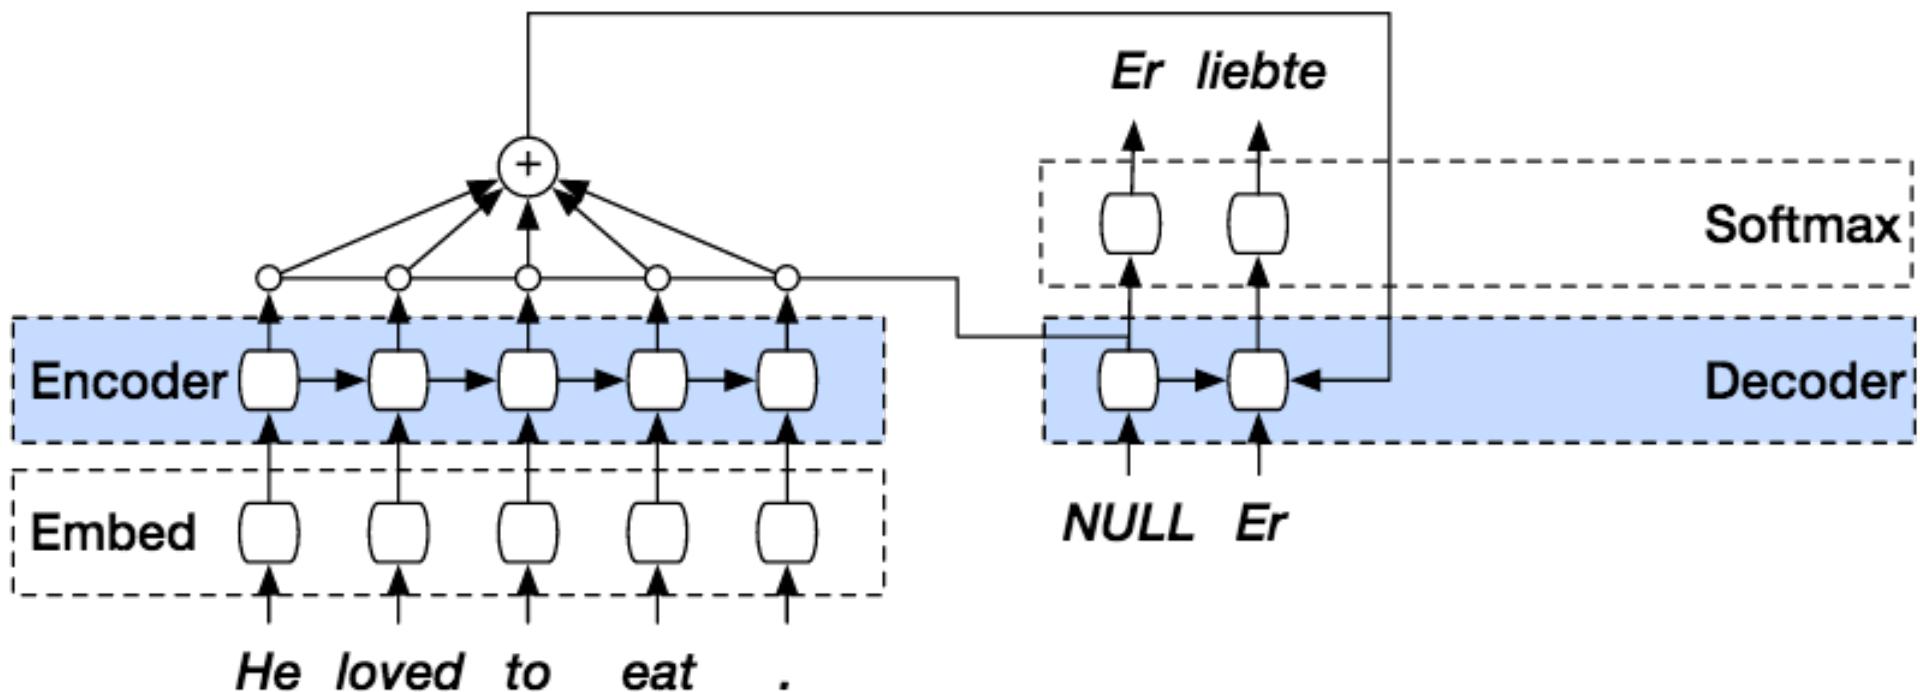
# Machine translation using encoder-decoder

- Use another LSTM as the decoder
- Input to the LSTM is the encoded sentence and the previously output word



# Attention model

- If the sentence is long, the model cannot remember long enough
  - Also problem of vanishing gradients
- Let the model pick where to focus



Dzmitry Bahdanau, Neural Machine Translation by Jointly Learning to Align and Translate, <https://arxiv.org/abs/1409.0473>

# Attention model for summarization

- Attention model can help with long context sentences for NLP

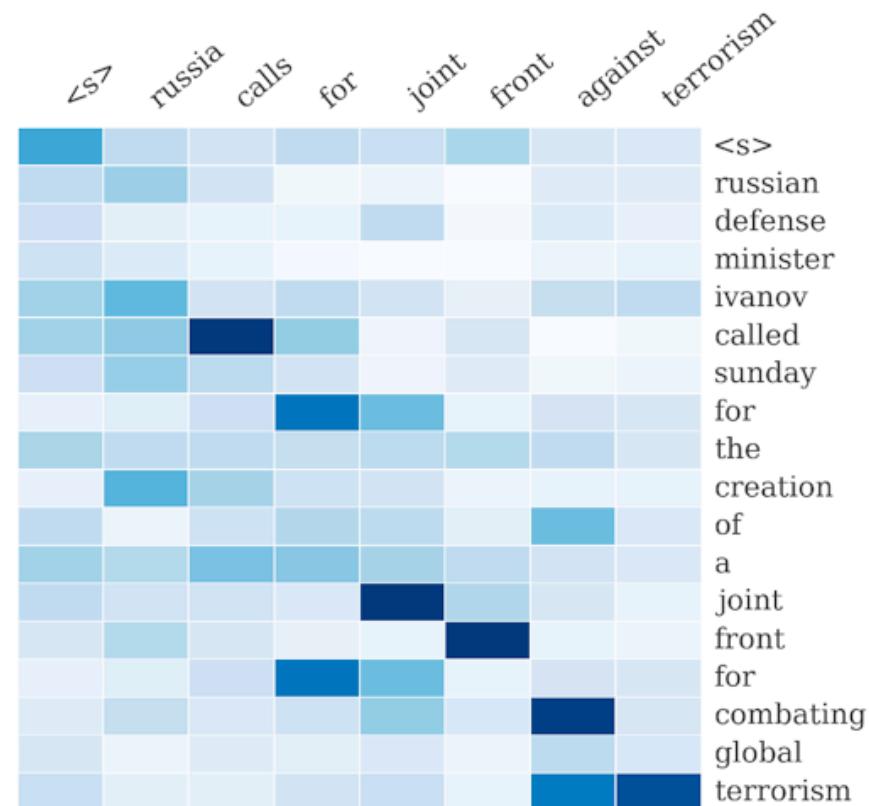


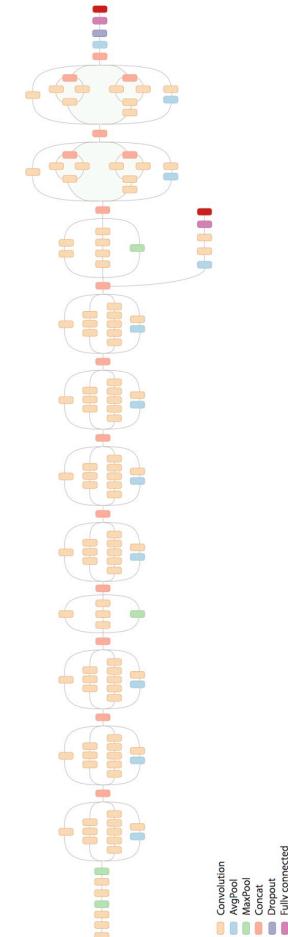
Figure 1: Example output of the attention-based summarization (ABS) system. The heatmap represents a soft alignment between the input (right) and the generated summary (top). The columns represent the distribution over the input after generating each word.

# Transfer Learning

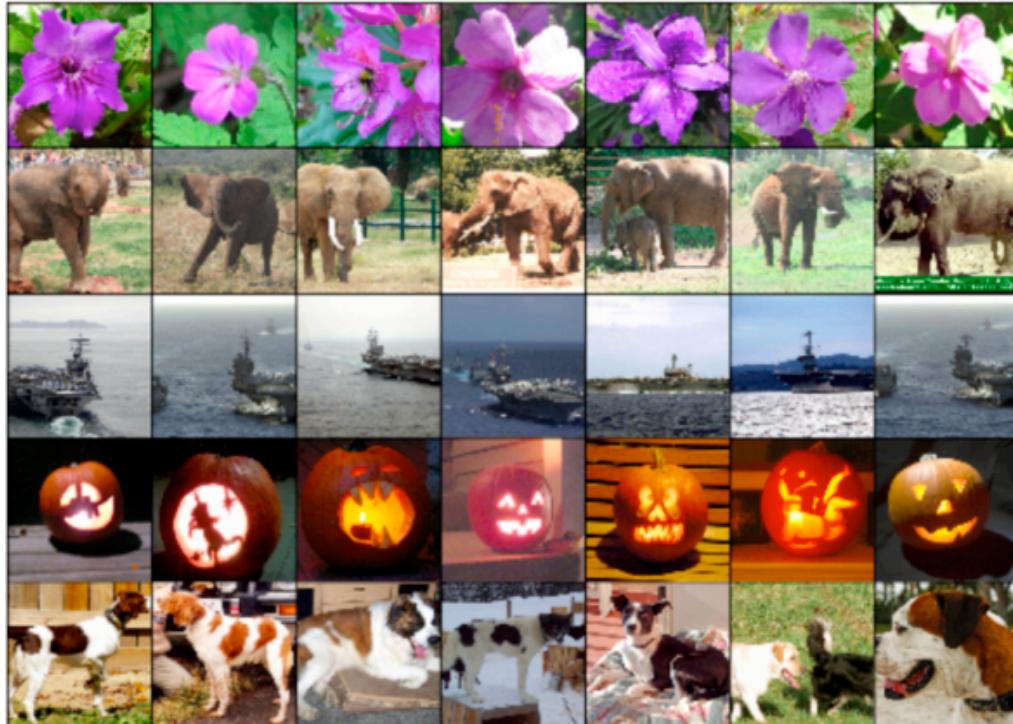
- Neural networks learns useful information about the data
    - Hierarchical features
  - These features are useful for other tasks
- 
- Learn on a big data set and apply on small data set
  - Many pre-trained models available
- 
- Replace softmax layer with your task

# Transfer Learning

Inception model



Imagenet : millions of pictures, 1000+ object classes



<http://image-net.org/>

<https://github.com/tensorflow/models/tree/master/inception>

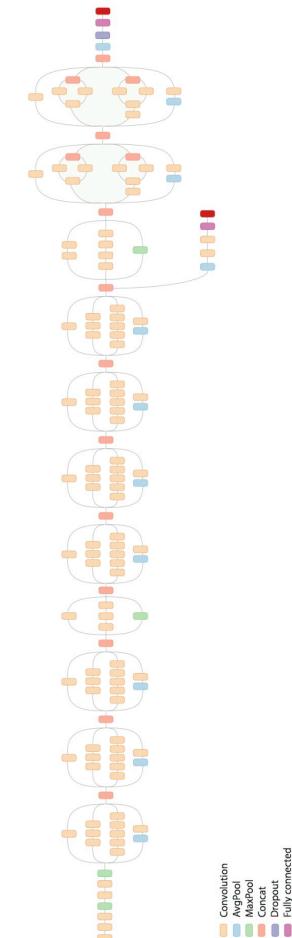
# Pose Estimation of Anime Characters



Pramook Khungurn, Pose estimation of anime/manga characters: a case for synthetic data

# Transfer Learning Note

- Remove more **layers at the top** if you believe lower level representations are more relevant to your task
- If you believe there is change in sensors, add **layers at the bottom**
  - Using a different microphone from the original set

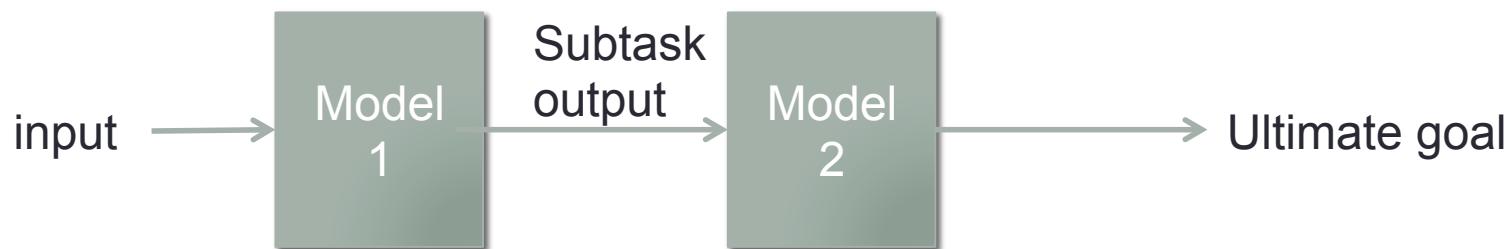


# Trends in Deep learning

- End-to-end models
- Multimodels (General purpose model)
- Speeding up simulations
- Reinforcement learning
- One button machines

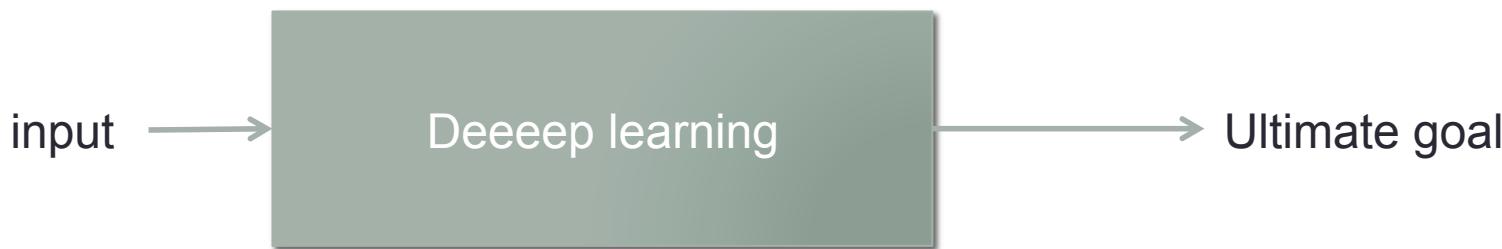
# End-to-End Models

- Traditional machine learning divides the main task into smaller sub-tasks
  - Speech-to-speech translation
    - L1 Speech -> L1 Text -> L2 Text -> L2 Speech
  - Text caption generation
    - Image segmentation -> objection recognition -> NLP
- Errors propagate
- Each part are not jointly optimized

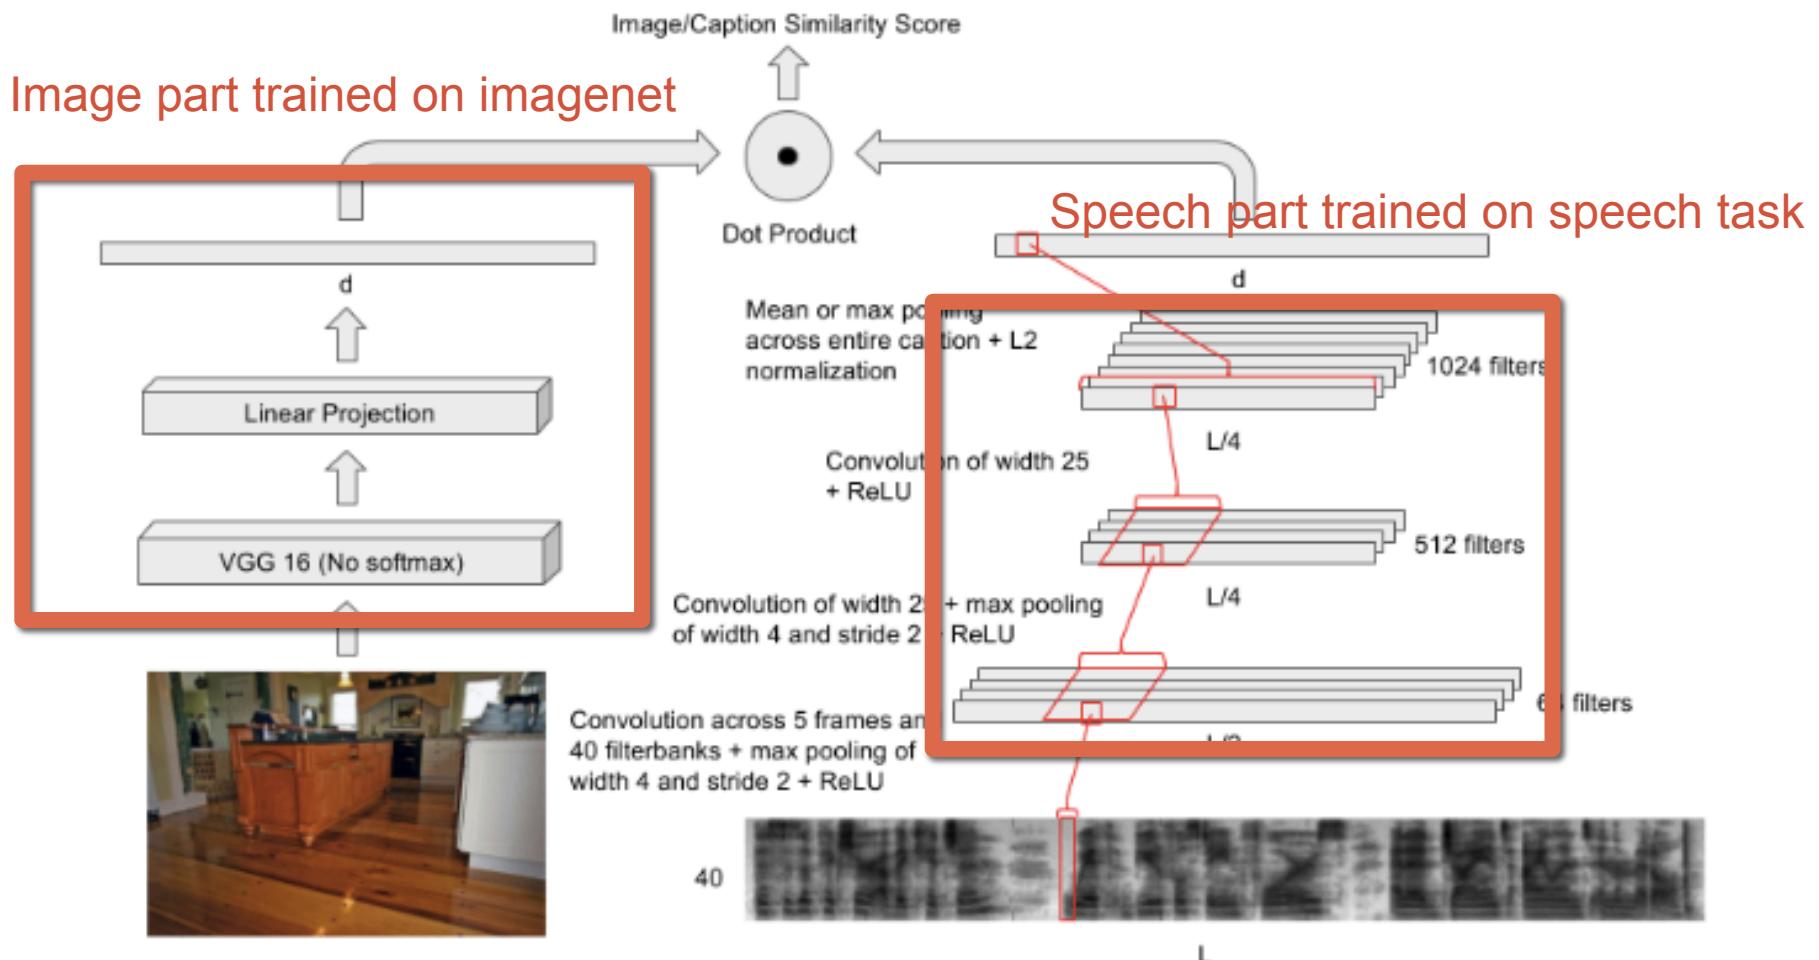


# End-to-End Models

- Traditional machine learning divides the main task into smaller sub-tasks
  - Speech-to-speech translation
    - L1 Speech -> L1 Text -> L2 Text -> L2 Speech
  - Text caption generation
    - Image segmentation -> objection recognition -> NLP
- Errors propagate
- Each part are not jointly optimized



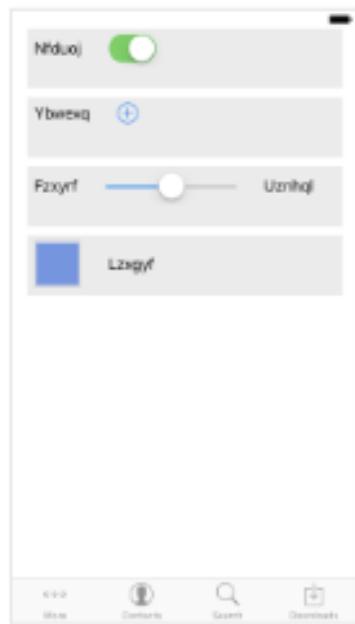
# Speech description for image querying



Whole model trained on a smaller parallel corpus

# Pix2Code

- Given image of UI, outputs platform specific code



input

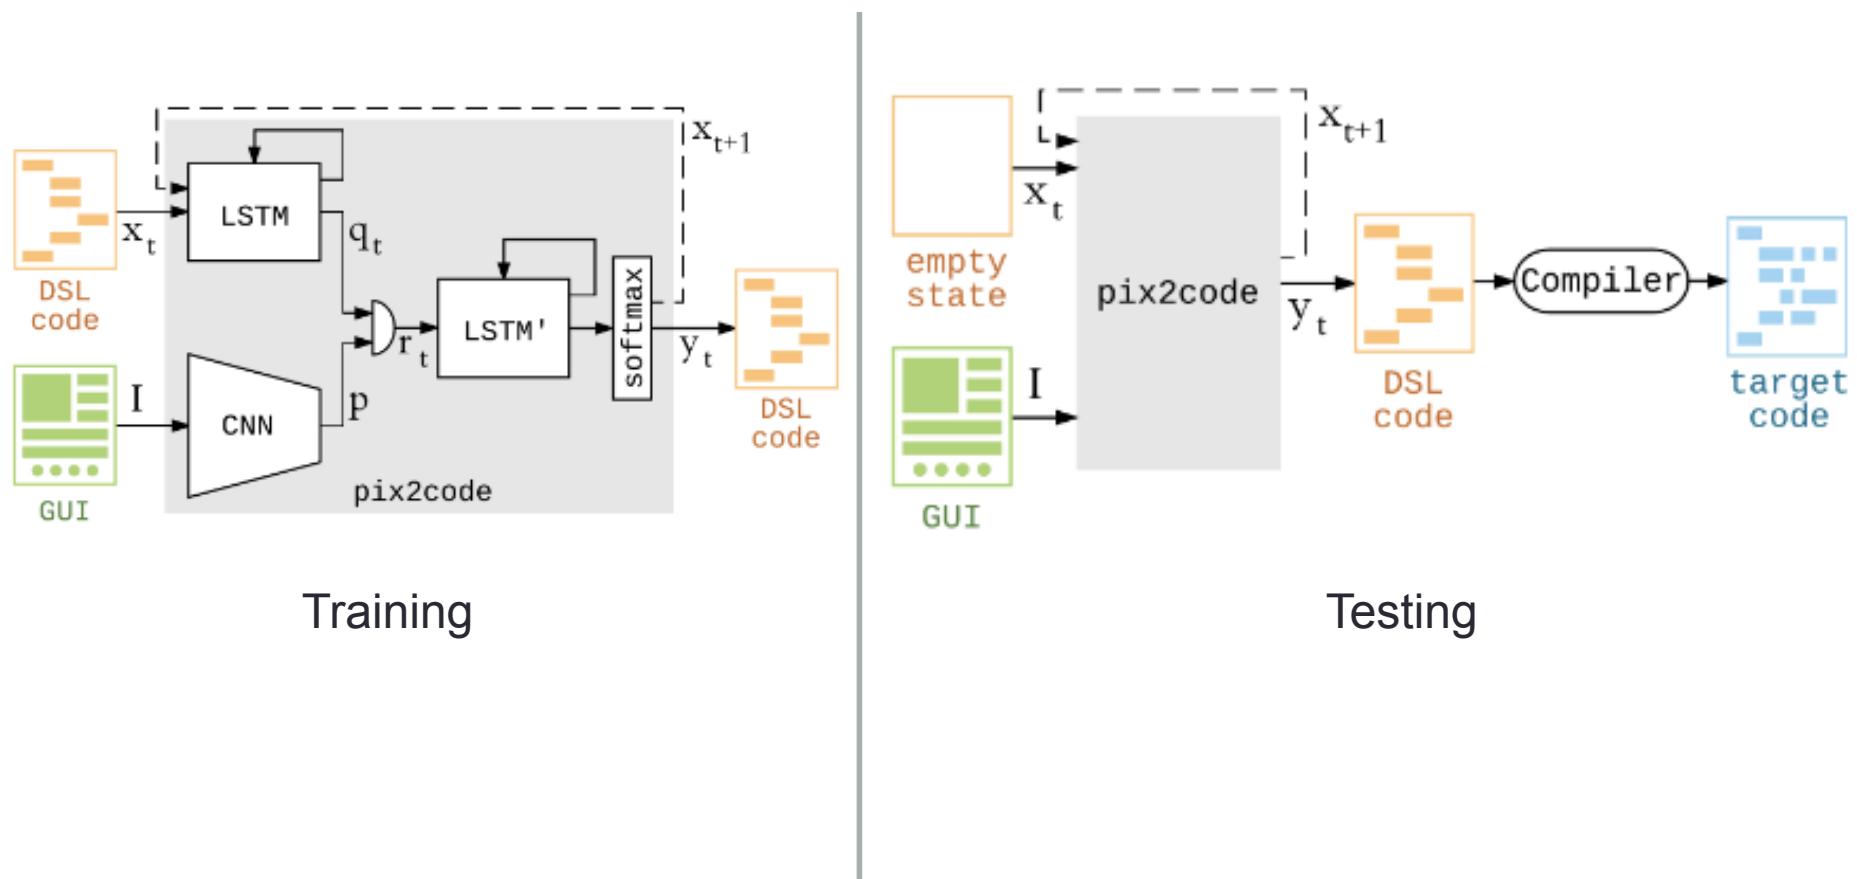


```
stack {
    row {
        label, switch
    }
    row {
        label, btn-add
    }
    row {
        label, slider, label
    }
    row {
        img, label
    }
}
footer {
    btn-more, btn-contact, btn-search, btn-download
}
```

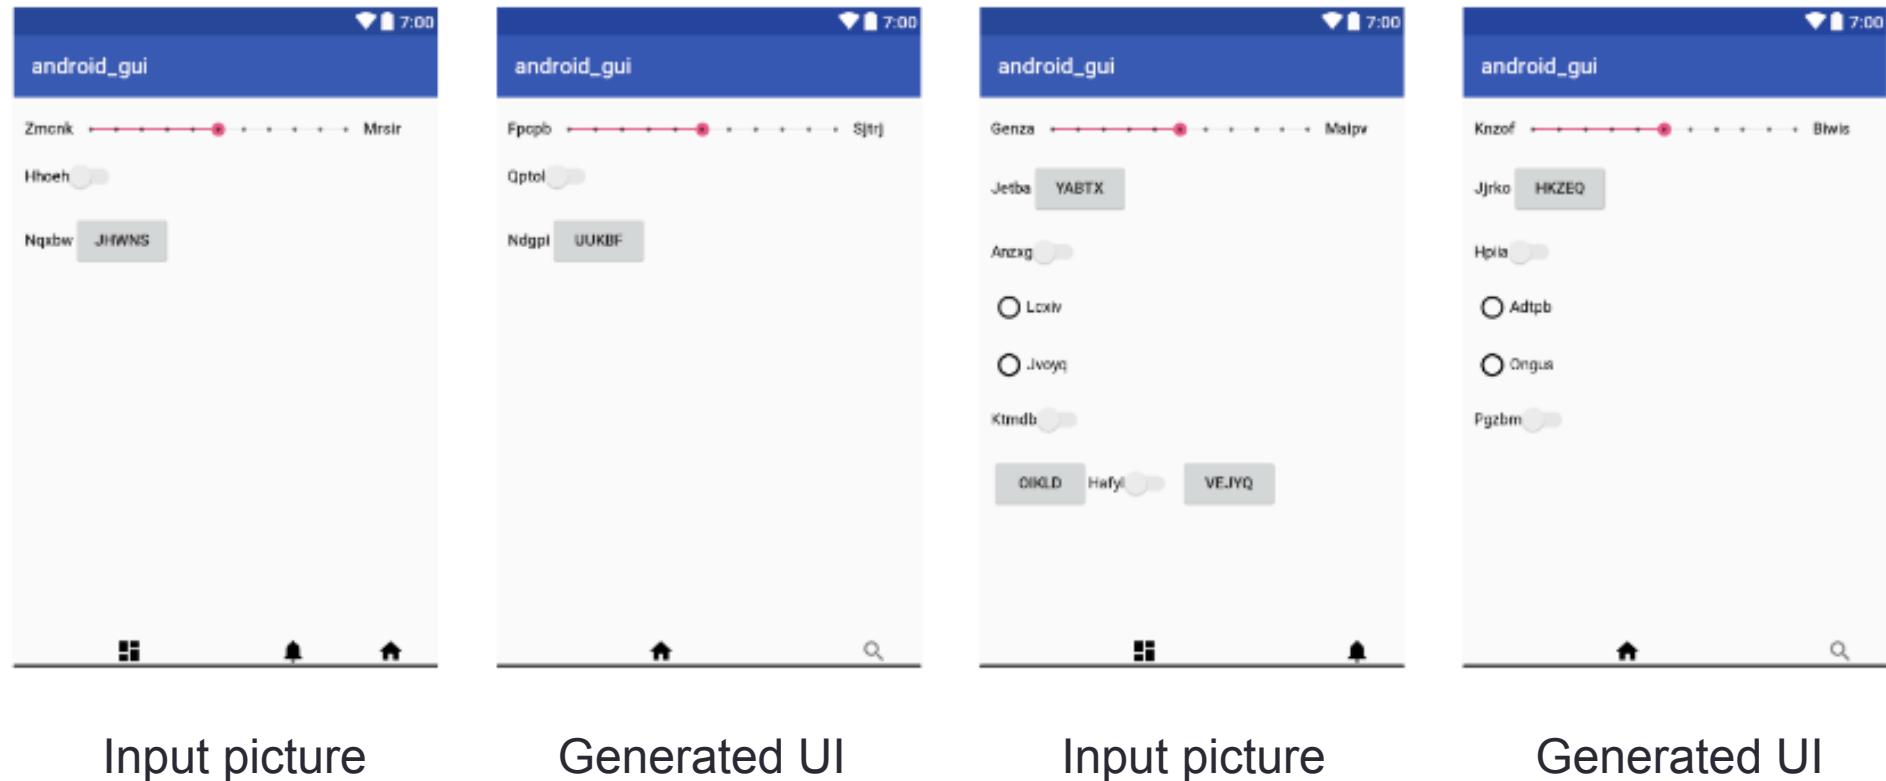
output

# Pix2Code

- Similar to the machine translation encoder-decoder but with image recognition as the encoder



# Pix2Code



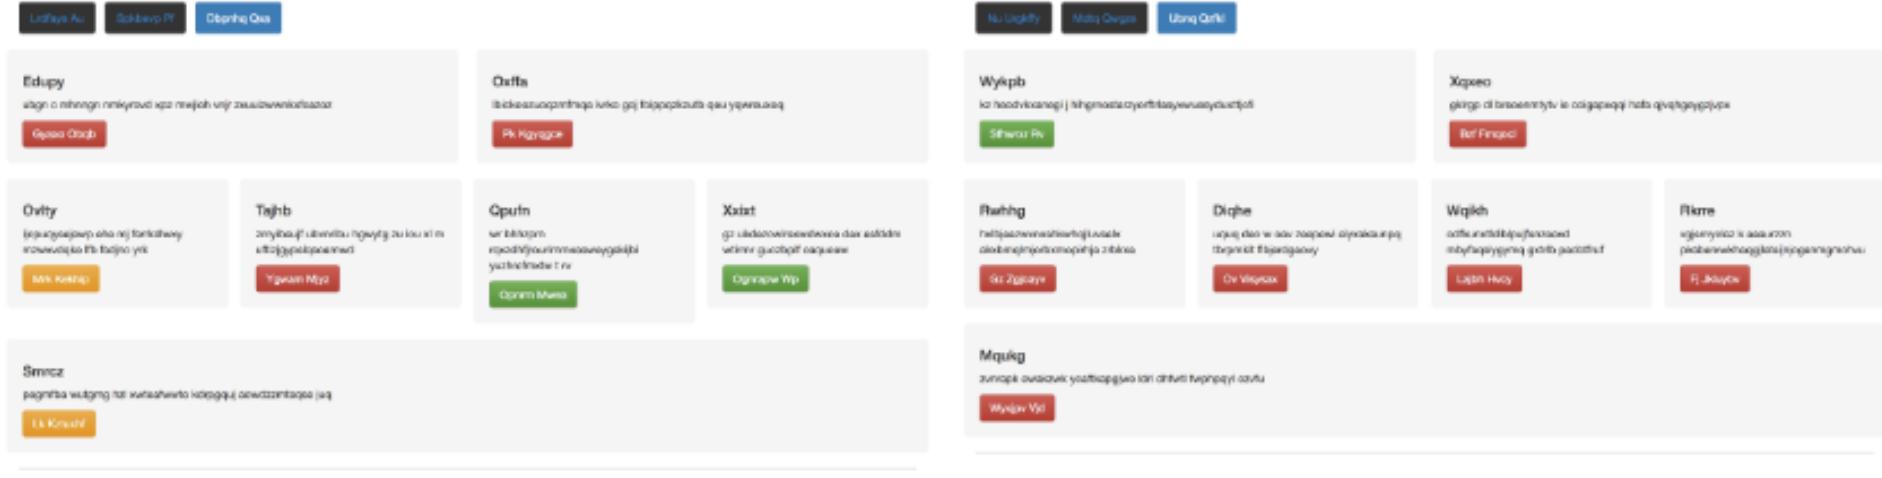
Input picture

Generated UI

Input picture

Generated UI

# Pix2Code

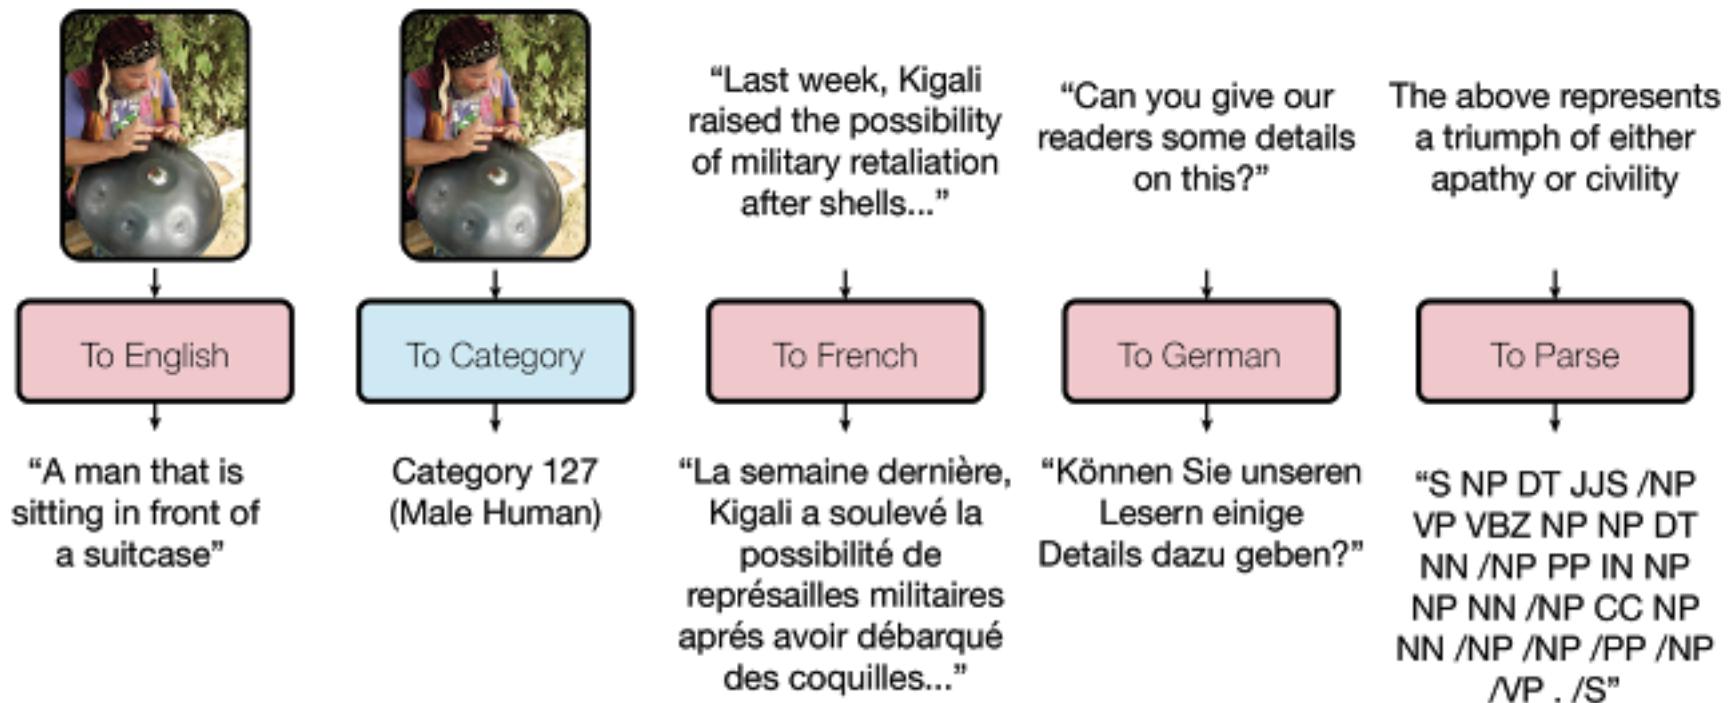


Input picture

Generated UI

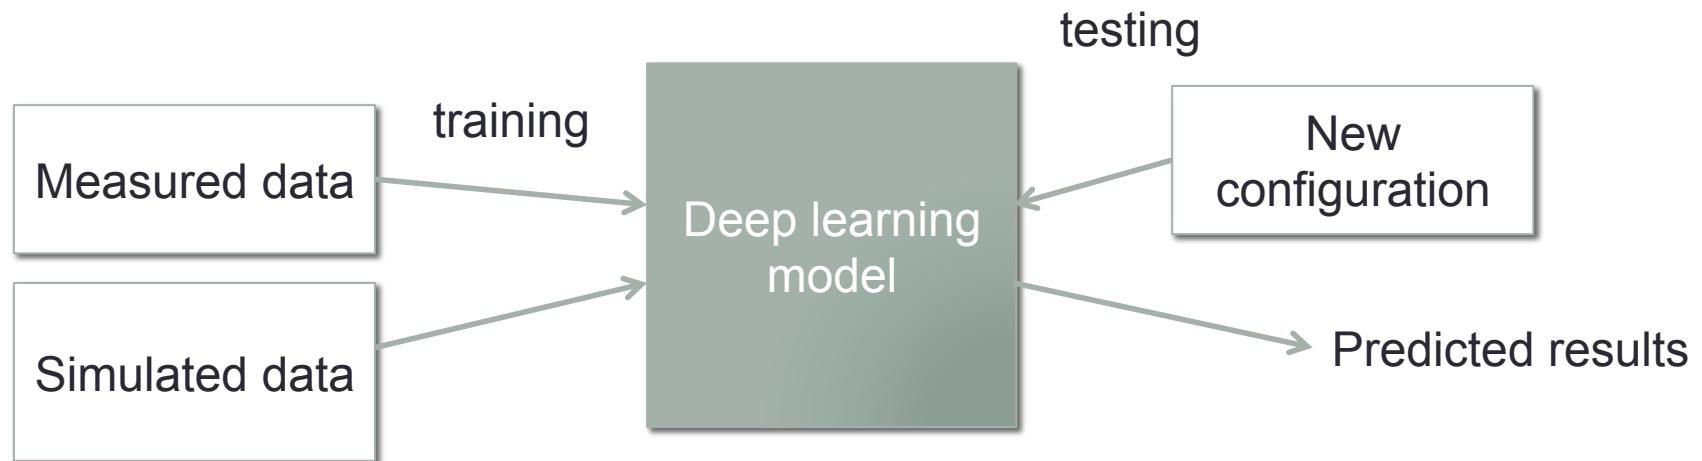
# Multimodels

- A single model that can handle different kinds of inputs
- Joint representation of meaning improves performance



# Speeding up simulations in physical science

- Several phenomena can be predicted by existing models in physics/chemistry/...
- Some models are slow
- Let deep learning learns the equations and give faster predictions at test time



# Accelerating Eulerian Fluid Simulation with Convolutional Networks

- <https://www.youtube.com/watch?v=9pXubxPsIJw>

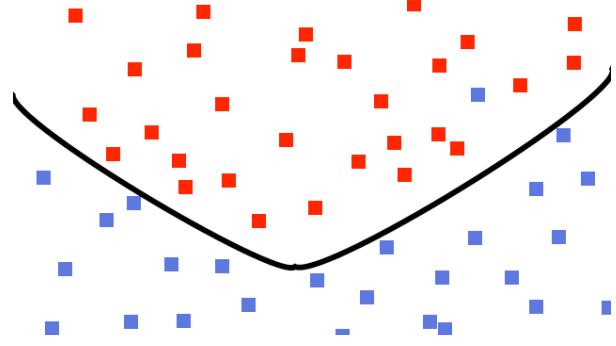


# Reinforcement learning

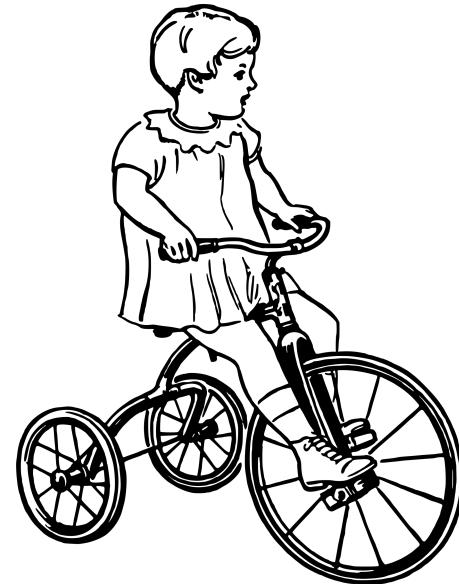
- Learning from experience



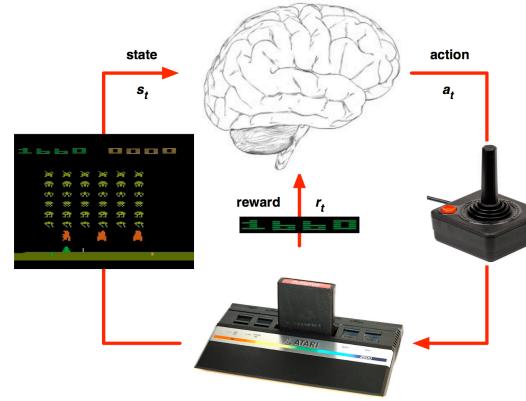
# Reinforcement learning



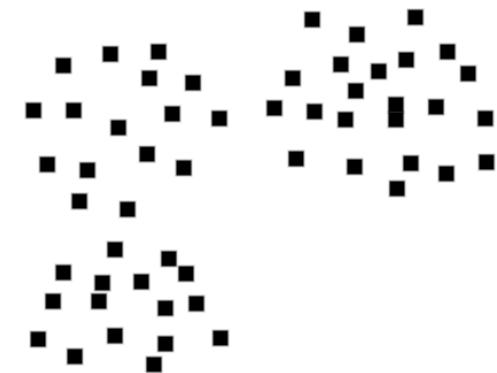
Supervised Learning



Reinforcement Learning



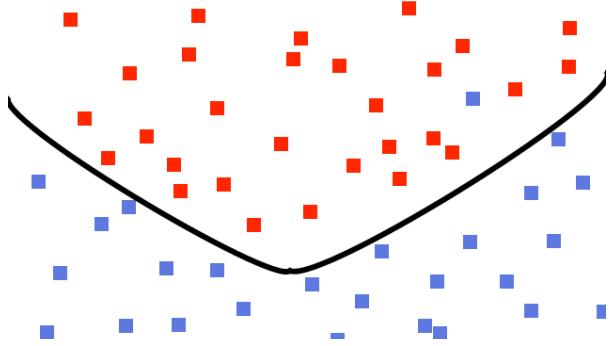
<https://keon.io/deep-q-learning/>



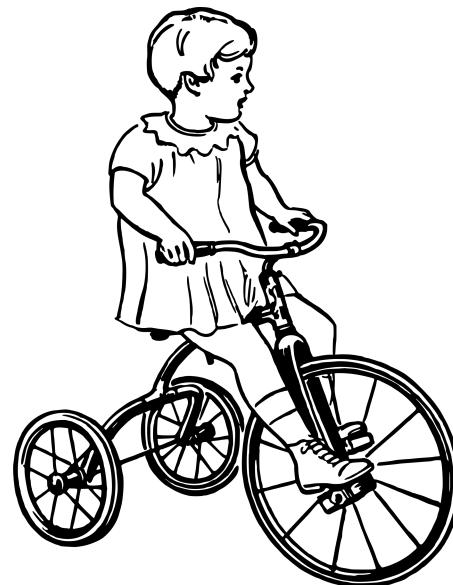
Unsupervised Learning



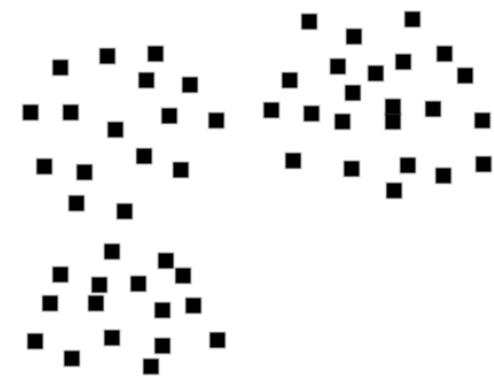
# Reinforcement learning



Supervised Learning



Reinforcement Learning



Unsupervised Learning

- Reinforcement learning tries to learn the effect of each action and how it effect the outcome
- Reinforcement learning has been around for a long time
- Deep learning simplifies the learning process

# Reinforcement learning for robotics

- <https://www.youtube.com/watch?v=HbHqC8Himol>



- Or use human supervision to help guide

# One button machines

- Machine learning as a tool for non-experts
- Can a non-expert just provide the data and let the machine decides how to proceed

**DataRobot** PRODUCT ▾ SOLUTIONS ▾ EDUCATION ▾ ABOUT ▾ WE'RE HIRING! ▾  CONTACT US

---

- ① Upload your data
- ② Select the target variable
- ③ Build 100s of models in one click
- ④ Explore top models and get insights
- ⑤ Deploy best model and make predictions

Summary

What would you like to predict?

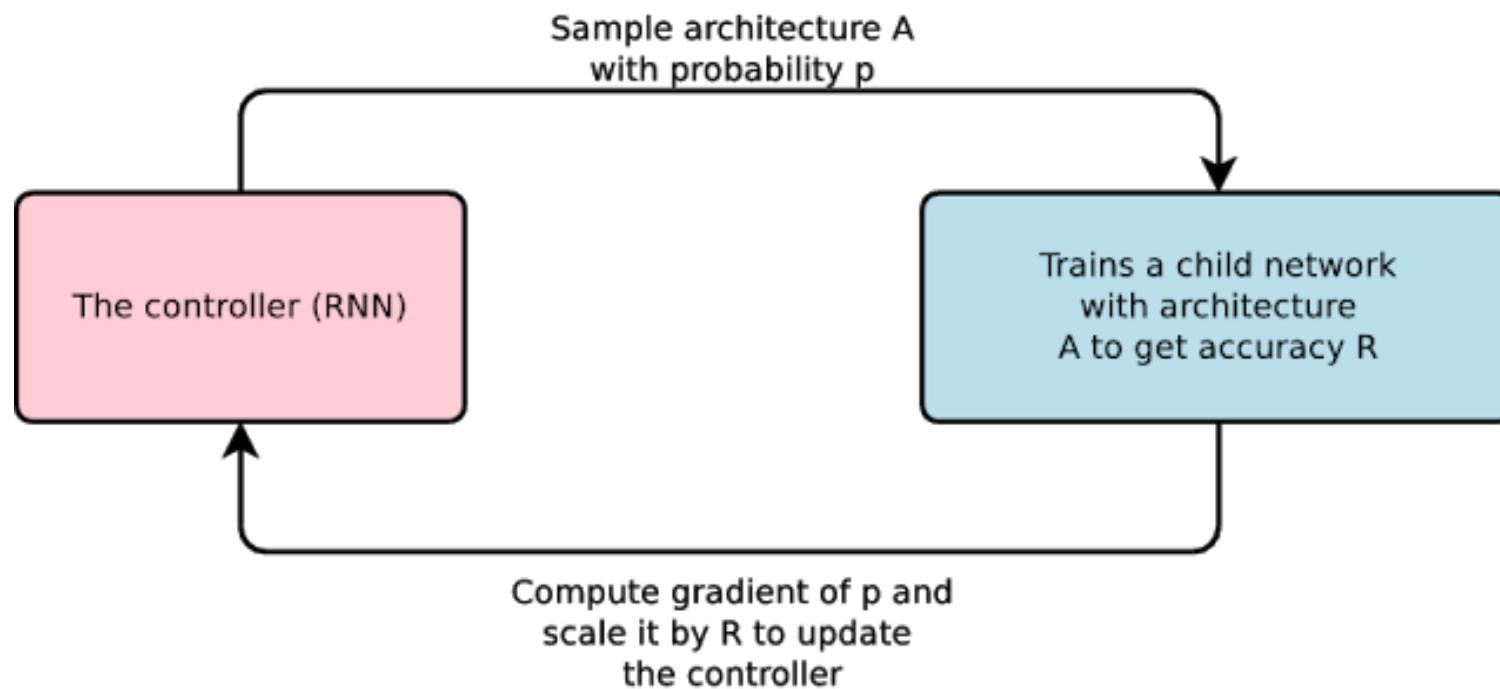
real

readmitted

Feature name	Var type	Unique	Missing	...
race	Categorical	5	221	
gender	Categorical	2	0	
age	Categorical	10	0	

# Reinforcement Learning for Model Selection

- Tuning a network takes time
- Let machine learning learns how to tune a network
- Matches or outperforms ML experts performance

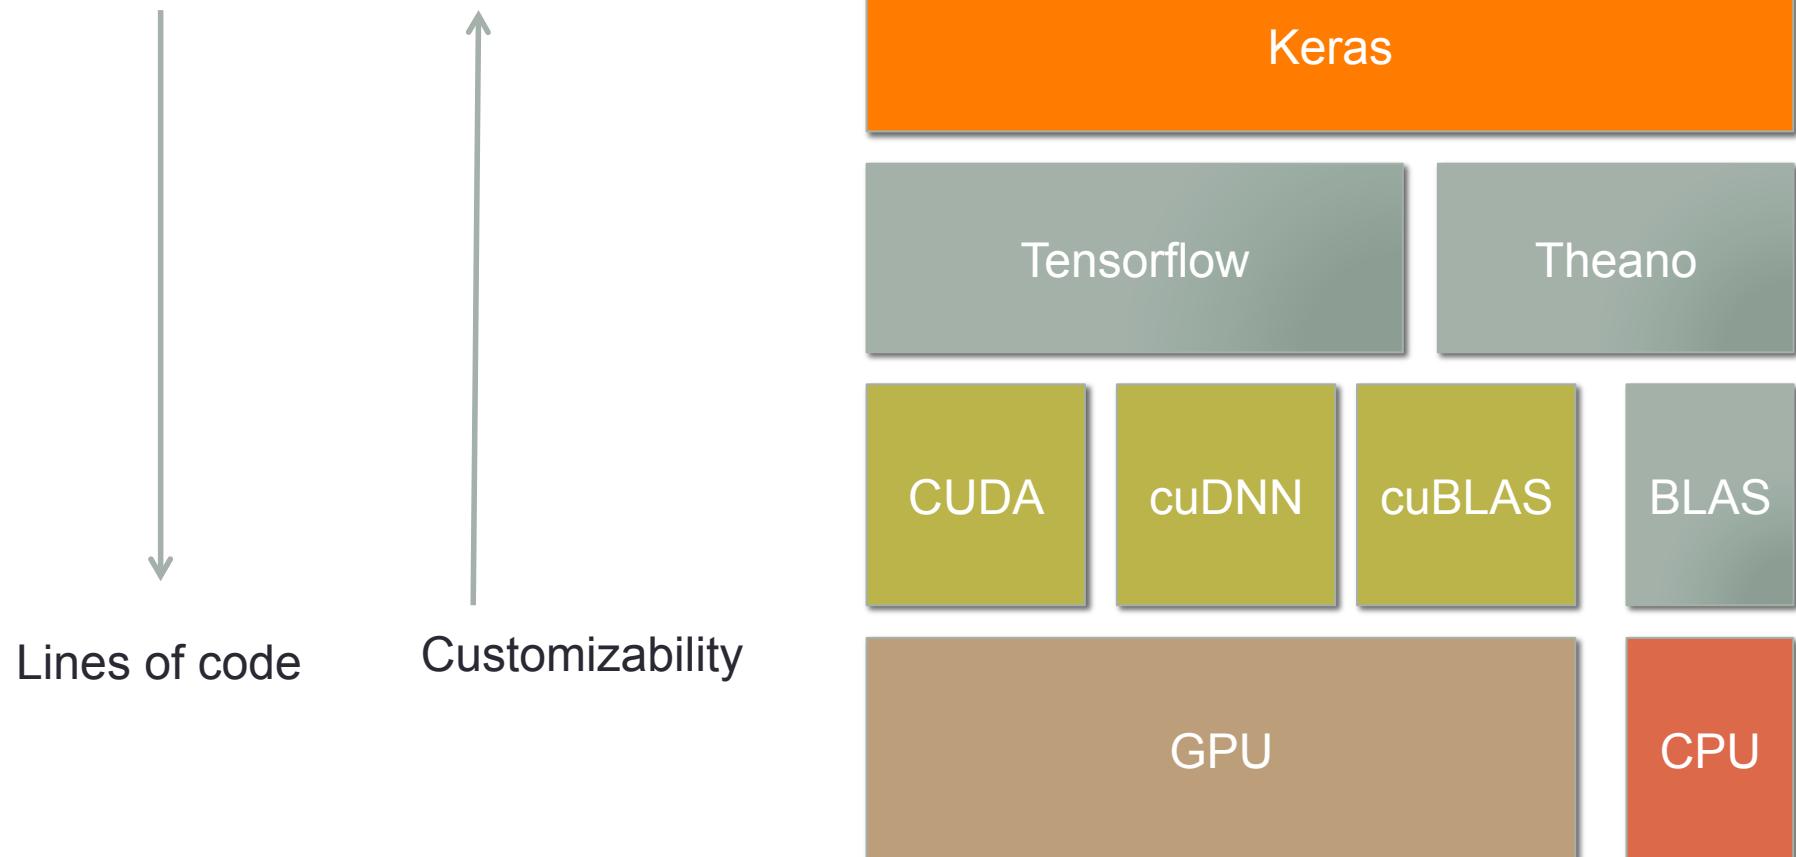


# What's a good candidate for DNN applications

- Anything human can complete in 1 second (Andrew Ng)
- Biggest bottleneck is still labeled data
  - Transfer learning
  - Multimodel learning
  - Unsupervised training (auto-encoders)
  - Crowd sourcing

# What toolkit

- Tensorflow – lower level
- Keras – higher level

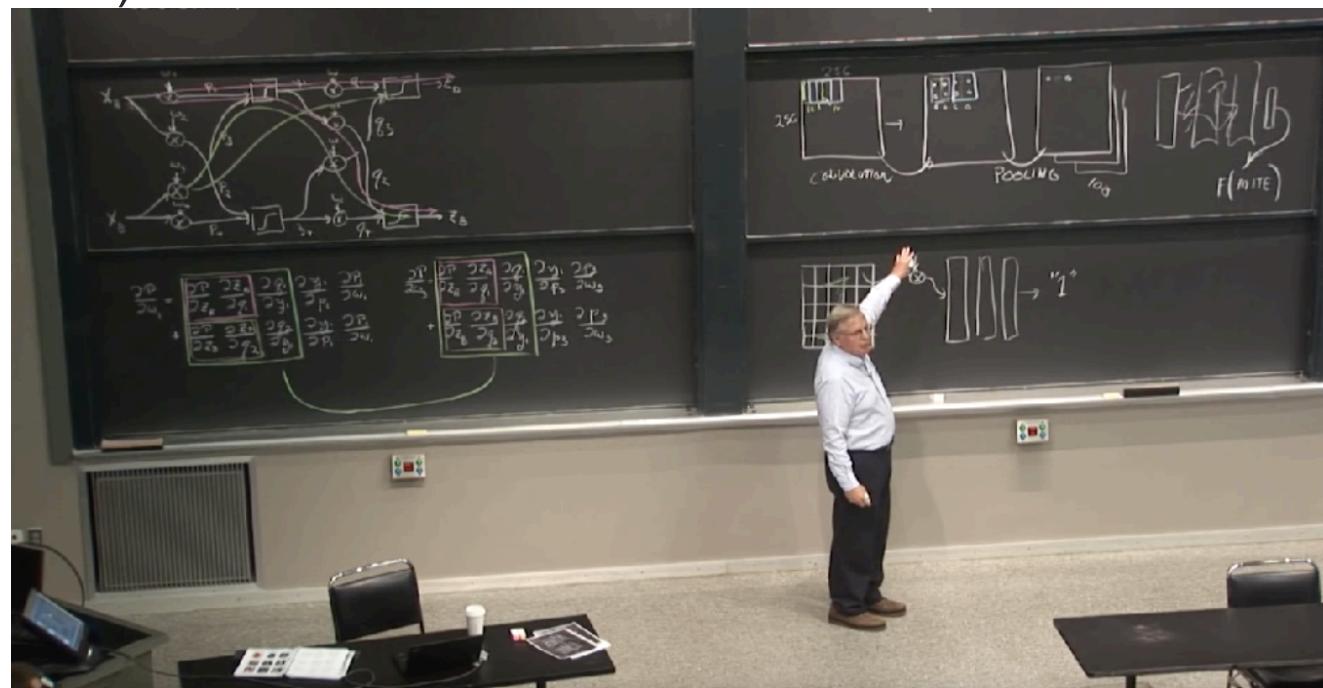


# Where to learn more

- Time commitment 2 hours:
  - Artificial Intelligence Complete Lectures (01-23)
  - The math in neural networks

<http://artificialbrain.xyz/artificial-intelligence-complete-lectures-01-23/>

(watch lecture 12a 12b)



# Where to learn more

- Time commitment 3 hours:
  - Tensorflow and deep learning - without a PhD by Martin Görner
  - <https://www.youtube.com/watch?v=vq2nnJ4g6N0>

The image is a composite of two parts. On the left, there is a screenshot of a code editor window titled "TensorFlow - run !". The code is written in Python and uses the TensorFlow library to train a model on the MNIST dataset. A handwritten note in red ink on the right side of the code says "running a Tensorflow computation, feeding placeholders". Below the code editor, a vertical line leads down to a "Tip" section which says "do this every 100 iterations". On the right side of the composite image, there is a video frame of a man in a white shirt and dark trousers standing on a stage and gesturing with his hands while speaking. The background of the video frame shows a screen with some text and logos.

```
sess = tf.Session()
sess.run(init)

for i in range(1000):
    # Load batch of images and correct answers
    batch_X, batch_Y = mnist.train.next_batch(100)
    train_data={X: batch_X, Y_: batch_Y}

    # train
    sess.run(train_step, feed_dict=train_data)

    # success ?
    a,c = sess.run([accuracy, cross_entropy], feed_dict=train_data)

    # success on test data ?
    test_data={X: mnist.test.images, Y_: mnist.test.labels}
    a,c = sess.run([accuracy, cross_entropy], feed=test_data)
```

running a Tensorflow computation, feeding placeholders

Tip  
do this every 100 iterations

= sess.run([accuracy, cross\_entropy], feed=

ccess on test data ?  
\_data={X: mnist.test.images, Y\_: mnist.test.labels}  
= sess.run([accuracy, cross\_entropy], feed=

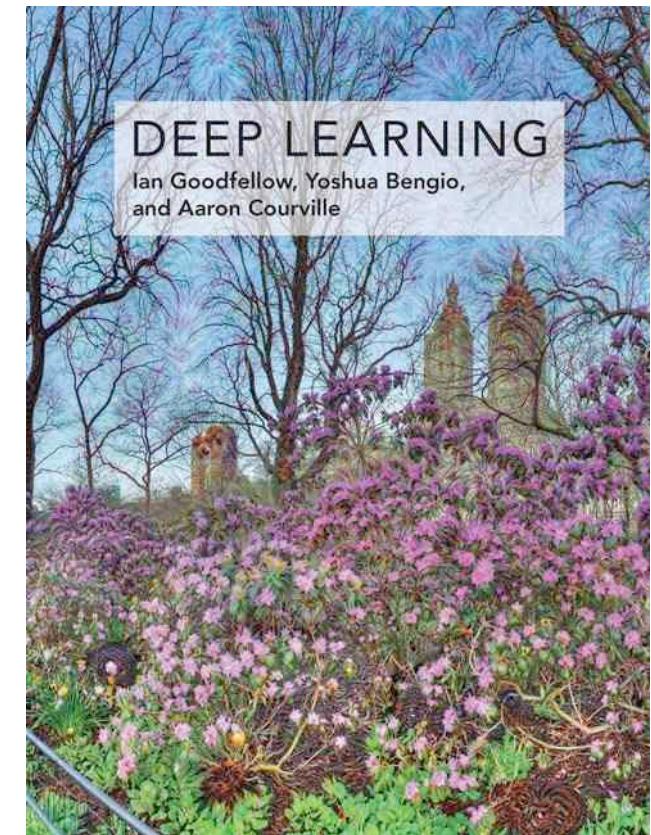
VOXX™

# Where to learn more

- Deep learning school
  - 2 day workshop with many talks and tutorial (videos on Youtube)
  - <https://www.bayareadl.school.org/>
- Fei-Fei Li and Andrew Karpathy's Computer Vision class (Stanford cs231n 2015/2016)
  - Full course about deep learning (vision oriented)
  - <http://cs231n.stanford.edu/>
- Richard Socher's Deep Learning and Natural Language Processing (Stanford cs224d)
  - Full course about deep learning (NLP oriented)
  - <http://cs224d.stanford.edu/>

# Where to learn more

- Hinton Neural networks course
  - If you like lots of math and really want to learn the theory
  - <https://www.coursera.org/learn/neural-networks>
- <http://www.deeplearningbook.org/>
  - The only deep learning book (currently)
  - Good coverage, quite terse



# Where to learn more

- For the latest developments subscribe to arxiv
  - Send email to [cs@arxiv.org](mailto:cs@arxiv.org)
  - Title: subscribe <name>
  - Content: add AI
- Reading group: TBA

# Demo

- Jupyter <http://jupyter.org/>
- <https://github.com/aymericdamien/TensorFlow-Examples>
- Picasso  
<https://medium.com/merantix/picasso-a-free-open-source-visualizer-for-cnns-d8ed3a35cfc5>