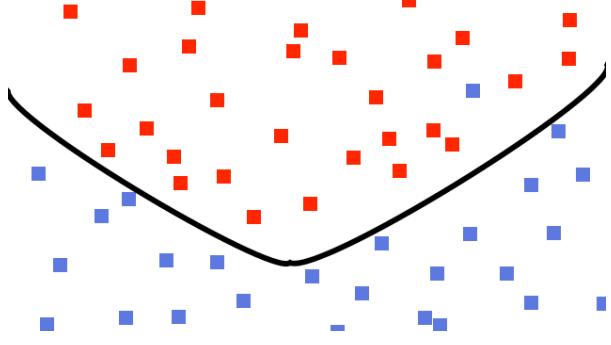


UNSUPERVISED LEARNING

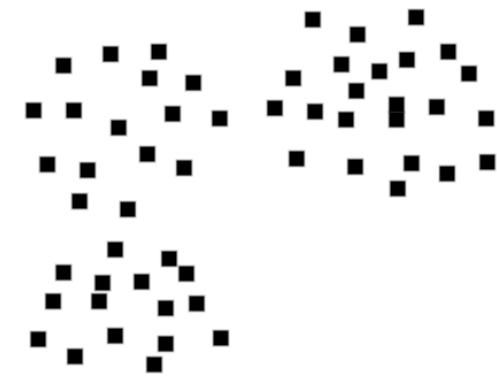
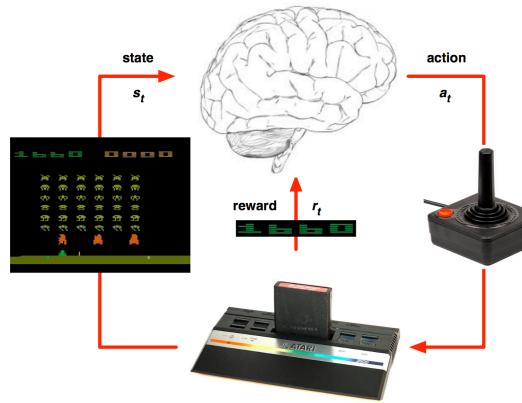
3 Modes of Learning



Supervised Learning



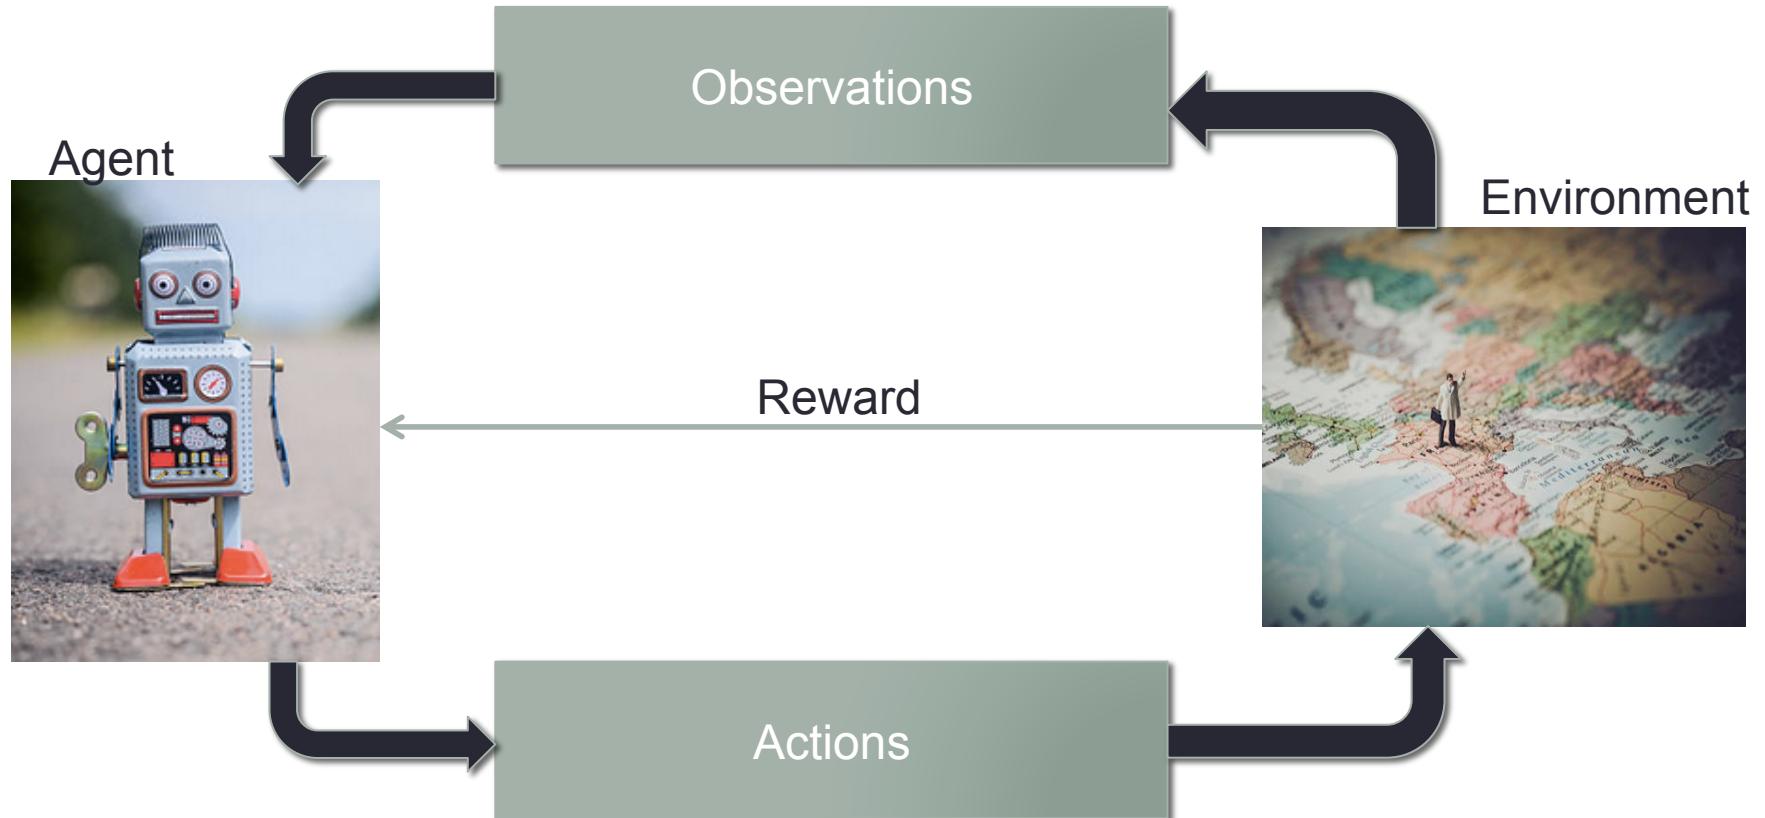
Reinforcement Learning



Unsupervised Learning



Reinforcement learning framework



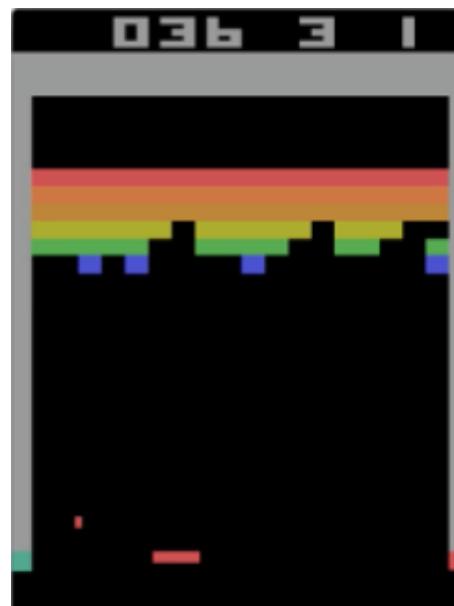
Learning through trial and error

Rewards-based learning

- Maximize the rewards



$$R_t = \Delta \text{distance}$$



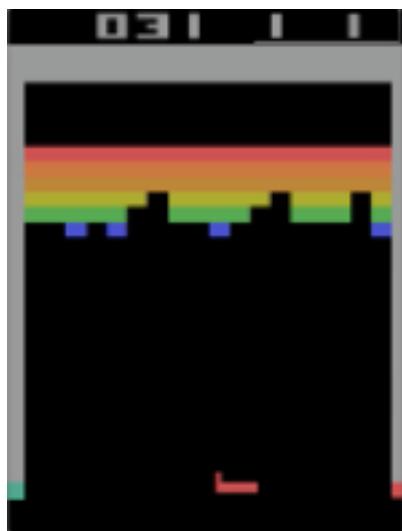
$$R_t = \Delta \text{score}$$



$$R_T = \begin{cases} 1 & , \text{win} \\ -1 & , \text{lose} \end{cases}$$

Credit assignment problem

- An action can have consequences further away in time
- Some movements might not have any effect on the outcome



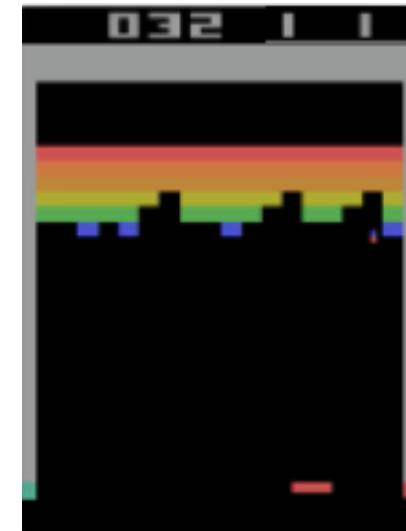
$$r_t = 0$$



10 time steps later



$$r_{t+10} = 1$$



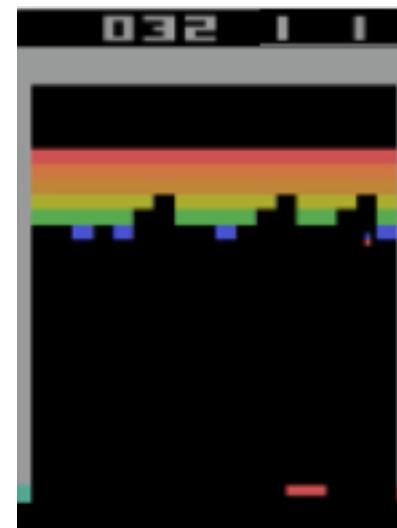
Cumulative rewards (Return)

- Maximize the expected **cumulative** rewards with discounting

$$G_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots = \sum_{t'=t}^{\infty} \gamma^{t'-t} r_t$$



$$r_t = 0$$



$$r_{t+10} = 1$$



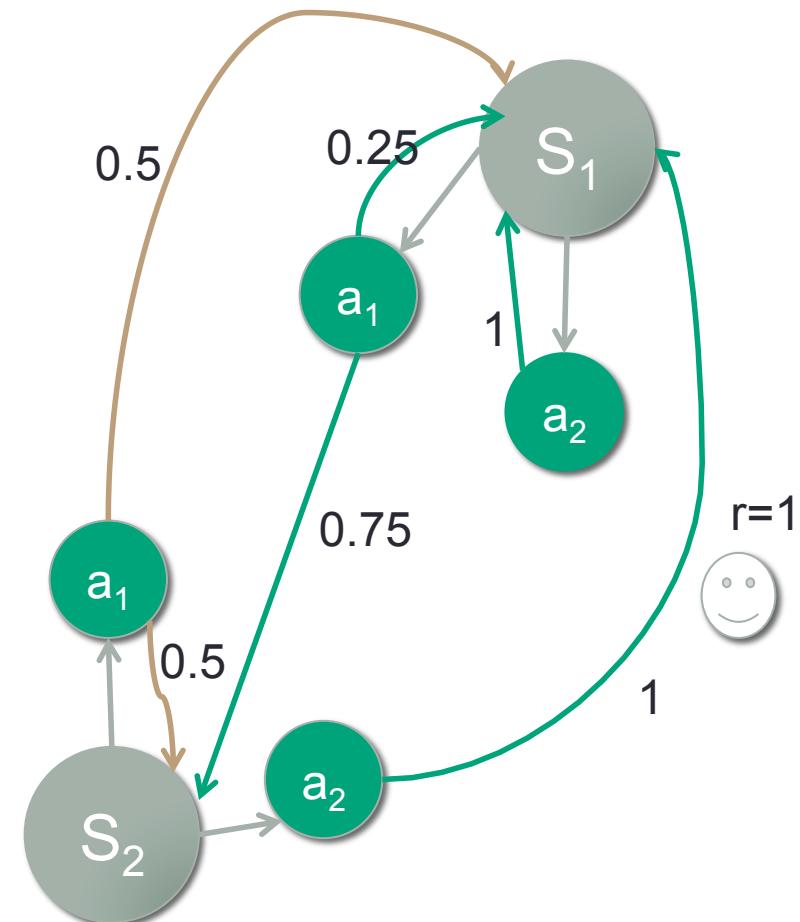
$$r_{t+34} = -1$$

Markov Decision Process (MDP)

- **S,A,P,R, γ**
- **S** – Set of states
- **A** – Set of actions
- **P** – Transition between states given an action

$$P_{s,s'}^a = \text{Prob}[s_{t+1} = s' \mid s_t = s, a_t = a]$$

- **R** – Rewards associated with actions and states
- **γ** – Discount factor



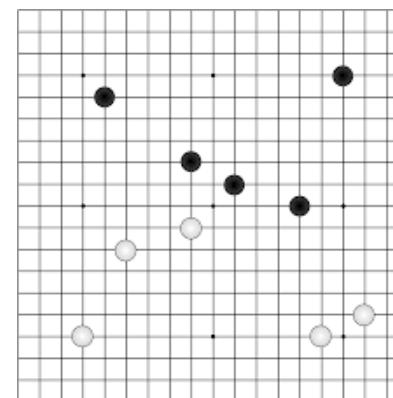
Markov State

- Markov property: All information from past is captured in the current state



A state

For Go, a board position
For simple video games, stack multiple frames



Policy

- A strategy to perform an action according to a particular state

$$\pi(s) = a$$

- A policy defines the behavior of the agent
- In MDP, we want to learn the best policy for the task
 - How?
 - Define some loss/reward function and minimize/maximize.

Action-Value Function (Q Function)

- Expected total rewards starting from state s taking action a then follows a given policy π

$$\begin{aligned} Q_\pi(s, a) &= E[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots | s_t = s, a_t = a] \\ &= E[r_t + \gamma Q_\pi(s', a') | s_t = s, a_t = a] \\ &= R_s^a + \gamma \sum_{s'} P_{s,s'}^a \max_{a' \in A} Q_\pi(s', a') \end{aligned}$$

- Note, Q is a function of s , a , and
- Q^* determines the optimal policy π^*

$$Q^*(s, a) = \max_{\pi} Q_{\pi}(s, a)$$

Q function and policy example

- For a policy, Q function gives the expected rewards

	Left	Right	Policy
	134	123	Left
	579	0	Left

- We want to find the best policy by maximizing the Q function.
- But we don't know Q. Need to estimate this table.
 - How? – Bellman Equation (iteratively update Q)

Q-learning

- Q-learning is used to learn Q^* by **sampling** the states s' and rewards r through events
 - *Starts with an original belief of Q values*
 - *Takes random actions to explore the possibilities*
 - *Update your belief of Q depending on the rewards from the environment*
- No need to explicitly learn the transition probability or the expected rewards
 - “**Learns by trial and error**”

Temporal difference Q-learning

One method to learn Q

Update Q by the reward received from taking action a in state s.

State s' is the state after taking action a.

$$\text{new } Q(s, a) = (1 - \alpha)Q(s, a) + \alpha[r + \gamma \max_{a' \in A} Q(s', a')]$$

Recall

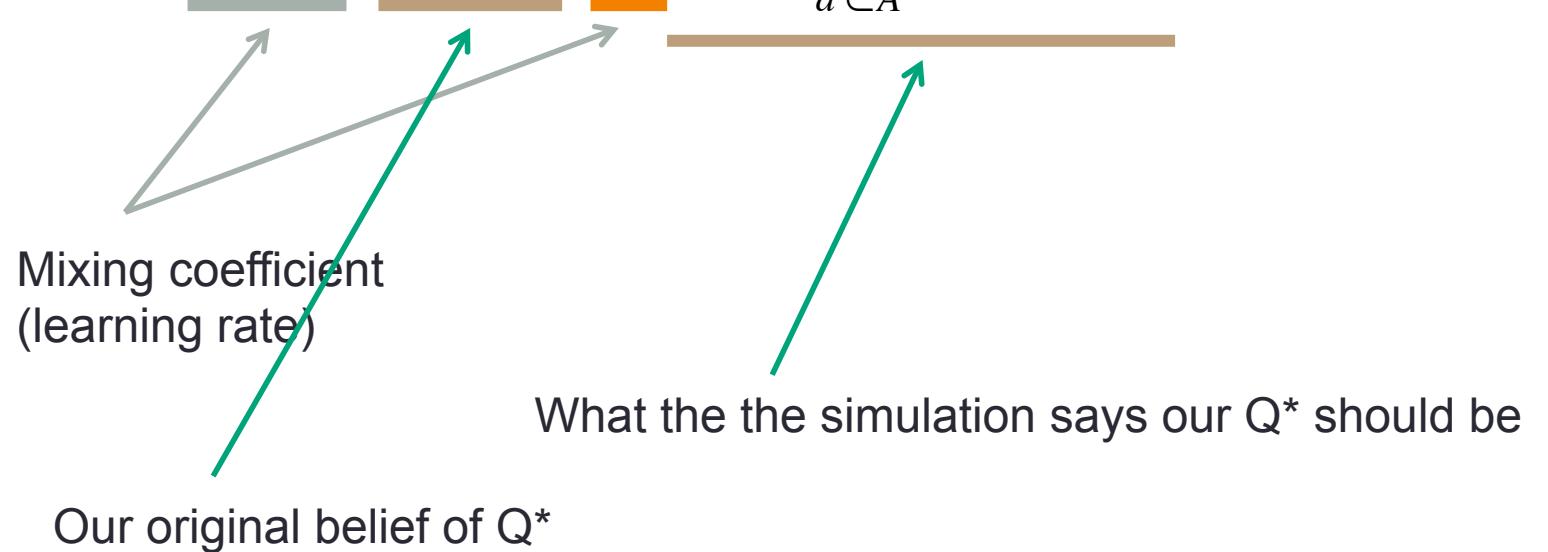
$$Q(s, a) = R_s^a + \gamma \sum_{s'} P_{s,s'}^a \max_{a' \in A} Q(s', a')$$

R (the expected rewards) becomes r (the sampled reward)
 The transition probability, P, is dropped because we sampled one outcome. This is a proxy for the sampled Q^*

Temporal difference Q-learning

Update Q by the reward received from taking action a in state s.

$$\text{new}Q(s, a) = (1 - \alpha)Q(s, a) + \alpha[r + \gamma \max_{a' \in A} Q(s', a')]$$



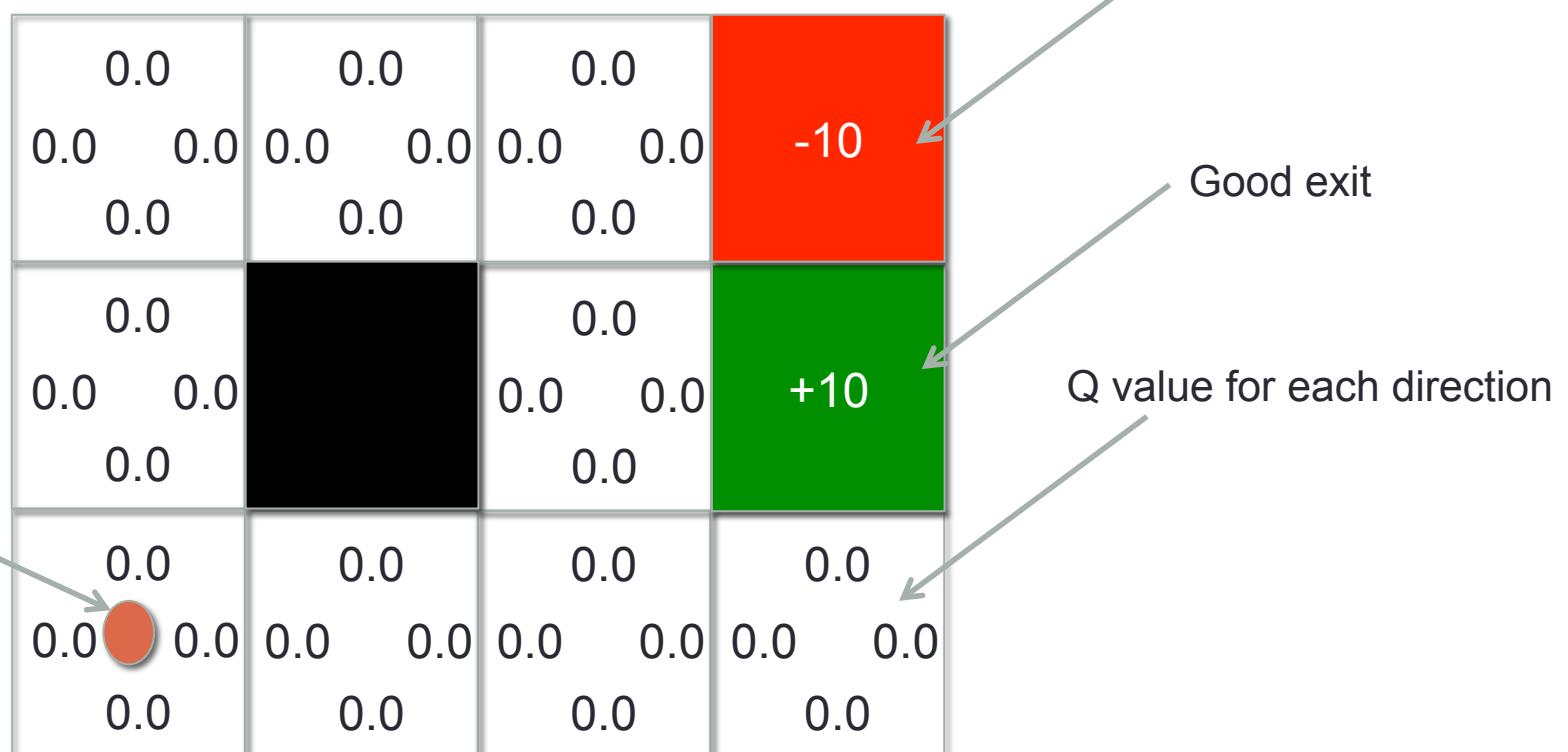
Toy example

The states are the different positions in the room
 4 actions: up, down, left, right

$$\gamma = 1.0$$

$$\alpha = 0.5$$

Agent



A reward of -1 for every move (penalty for time wasted)

Toy example

$$\text{new } Q(s, a) = (1 - \alpha)Q(s, a) + \alpha[r + \gamma \max_{a' \in A} Q(s', a')]$$

$$\begin{aligned}\gamma &= 1.0 \\ \alpha &= 0.5\end{aligned}$$

0.0	0.0	0.0	-10
0.0	0.0	0.0	-10
0.0	0.0	0.0	-10
0.0	0.0	0.0	+10
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

A 7x7 grid representing a state-action value function (Q-table). The last row and column are shaded grey, indicating they are not part of the current state-action pair being updated. The cell at (0,0) is highlighted with a red circle and has an arrow pointing to it from the left.

$$(1-0.5)0+0.5[-1 + 1 * 0] = -0.5$$

Toy example

$$\text{new } Q(s, a) = (1 - \alpha)Q(s, a) + \alpha[r + \gamma \max_{a' \in A} Q(s', a')]$$

$$\begin{aligned}\gamma &= 1.0 \\ \alpha &= 0.5\end{aligned}$$

0.0	0.0	0.0	-10
0.0	0.0	0.0	-10
0.0	0.0	0.0	-10
0.0	0.0	0.0	+10
0.0	0.0	0.0	+10
-0.5	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

Toy example

$$\text{new } Q(s, a) = (1 - \alpha)Q(s, a) + \alpha[r + \gamma \max_{a' \in A} Q(s', a')]$$

$$\begin{aligned}\gamma &= 1.0 \\ \alpha &= 0.5\end{aligned}$$

0.0	0.0	0.0	0.0	-10
0.0	0.0	0.0	0.0	
0.0	0.0	0.0	0.0	+10
0.0	0.0	0.0	0.0	
-0.5	0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0	0.0

$$(1-0.5)0+0.5[-1 + 1 * 0] = -0.5$$

Toy example

$$\text{new } Q(s, a) = (1 - \alpha)Q(s, a) + \alpha[r + \gamma \max_{a' \in A} Q(s', a')]$$

$$\begin{aligned}\gamma &= 1.0 \\ \alpha &= 0.5\end{aligned}$$

0.0	0.0	0.0	0.0	-10
0.0	0.0	0.0	0.0	-10
-0.5		0.0	0.0	+10
0.0	0.0	0.0	0.0	+10
-0.5	0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0	0.0

Toy example

$$\text{new } Q(s, a) = (1 - \alpha)Q(s, a) + \alpha[r + \gamma \max_{a' \in A} Q(s', a')]$$

$$\begin{aligned}\gamma &= 1.0 \\ \alpha &= 0.5\end{aligned}$$

0.0	0.0	0.0	0.0	0.0	0.0	-10
0.0	0.0	0.0	0.0	0.0	0.0	
-0.5				0.0	0.0	+10
0.0	0.0			0.0	0.0	
0.0				0.0	0.0	
-0.5	0.0	0.0	0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0	0.0	0.0	0.0
0.0				0.0	0.0	

$$(1-0.5)0+0.5[-1 + 1 * 0] = -0.5$$

Toy example

$$\text{new } Q(s, a) = (1 - \alpha)Q(s, a) + \alpha[r + \gamma \max_{a' \in A} Q(s', a')]$$

$$\begin{aligned}\gamma &= 1.0 \\ \alpha &= 0.5\end{aligned}$$

0.0	0.0	0.0	0.0	-10
0.0	-0.5	0.0	0.0	0.0
0.0	0.0	0.0	0.0	+10
-0.5			0.0	
0.0	0.0	0.0	0.0	
0.0	0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0	0.0

$$(1-0.5)0+0.5[-1 + 1 * 0] = -0.5$$

Toy example

$$\text{new } Q(s, a) = (1 - \alpha)Q(s, a) + \alpha[r + \gamma \max_{a' \in A} Q(s', a')]$$

$$\begin{aligned}\gamma &= 1.0 \\ \alpha &= 0.5\end{aligned}$$

0.0	0.0	0.0	
0.0	-0.5	0.0	-0.5
0.0	0.0	0.0	0.0
-0.5			0.0
0.0	0.0	0.0	0.0
0.0			+10
-0.5			0.0
0.0	0.0	0.0	0.0
0.0			0.0

$$(1-0.5)0+0.5[-1 + 1 * -10] = -5.5$$

Toy example

$$\text{new}Q(s,a) = (1-\alpha)Q(s,a) + \alpha[r + \gamma \max_{a' \in A} Q(s',a')]$$

$$\gamma = 1.0$$

0.0	0.0	0.0	
0.0	-0.5	0.0	-0.5
0.0	0.0	0.0	0.0
-0.5			0.0
0.0	0.0	0.0	0.0
0.0			0.0
-0.5	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

Toy example

$$\text{new } Q(s, a) = (1 - \alpha)Q(s, a) + \alpha[r + \gamma \max_{a' \in A} Q(s', a')]$$

$$\begin{aligned}\gamma &= 1.0 \\ \alpha &= 0.5\end{aligned}$$

0.0	0.0	0.0	-5.5	-10
0.0	-0.5	0.0	-0.5	0.0
0.0	0.0	0.0	0.0	0.0
-0.5			0.0	+10
0.0	0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0	0.0
-0.5	0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0	0.0

Toy example

$$\text{new } Q(s, a) = (1 - \alpha)Q(s, a) + \alpha[r + \gamma \max_{a' \in A} Q(s', a')]$$

$$\begin{aligned}\gamma &= 1.0 \\ \alpha &= 0.5\end{aligned}$$

0.0	0.0	0.0	-5.5	-10
0.0	-0.5	0.0	-0.5	0.0
0.0	0.0	0.0	0.0	0.0
-0.5			0.0	+10
0.0	0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0	0.0
-0.5	0.0	0.0	0.0	0.0
0.0	-0.5	0.0	0.0	0.0
0.0	0.0	0.0	0.0	0.0

Toy example

$$\text{new } Q(s, a) = (1 - \alpha)Q(s, a) + \alpha[r + \gamma \max_{a' \in A} Q(s', a')]$$

$$\begin{aligned}\gamma &= 1.0 \\ \alpha &= 0.5\end{aligned}$$

0.0	0.0	0.0	-5.5	-10
0.0	-0.5	0.0	-0.5	0.0
0.0	0.0	0.0	0.0	0.0
-0.5			0.0	+10
0.0	0.0	0.0	0.0	0.0
0.0			0.0	0.0
-0.5	0.0	0.0	0.0	0.0
0.0	-0.5	0.0	-0.5	0.0
0.0	0.0	0.0	0.0	0.0

Toy example

$$\text{new } Q(s, a) = (1 - \alpha)Q(s, a) + \alpha[r + \gamma \max_{a' \in A} Q(s', a')]$$

$$\begin{aligned}\gamma &= 1.0 \\ \alpha &= 0.5\end{aligned}$$

0.0	0.0	0.0	
0.0	-0.5	0.0	-0.5
0.0	0.0	0.0	0.0
-0.5			0.0
0.0	0.0	0.0	0.0
0.0			0.0
-0.5	0.0	0.0	0.0
0.0	-0.5	0.0	-0.5
0.0	0.0	0.0	0.0

The table shows a 9x4 grid of values. The last row contains a red box with a value of -10, a green box with a value of +10, and two white boxes with a value of 0.0. An arrow points from the +10 box to the bottom right corner of the green box.

$$(1-0.5)0+0.5[-1 + 1 * 10] = 4.5$$

Toy example

$$\text{new } Q(s, a) = (1 - \alpha)Q(s, a) + \alpha[r + \gamma \max_{a' \in A} Q(s', a')]$$

$$\begin{aligned}\gamma &= 1.0 \\ \alpha &= 0.5\end{aligned}$$

0.0	0.0	0.0	-5.5	-10
0.0	-0.5	0.0	-0.5	0.0
0.0	0.0	0.0	0.0	0.0
-0.5			0.0	
0.0	0.0	0.0	0.0	0.0
0.0			0.0	
-0.5	0.0	0.0	0.0	4.5
0.0	-0.5	0.0	-0.5	0.0
0.0	0.0	0.0	0.0	0.0

Toy example

$$\text{new } Q(s, a) = (1 - \alpha)Q(s, a) + \alpha[r + \gamma \max_{a' \in A} Q(s', a')]$$

$$\begin{aligned}\gamma &= 1.0 \\ \alpha &= 0.5\end{aligned}$$

0.0	0.0	0.0				
0.0	-0.5	0.0	-0.5	0.0	-5.5	-10
0.0	0.0	0.0		0.0		
-0.5				0.0		
0.0	0.0			0.0	0.0	+10
0.0				0.0		
-0.5	0.0	0.0	0.0	0.0	4.5	
0.0	-0.5	0.0	-0.5	0.0	-0.5	0.0
0.0	0.0	0.0		0.0		

$$(1-0.5)(-0.5)+0.5[-1 + 1^* 4.5] = 1.5$$

Toy example

$$\text{new}Q(s,a) = (1 - \alpha)Q(s,a) + \alpha[r + \gamma \max_{a' \in A} Q(s',a')]$$

$$\begin{aligned}\gamma &= 1.0 \\ \alpha &= 0.5\end{aligned}$$

0.0	0.0	0.0				
0.0	-0.5	0.0	-0.5	0.0	-5.5	-10
0.0	0.0	0.0		0.0		
-0.5				0.0		
0.0	0.0			0.0	0.0	+10
0.0				0.0		
-0.5	0.0	0.0	0.0		4.5	
0.0	-0.5	0.0	-0.5	0.0	1.5	0.0
0.0	0.0	0.0		0.0		0.0

Toy example

$$\text{new } Q(s, a) = (1 - \alpha)Q(s, a) + \alpha[r + \gamma \max_{a' \in A} Q(s', a')]$$

$$\begin{aligned}\gamma &= 1.0 \\ \alpha &= 0.5\end{aligned}$$

0.0	0.0	0.0					
0.0	-0.5	0.0	-0.5	0.0	-5.5	-10	
0.0		0.0		0.0			
-0.5				0.0			
0.0	0.0			0.0	0.0	+10	
0.0				0.0			
-0.5		0.0		0.0		4.5	
0.0	-0.5	0.0	-0.5	0.0	1.5	0.0	0.0
0.0		0.0		0.0		0.0	

$$(1-0.5)(0.0)+0.5[-1 + 1 * 1.5] = 0.25$$

Toy example

$$\text{new } Q(s, a) = (1 - \alpha)Q(s, a) + \alpha[r + \gamma \max_{a' \in A} Q(s', a')]$$

$$\begin{aligned}\gamma &= 1.0 \\ \alpha &= 0.5\end{aligned}$$

0.0	0.0	0.0	
0.0	-0.5	0.0	-0.5
0.0	0.0	0.0	0.0
-0.5			0.0
0.0	0.0	0.0	+10
0.0			0.0
-0.5	0.0	0.0	4.5
0.0	-0.5	0.0	0.25
0.0	0.0	0.0	0.0

$$(1-0.5)(0.0)+0.5[-1 + 1 * 1.5] = 0.25$$

Toy example

$$\text{new } Q(s, a) = (1 - \alpha)Q(s, a) + \alpha[r + \gamma \max_{a' \in A} Q(s', a')]$$

$$\begin{aligned}\gamma &= 1.0 \\ \alpha &= 0.5\end{aligned}$$

0.0	0.0	0.0	
0.0	-0.5	0.0	-0.5
0.0	0.0	0.0	0.0
-0.5			0.0
0.0	0.0	0.0	+10
0.0			0.0
-0.5	0.0	0.0	4.5
0.0	-0.5	0.0	0.25
0.0	0.0	0.0	0.0

A 4x4 grid representing a state-action value function (Q-table). The values are as follows:

- Row 1: 0.0, 0.0, 0.0, (empty)
- Row 2: 0.0, -0.5, 0.0, -0.5
- Row 3: 0.0, 0.0, 0.0, 0.0
- Row 4: -0.5, (black), 0.0, 0.0
- Row 5: 0.0, 0.0, 0.0, +10
- Row 6: 0.0, (black), (black), 0.0
- Row 7: -0.5, 0.0, 0.0, 4.5
- Row 8: 0.0, -0.5, 0.0, 0.25
- Row 9: 0.0, 0.0, 0.0, 0.0

An arrow points from the value 0.0 in the fourth row to the value 0.0 in the fifth row, indicating the update step.

$$(1-0.5)(0.0)+0.5[-1 + 1 * 0.0] = -0.5$$

Toy example

$$\text{new } Q(s, a) = (1 - \alpha)Q(s, a) + \alpha[r + \gamma \max_{a' \in A} Q(s', a')]$$

$$\begin{aligned}\gamma &= 1.0 \\ \alpha &= 0.5\end{aligned}$$

0.0	0.0	0.0					
0.0	-0.5	0.0	-0.5	0.0	-5.5	-10	
0.0		0.0		0.0			
-0.5				0.0			
0.0	0.0		0.0	0.0	0.0	+10	
0.0							
-0.5	0.0		-0.5	4.5			
0.0	-0.5	0.0	-0.5	0.0	1.5	0.25	0.0
0.0		0.0		0.0			

$$(1-0.5)(0.0)+0.5[-1 + 1 * 10] = 4.5$$

Toy example

$$\text{new } Q(s, a) = (1 - \alpha)Q(s, a) + \alpha[r + \gamma \max_{a' \in A} Q(s', a')]$$

$$\begin{aligned}\gamma &= 1.0 \\ \alpha &= 0.5\end{aligned}$$

0.0	0.0	0.0	
0.0	-0.5	0.0	-0.5
0.0	0.0	0.0	0.0
-0.5			0.0
0.0	0.0	0.0	4.5
0.0			+ 0.5
-0.5	0.0	-0.5	4.5
0.0	-0.5	0.0	0.25
0.0	0.0	0.0	0.0

$$(1-0.5)(0.0)+0.5[-1 + 1 * 10] = 4.5$$

Toy example

$\gamma = 1.0$
 $\alpha = 0.5$

Many moves later

0.0	0.0	0.0	-10
0.0	-0.5	0.0	-8.25
0.0	0.0	2.625	
-0.5		0.375	
0.0	0.0	0.0	+10
0.0		0.25	
-0.5	0.0	1.5	6.75
0.0	-0.5	0.0	0.0
0.0	0.0	0.0	0.0

Toy example

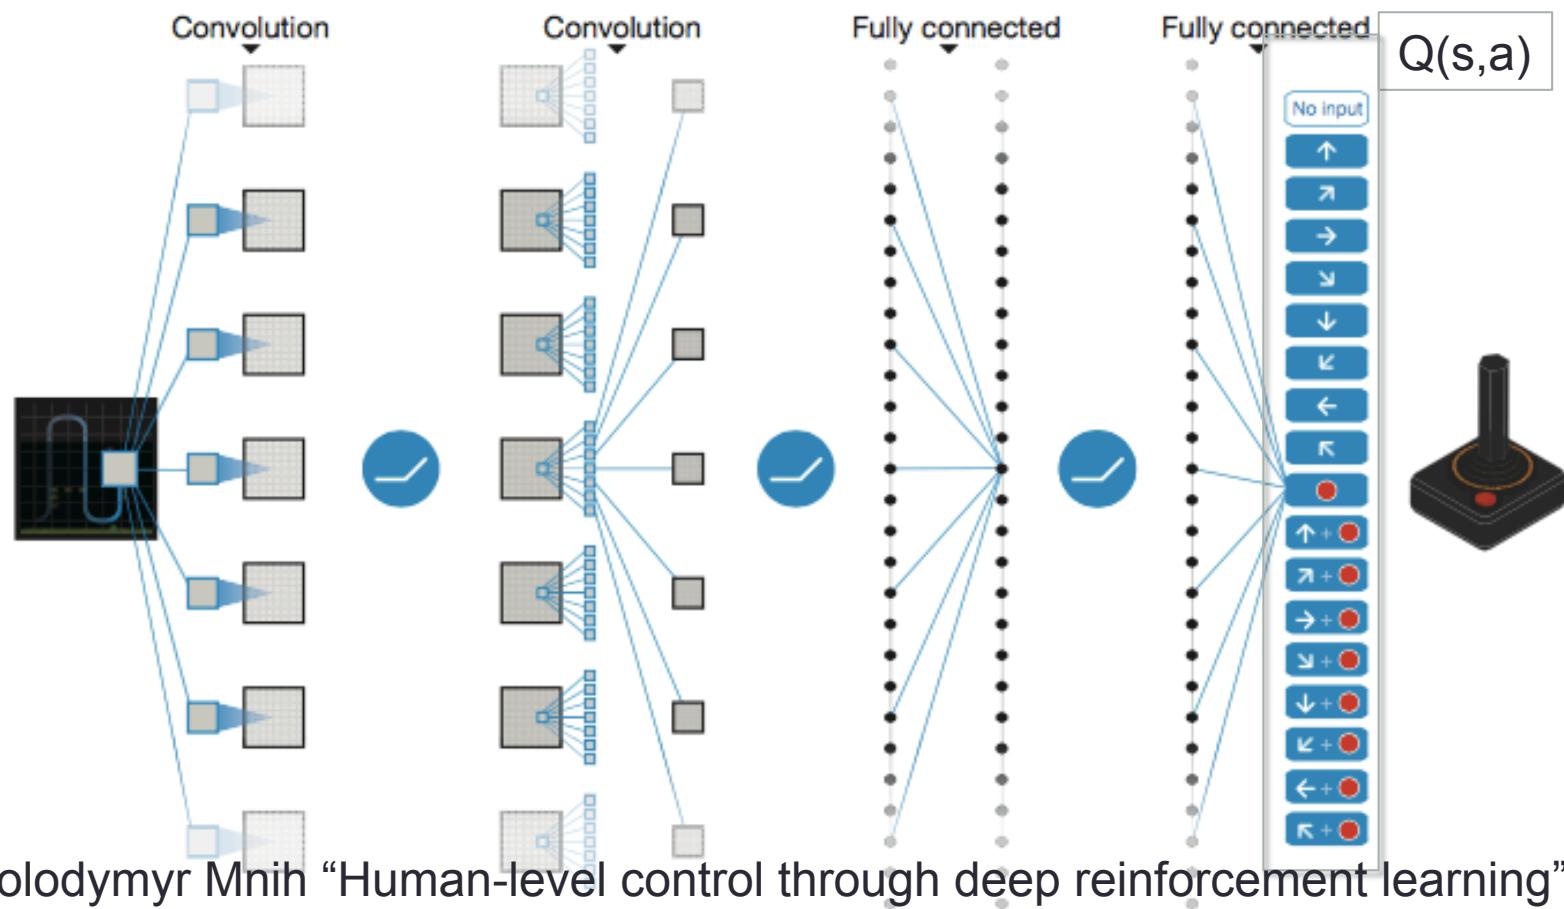
Many many more moves later

$$\begin{aligned}\gamma &= 1.0 \\ \alpha &= 0.5\end{aligned}$$

0.0	0.0	0.0	-10
0.0	6.0	5.0	7.0
4.0	0.0	8.0	
5.0		7.0	
0.0	0.0	0.0	+10
3.0		7.0	
4.0	0.0	8.0	9.0
0.0	6.0	5.0	7.0
0.0	0.0	0.0	0.0

Deep Q learning (DQN)

- Use a DNN to learn the Q function. Outputs are the possible actions.



DQN Loss Function

- Similar to the update equation of Q-learning

$$\text{new } Q(s, a) = (1 - \alpha)Q(s, a) + \alpha[r + \gamma \max_{a' \in A} Q(s', a')]$$

- We want to move our current belief (the Q generated from the network) towards the sampled Q

$$\text{Loss} = \{Q(s, a) - [r + \gamma \max_{a' \in A} Q(s', a')]\}^2$$

A couple more tricks

- Experience replay
 - Adjacent frames in a game are highly correlated
 - Saves past experiences and re-use them in the mini-batch
- Target network
 - Updating the Q belief too quickly leads to instability
 - Keeps a target network that updates less frequently.

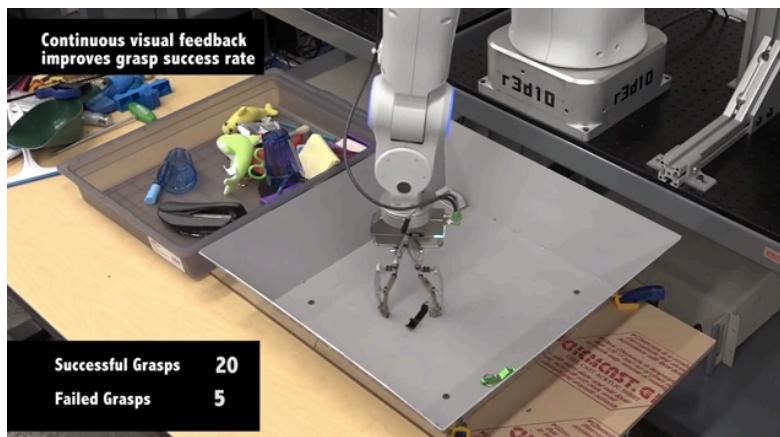
$$\text{Loss} = \{Q(s, a) - [r + \gamma \max_{a' \in A} \hat{Q}(s', a')] \}^2$$

Other applications using reinforcement learning

Robotics



<http://spectrum.ieee.org/automaton/robotics/drones/drone-uses-ai-and-11500-crashes-to-learn-how-to-fly>

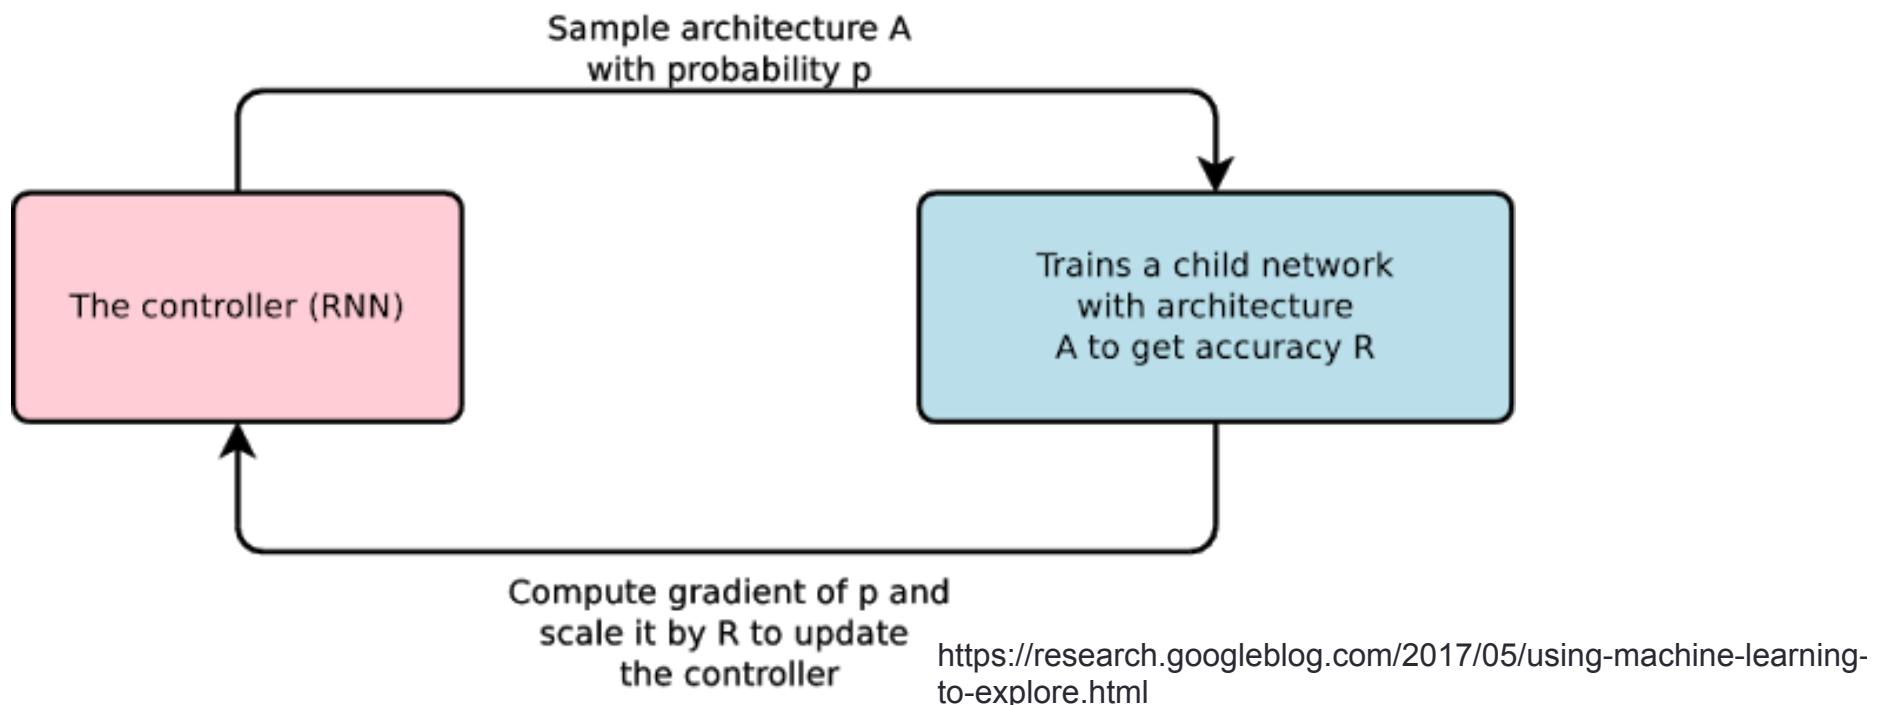


<https://www.youtube.com/watch?v=HbHqC8Himol>

<https://research.googleblog.com/2016/03/deep-learning-for-robots-learning-from.html>

Model selection for deep architectures

- Tuning a network takes time
- Let machine learning learns how to tune a network
- Matches or outperforms ML experts performance



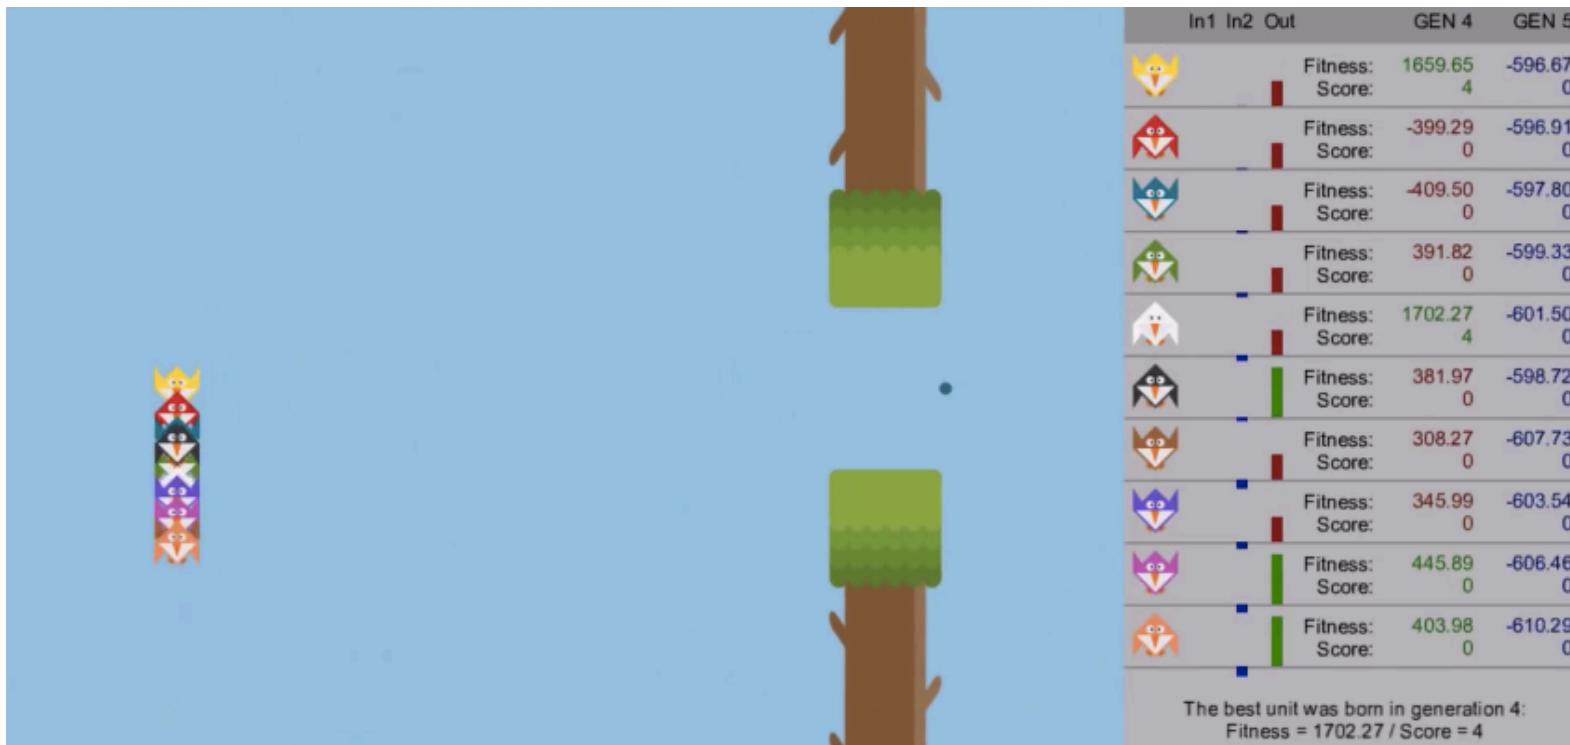
<https://research.googleblog.com/2017/05/using-machine-learning-to-explore.html>

Genetic Algorithm for Reinforcement learning

- Updating based on the gradients for RL is hard
 - Credit assignment problem
 - Target network as a workaround for too fast updates
- Estimate the direction to minimize the loss by random perturbations
 - Perturb the weights and let the new agents roam in the environment
- Highly parallelizable

Flappy bird genetic algorithm

- <https://www.youtube.com/watch?v=aeWmdojEJf0>



Yann LeCun's comment



Yann LeCun

March 14, 2016 ·

Follow

Statement from a Slashdot post about the AlphaGo victory: "We know now that we don't need any big new breakthroughs to get to true AI"

That is completely, utterly, ridiculously wrong.

As I've said in previous statements: most of human and animal learning is unsupervised learning. If intelligence was a cake, unsupervised learning would be the cake, supervised learning would be the icing on the cake, and reinforcement learning would be the cherry on the cake. We know how to make the icing and the cherry, but we don't know how to make the cake.

We need to solve the unsupervised learning problem before we can even think of getting to true AI. And that's just an obstacle we know about. What about all the ones we don't know about?

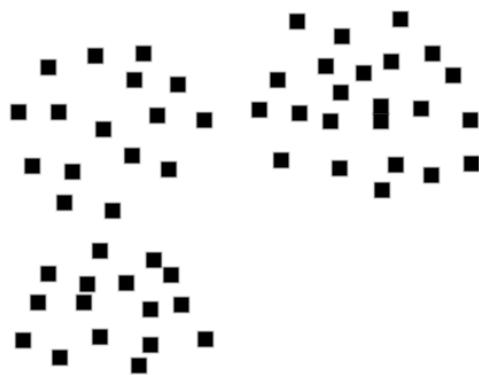
#deeplearning #AI #AlphaGo

Unsupervised Learning

- Clustering
- Autoencoders
- Generative Adversarial Networks

Clustering

- Discover underlying structure of data



- Learn the hidden class labels in the data.

Simple clustering

- K-means
 - Find the means, and assign
- GMM
 - Find the means, covariance, then soft assignment
- K-means vs GMM
 - Spherical covariance vs any covariance
 - Hard vs soft assignment
- Need to pick k

Latent Dirichlet Allocation (LDA)

- Another method that learns the hidden labels.
- Mostly used for topic modeling and document classification
- Let's cover some basic knowledge about graphical models

Topic modeling motivation

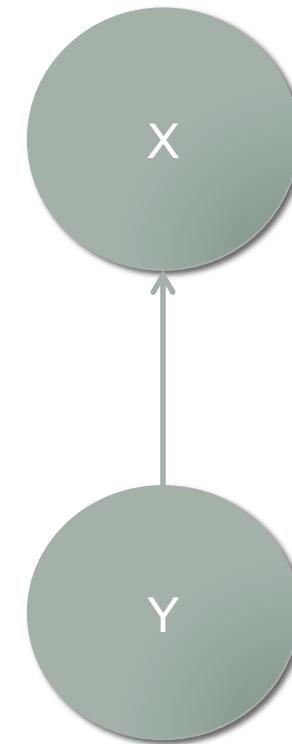
- I want to be able to give probabilities to words/sentences/documents.
- $P(\text{"Roses are red. Violets are blue."}) = ?$
- $P(X_1 = \text{Rosses}, X_2 = \text{are}, X_3 = \text{red}, \dots)$
 - This probability distribution is intractable

Introduction to graphical models

- Given rain in the morning predict sunny or cloudy in the afternoon

Graphical representation of the relationship

$P(X Y)$	$Y =$ Rain	$Y =$ No rain
$X =$ cloudy		
$X =$ sunny		

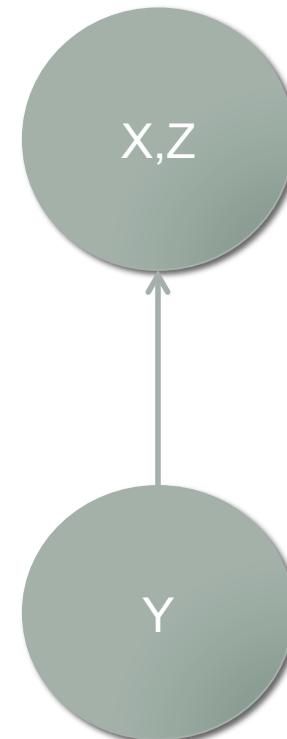


Curse of dimensionality

- Given rain in the morning predict sunny or cloudy and windy or cloudy in the afternoon

Graphical representation of the relationship

$P(X,Z Y)$	$Y =$ Rain	$Y =$ No rain
$X =$ cloudy $Z =$ windy		
$X =$ sunny $Z =$ windy		
$X =$ cloudy $Z =$ no wind		
$X =$ sunny $Z =$ no wind		

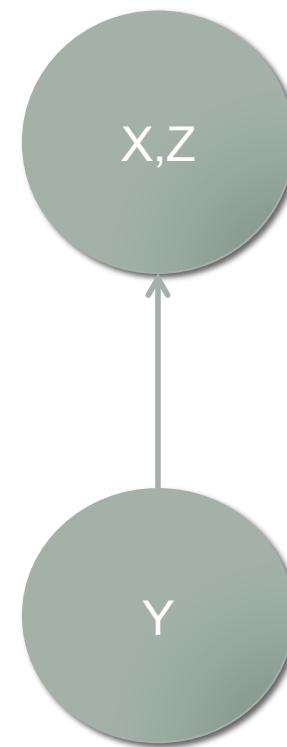


Curse of dimensionality

- If there are more variables, this table keeps getting bigger and bigger

Graphical representation of the relationship

$P(X,Z Y)$	$Y =$ Rain	$Y =$ No rain
X = cloudy Z = windy		
X = sunny Z = windy		
X = cloudy Z = no wind		
X = sunny Z = no wind		



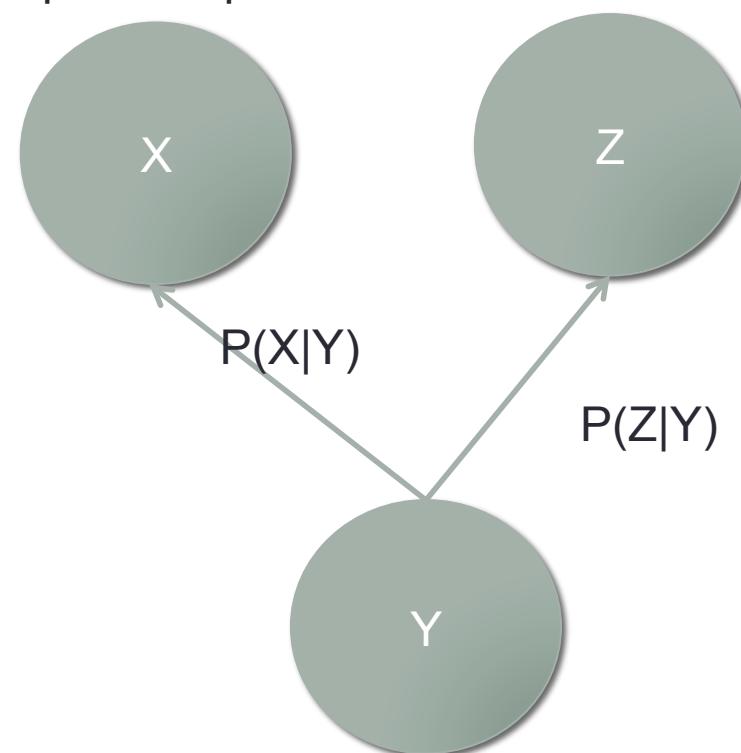
Naïve Bayes

- Assume conditional independence

$P(X Y)$	$Y =$ Rain	$Y =$ No rain
$X =$ cloudy		
$X =$ sunny		

$P(Z Y)$	$Y =$ Rain	$Y =$ No rain
$Z =$ windy		
$Z =$ no wind		

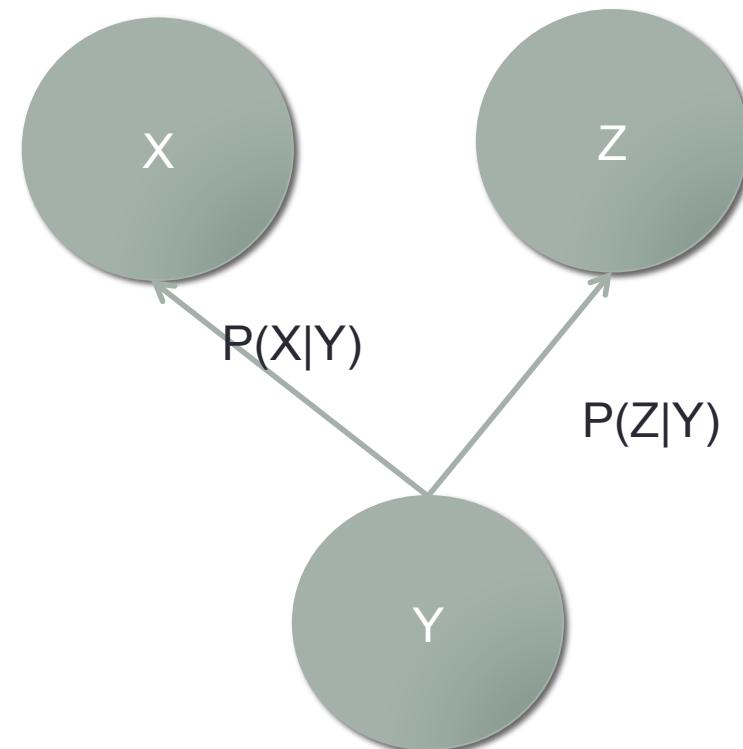
Graphical representation of the relationship



Probabilistic models and independence assumptions

- Most probabilistic models will assume some independence assumption in order to reduce the parameter space
- Graphical representations helps summarize the relationship
 - Arrows denote a conditional relationship

Graphical representation of the relationship



Probabilistic models and independence assumptions

- Also implies
- $P(X, Y, Z) = P(Y)P(X|Y)P(Z|Y)$

Graphical representation of the relationship

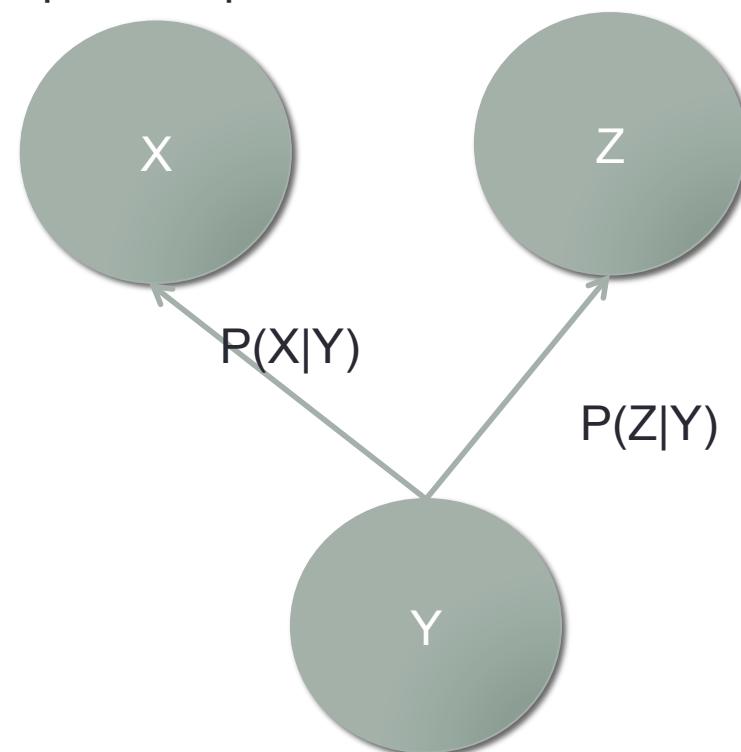


Plate notation

- Sometimes you have too many relationships.

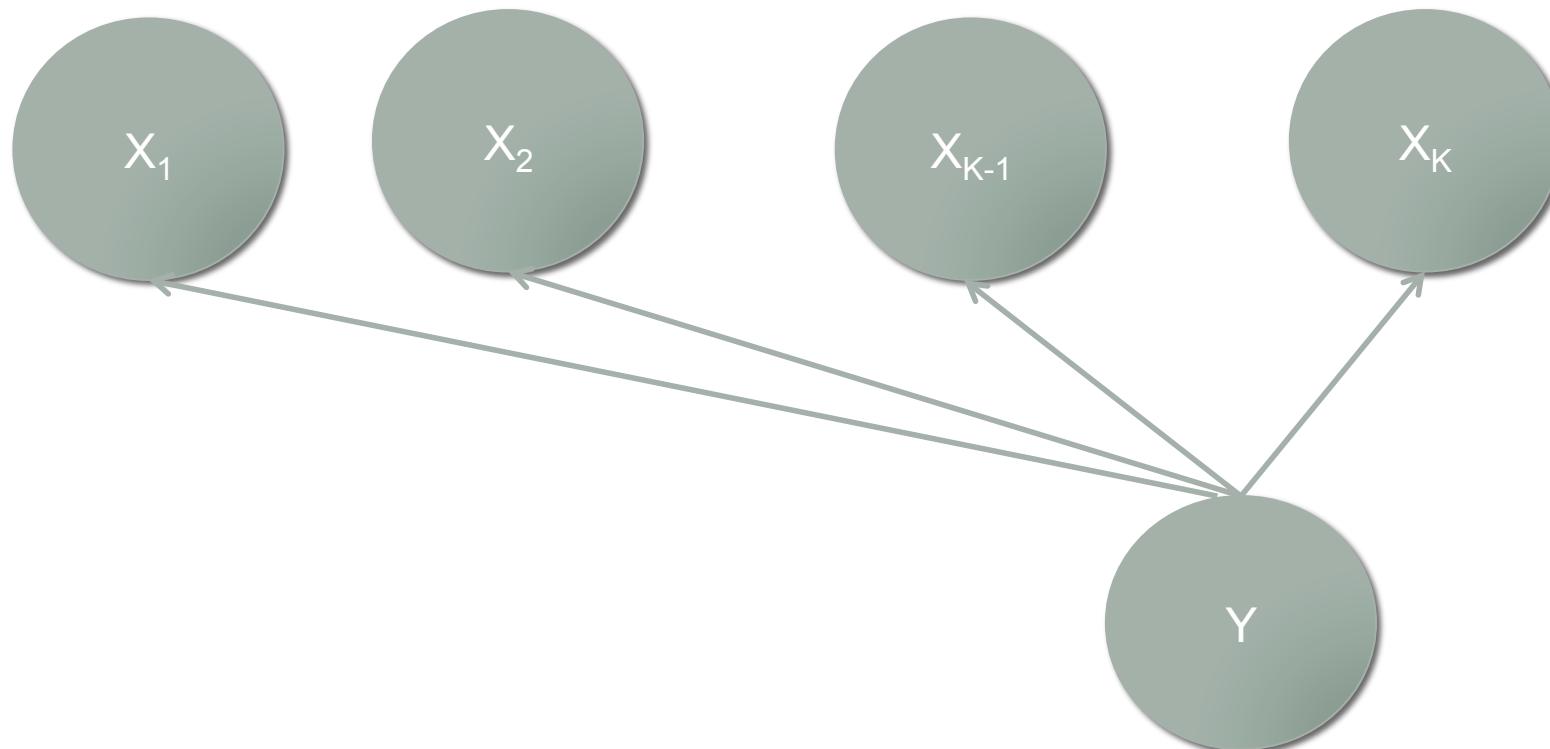
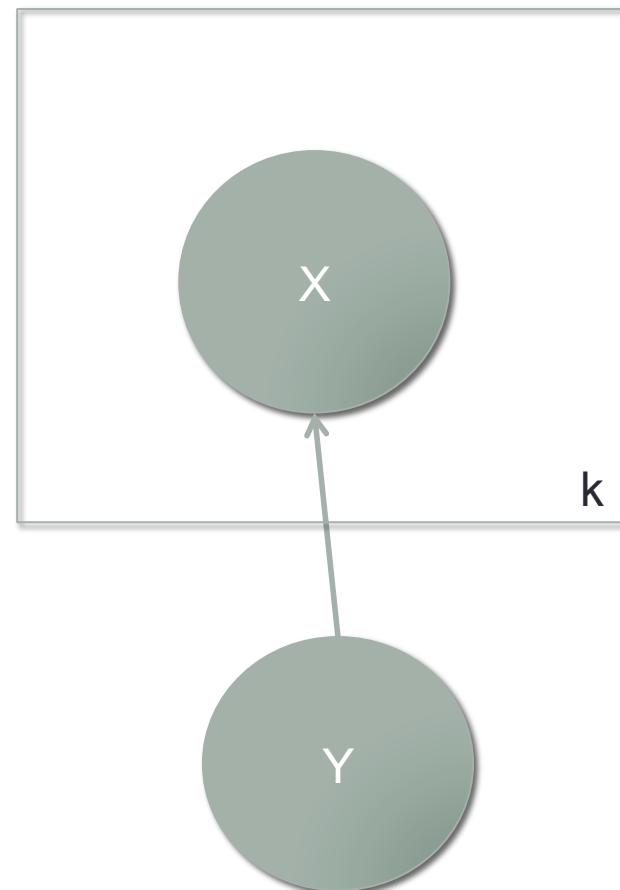


Plate notation

- Summarize by using a square box with number



Topic modeling motivation

- I want to be able to give probabilities to words/sentences/documents.
- $P(\text{"Roses are red. Violets are blue."}) = ?$
- $P(X_1 = \text{Rosses}, X_2 = \text{are}, X_3 = \text{red}, \dots)$
 - This probability distribution is intractable
- Assume independence
 - $P(X_1) P(X_2) P(X_3) \dots$
 - Lost all relational structure. Can we do better?
 - Assume words come from topics. How?

Topic modeling example

- I want to write about Pattern Recognition
- Topics I want to cover
 - GMMs
 - Neural networks
 - Linear regression
- Let's assign the proportion of topics
 - $P(\text{GMMs}) = 0.3$
 - $P(\text{Neural networks}) = 0.6$
 - $P(\text{Linear regression}) = 0.1$

Topic modeling example

- I want to write about Pattern Recognition
- Topics I want to cover
 - GMMs
 - Neural networks
 - Linear regression
- Let's assign the proportion of topics
 - $P(\text{GMMs}) = 0.3$
 - $P(\text{Neural networks}) = 0.6$
 - $P(\text{Linear regression}) = 0.1$

Topic modeling example

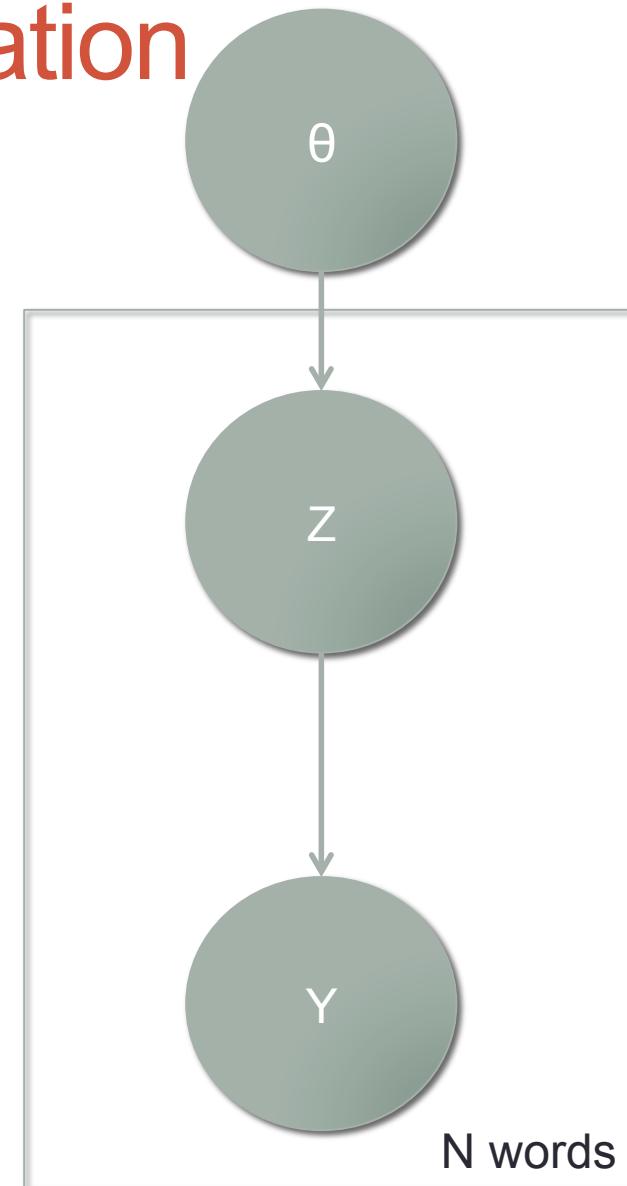
- Each topic can have a word proportion
 - GMM
 - $P(\text{Guassian}) = 0.7, P(\text{MLE}) = 0.1, \dots$
 - Neural networks
 - $P(\text{neuron}) = 0.3, P(\text{softmax}) = 0.5, \dots$
 - Linear regression
 -

Topic modeling

- $P(W_1=\text{Guassian}, W_2=\text{neuron}, \dots)$
- Let Z_n be the topic of W_n
 - $= P(W_1 = \text{Gaussian} | Z_1 = \text{GMM}) P(W_2 = \text{neuron} | Z_2 = \text{neural networks}) * \dots * P(Z_1 = \text{GMM}) P(Z_2 = \text{neural networks}) * \dots$
 - $= 0.7 * 0.3 * \dots$
- * $0.3 * 0.6 * \dots$

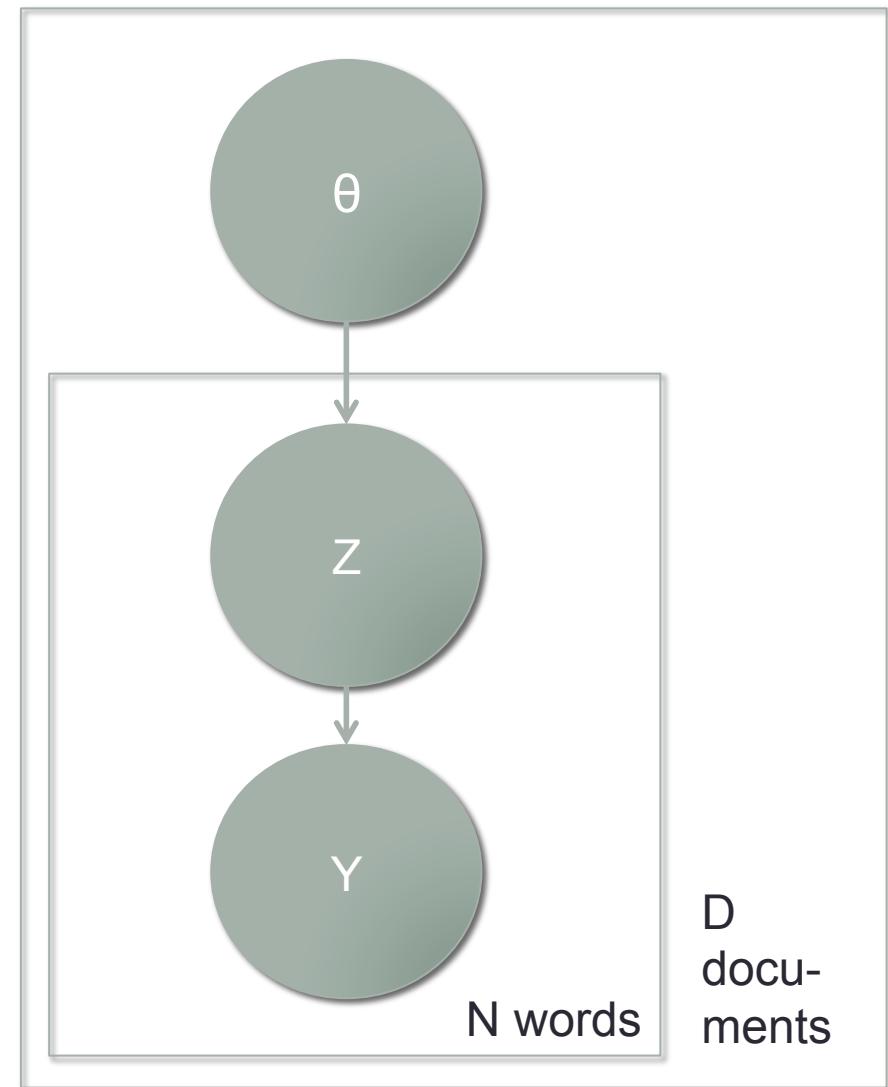
Graphical representation

θ is the distribution of topics for each word



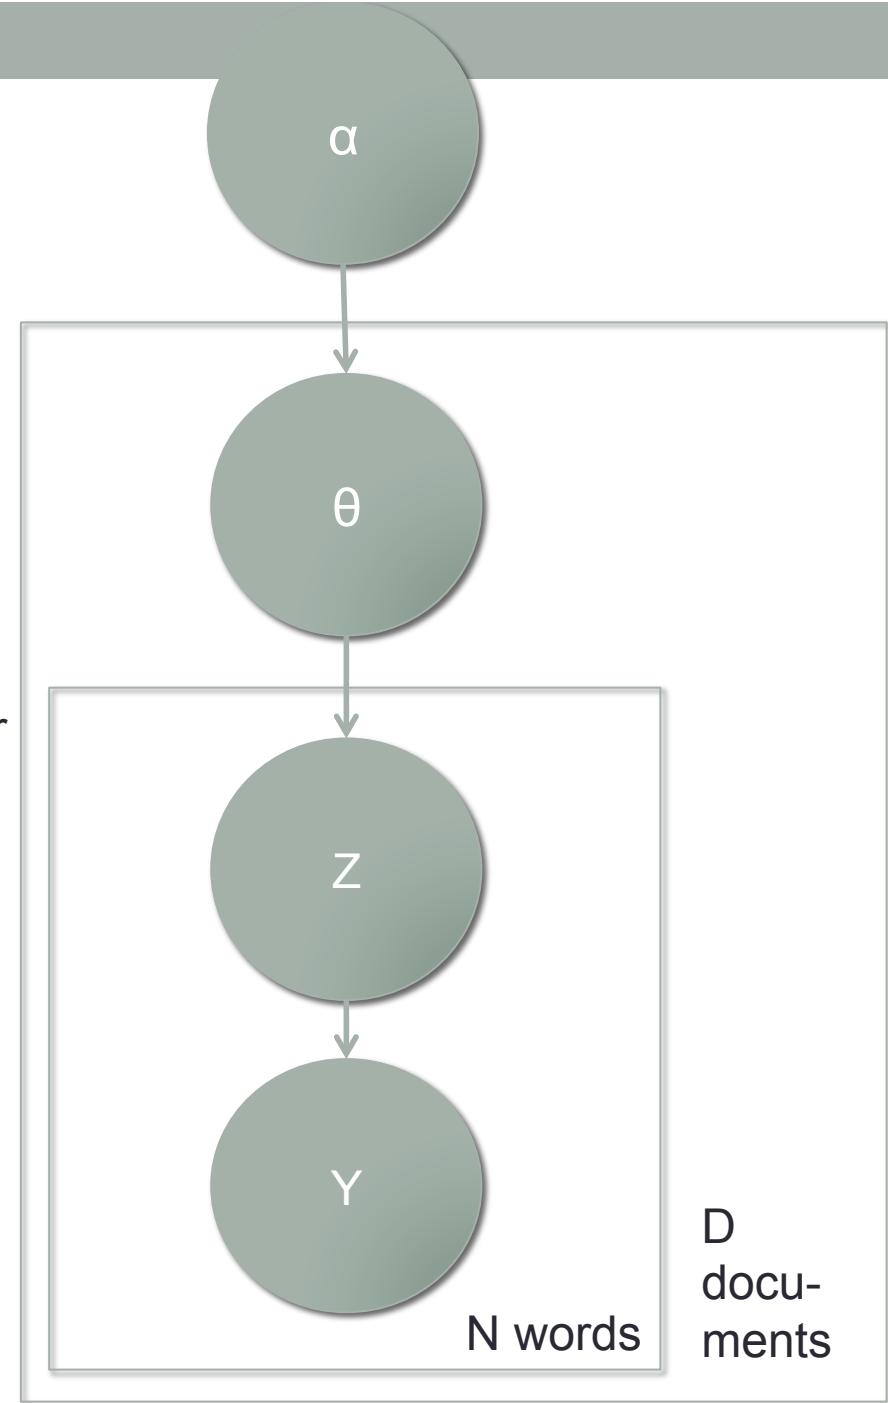
Multiple documents

θ is the distribution of topics for each word
in the document



Multiple documents

θ is the distribution of topics for each word
in the document
 α is the distribution of documents topics for
each document



α

- α is a **Dirichlet distribution**
- Or a distribution of distributions...
- ??????
- Example: rolling a die
- $P(x=1) = 0.3, P(x=2) = 0.1, \dots$
- This is a distribution. (**Multinomial distribution**)
- But what if I have a bag of dices, each with different distributions.
- α tells me how probability of what kind of die I will get

Dirichlet distribution

- pdf

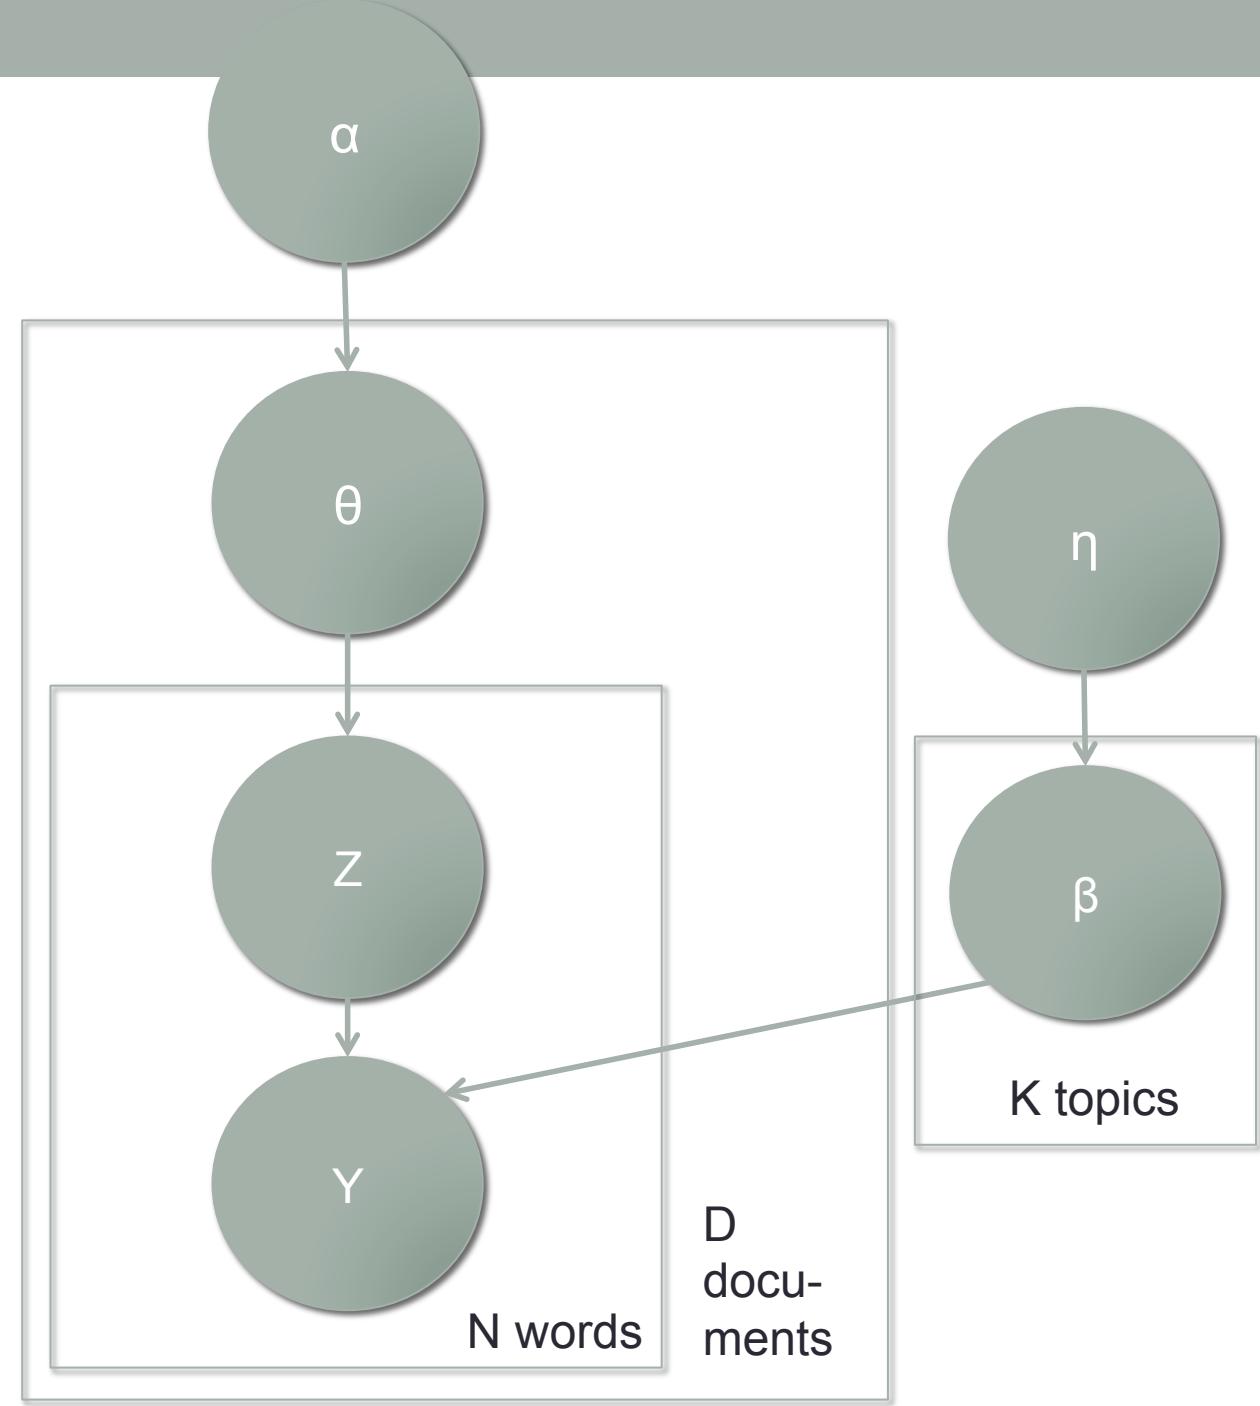
Parameters giving preference to each side of die/topic

$$f(x_1, \dots, x_K; \alpha_1, \dots, \alpha_K) = \frac{1}{B(\boldsymbol{\alpha})} \prod_{i=1}^K x_i^{\alpha_i - 1}$$

Probability of each side of die, probability of each topic

Learning topic distributions

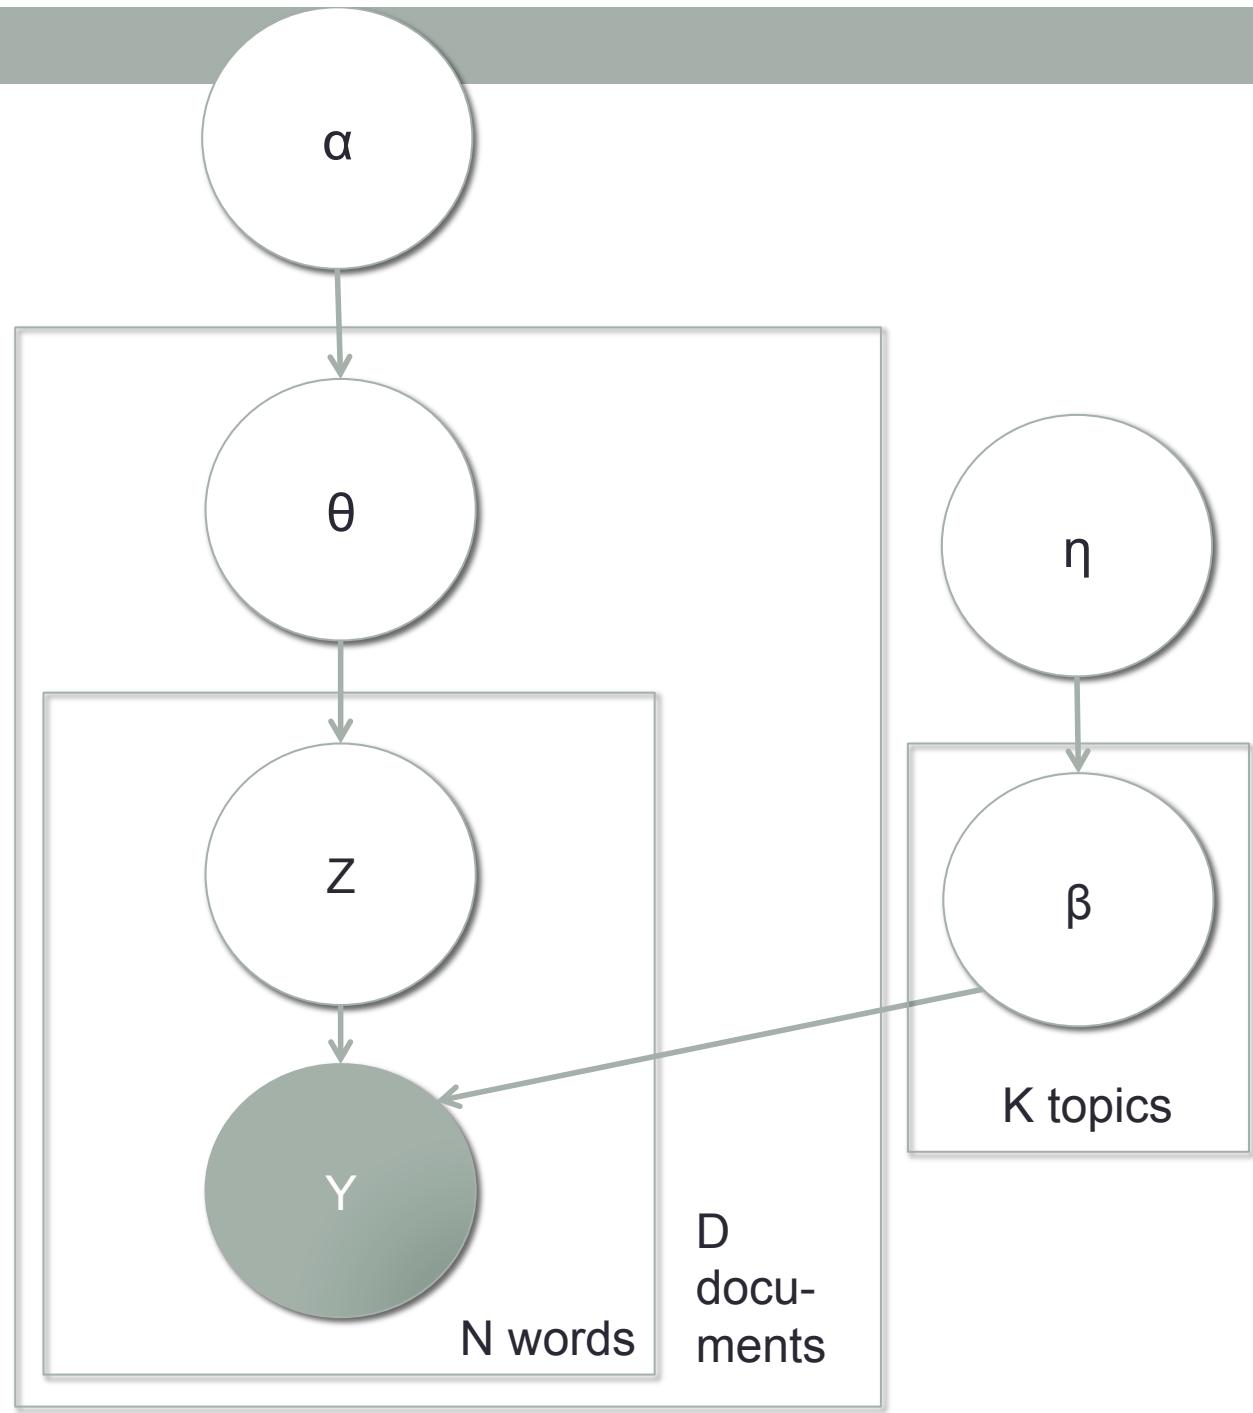
β is probability of word
given topic
 η is dirchlet for β



Latent variables

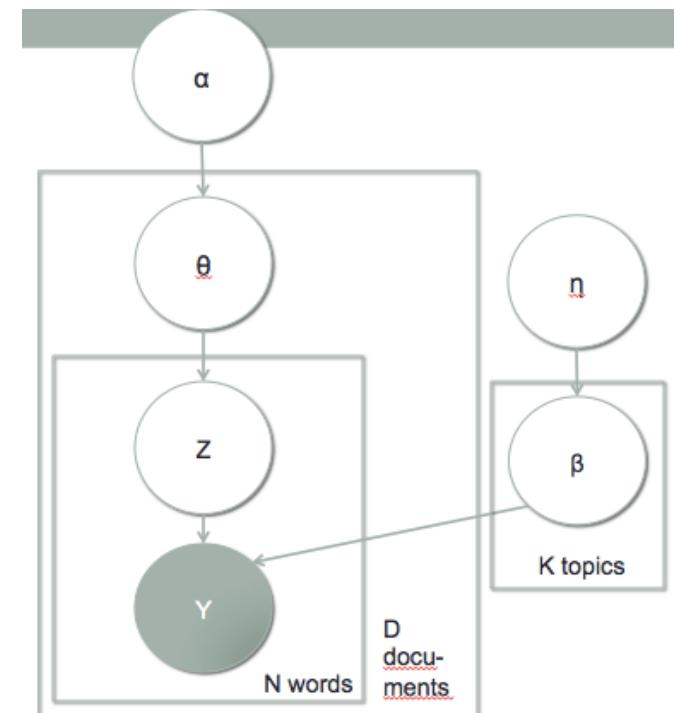
Graphical representation also use coloring to represent observed variables (fixed). Shaded are observed. While non-shaded are latent variables (z, β, θ) or hyperparameters (α, η)

Latent Dirchlet Allocation



Generative process

- For each topic $k \in \{1, \dots, K\}$:
 - Draw word distributions $\beta_k \sim \text{Dir}(\eta)$
- For each document $d \in \{1, \dots, D\}$:
 - Draw topic proportions $\theta_d \sim \text{Dir}(\alpha)$
 - For each word in a document $n \in \{1, \dots, N\}$:
 - Draw a topic index $z_{dn} \sim \text{Mult}(\theta_d)$
 - Generate word from chosen topic
 - $w_{dn} \sim \text{Mult}(\beta_{z_{dn}})$



Latent Dirichlet Allocation

- Guess the topic distributions of the document
 - Guess the topic assignment of each word
 - Guess the distribution of words for each topic
-
- How do we actually learn these parameters and latent variables?
 - Gibbs sampling
 - Variational methods
 - Totally beyond the scope of this class

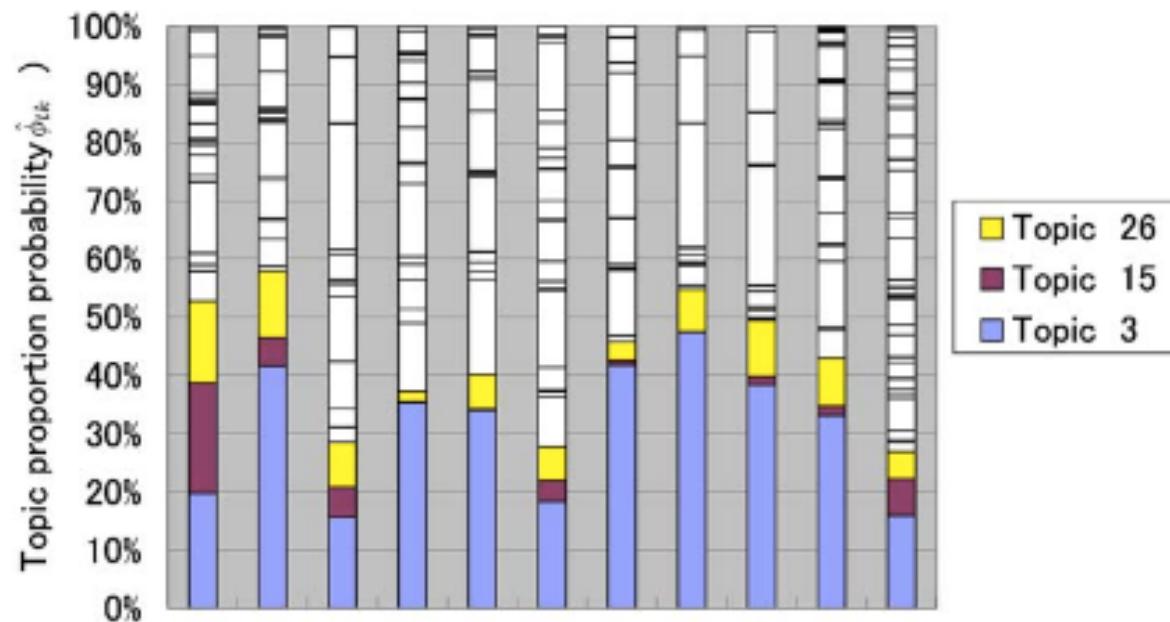
Example usage

- Lots of lectures video. Can we discover the topics? Can we classify the videos? No topic labels given

Top 10 high probability nouns in word probabilities of Topics 3 ($\theta_{k=3}$), 15 ($\theta_{k=15}$), and 26 ($\theta_{k=26}$).

Topic 3 (~ classical mechanics)	Topic 15 (~astronomy)	Topic 26 (~ (time) unit)
m	Light	Percent
Energy	Degrees	Time
Force	Angle	Dollars
Mass	Frequency	Times
Point	Energy	Minutes
Velocity	Direction	Day
Direction	Waves	Bit
v	Sun	Year
Times	Star	Hour
Speed	Speed	Half

Topic distribution of each document

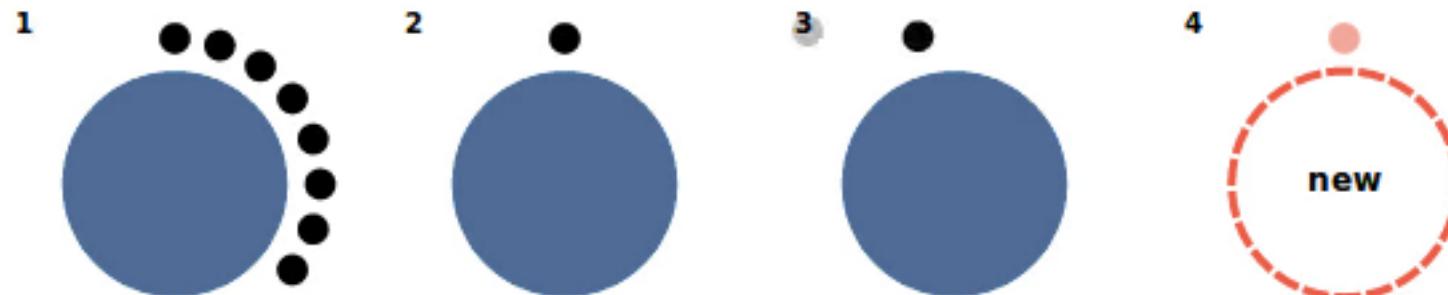


Discovering K

- In K-means, GMMs, LDA, we need K, the number of latent classes.
- Can we find it automatically?

Chinese Restaurant Process

- A random process is analogous to seating customers in a restaurant with infinite tables.
 - Each table has infinite seats
- First person sits in the first table
- n^{th} person sits at a table based on
 - Join existing table according to
 - number of person in that table, and table characteristics
 - Join a new empty table αc (some hyperparameter)



Chinese Restaurant Process

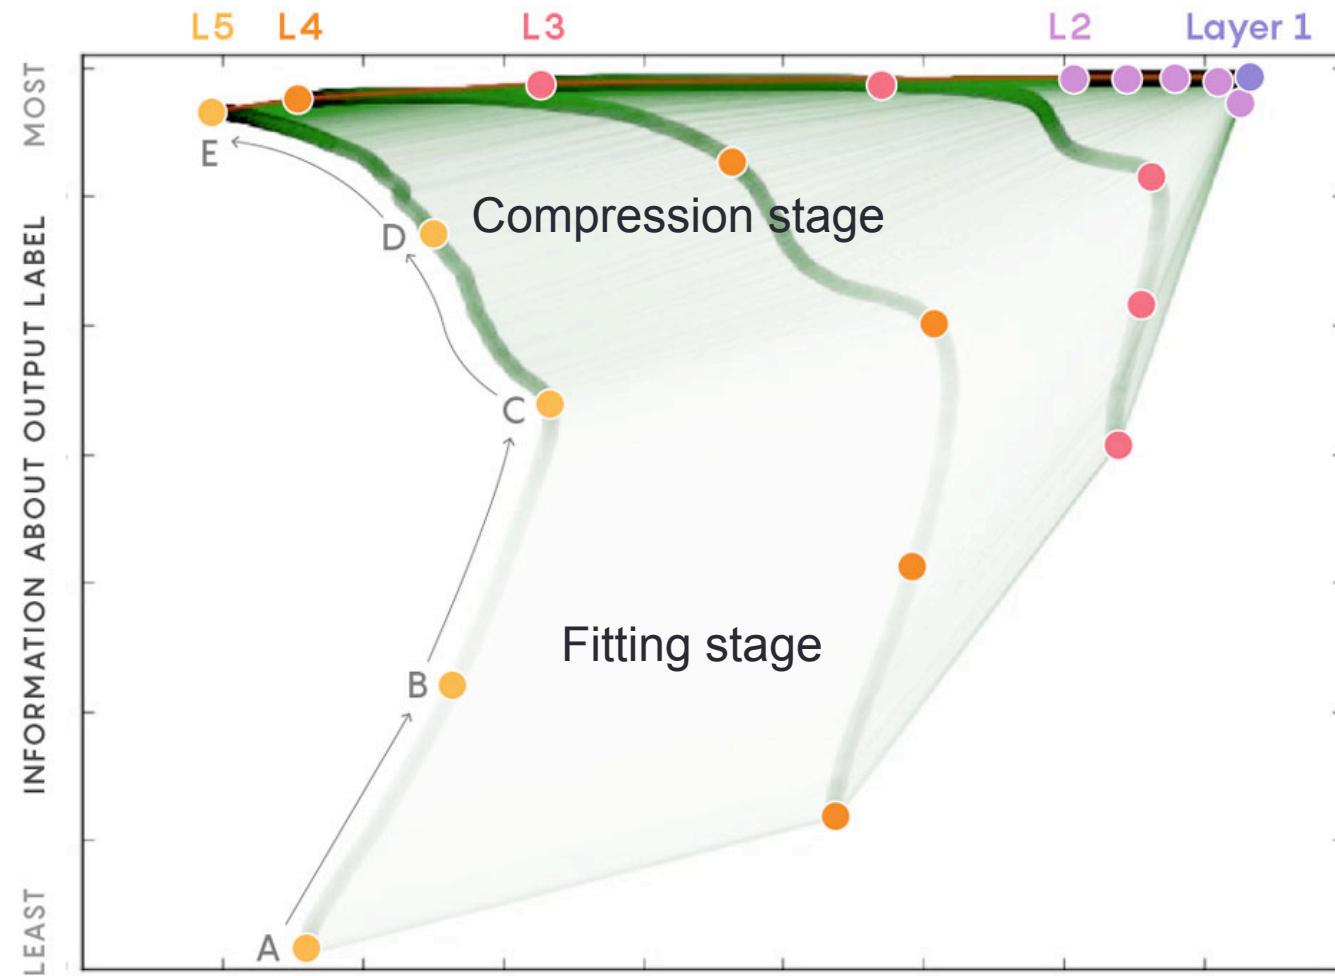
- After an assignment to a table, the seat is fixed. The table distribution is also updated.
 - CRP automatically starts a new class when it thinks there should be a new class.
 - Hyperparameter c control how often to start a new class.
-
- Note: This is theoretically nice, but for most applications just pick a K .

Unsupervised Learning

- Clustering
- Autoencoders
- Generative Adversarial Networks

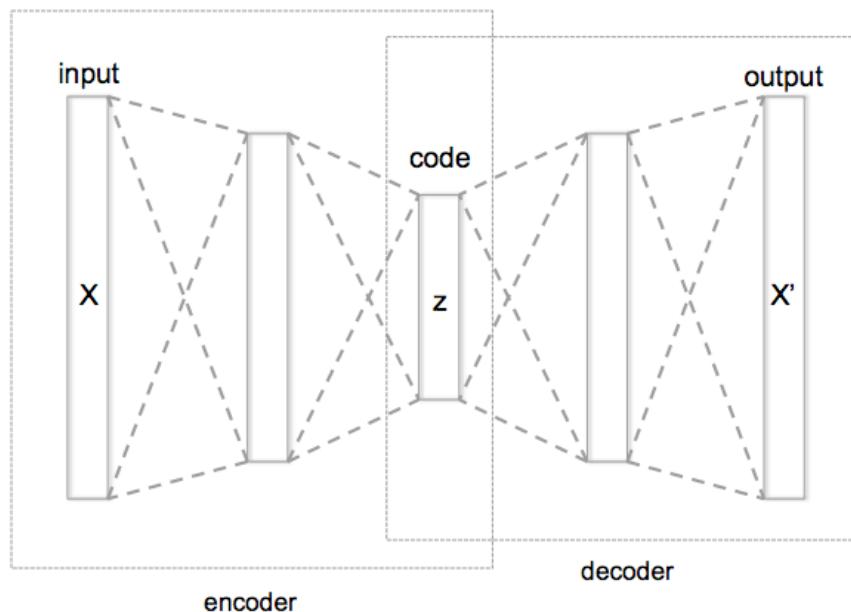
Inside Deep Learning

New experiments reveal how deep neural networks evolve as they learn.



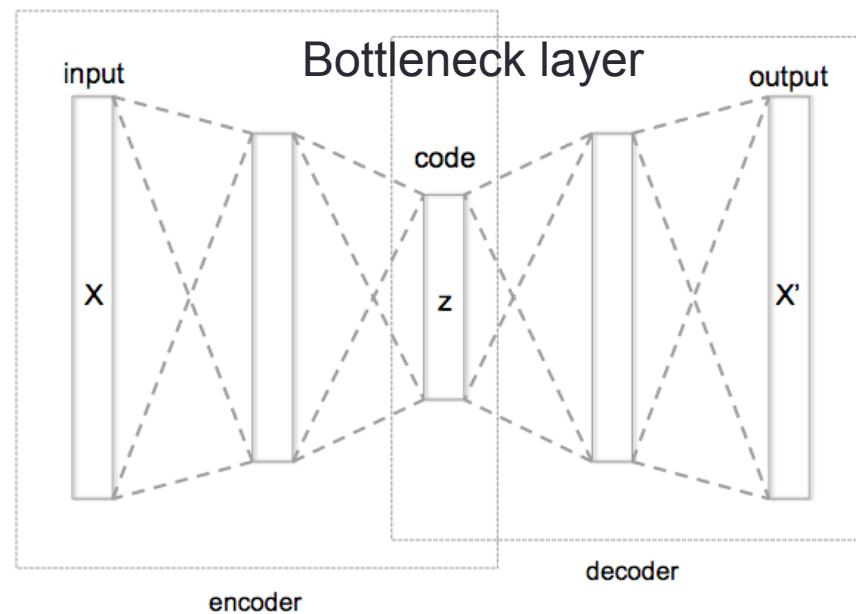
Autoencoders

- Use neural network properties to compress data
 - Automatically learn features



Simplest model

- Input X , and wants to output the same X
- Needs a bottleneck layer to enforce non-trivial solutions (identity matrix)
- Loss function = $\|x-x'\|^2$



Denoising autoencoder

- An autoencoder that denoise corrupted inputs
 - Blur image, sound corrupted by noise
- How?
 - Corrupt your input $x \rightarrow x''$
 - Use x'' as the input, and minimize the same loss
 - Loss = $\|x-x'\|^2$

Example

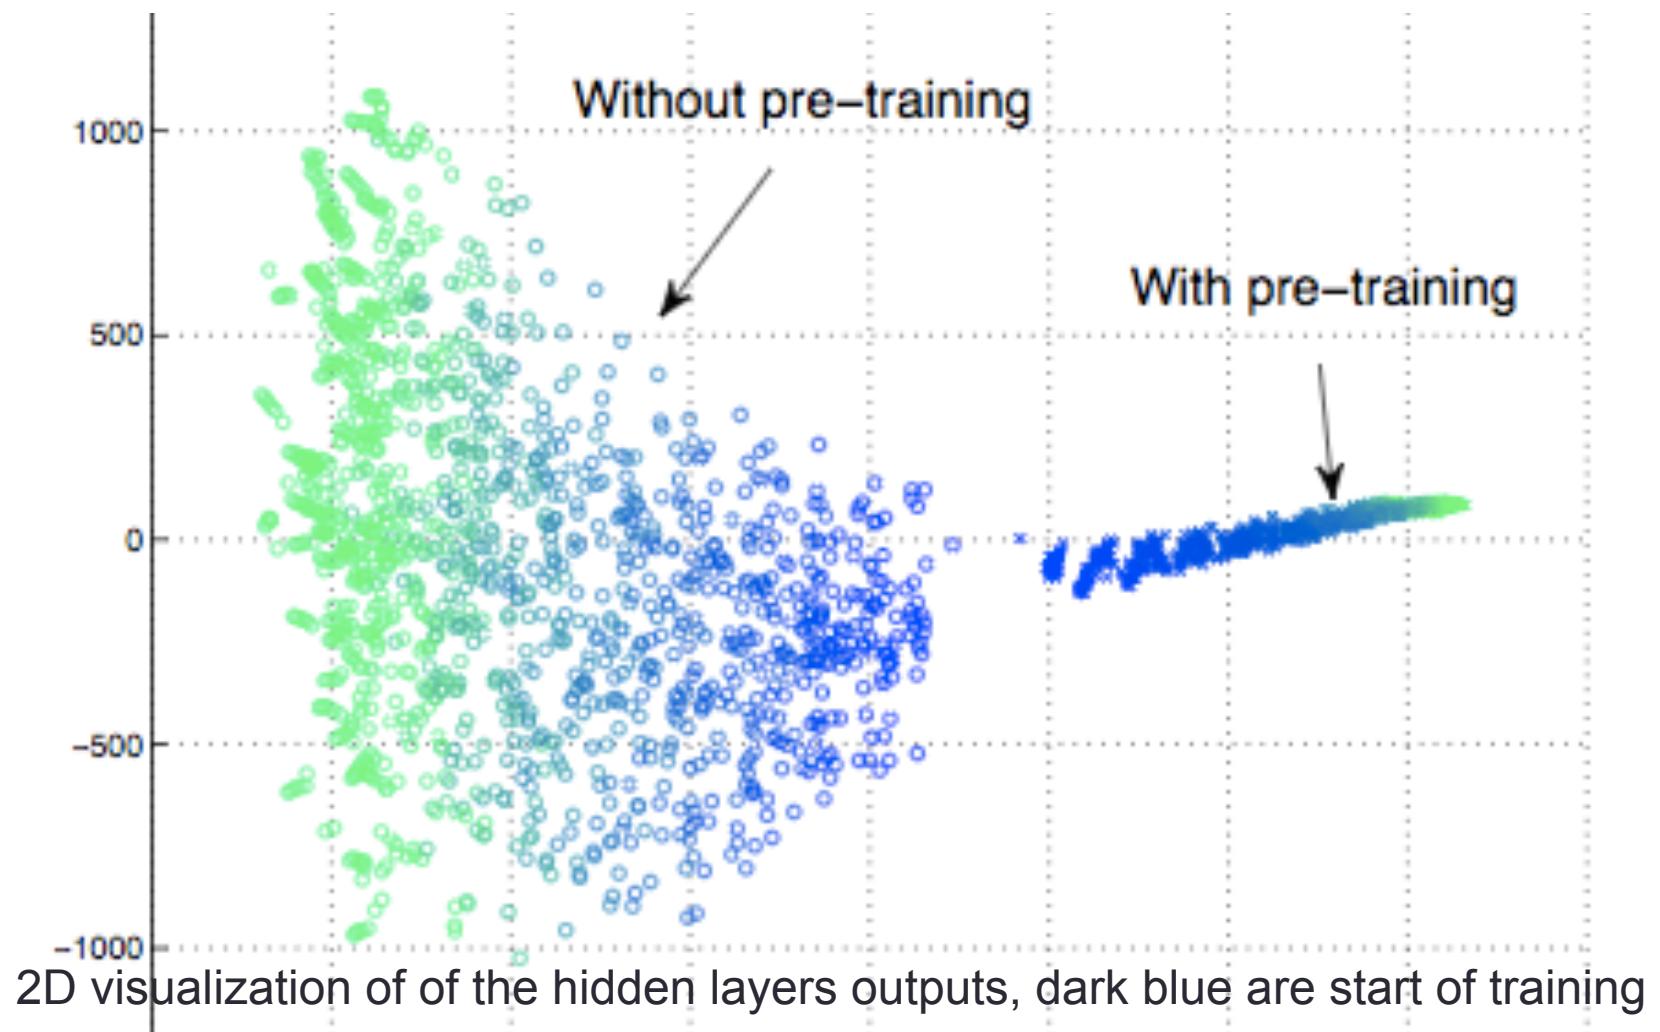
Original	Corrupted	Generated
7 2 1 0 4	7 2 1 0 4	7 2 1 0 4
1 4 9 5 9	1 4 9 5 9	1 4 9 5 9
0 6 9 0 1	0 6 9 0 1	0 6 9 0 1
5 9 7 3 4	5 9 7 3 4	5 9 7 3 4
9 6 4 5 4	9 6 4 5 4	9 6 4 5 4

<https://www.doc.ic.ac.uk/~js4416/163/website/autoencoders/denoising.html>

Other uses

- Besides denoising, can also be used to make your “code” more robust to the type of corruption you introduced.
- Another usage is for pre-training the neural network when you have lots of un-labeled data.
 - Use autoencoder on un-labeled data.
 - Do it layer by layer: Train with 1 layer, then train with 2 layer, ... until finish
 - Do regular SGD on labeled data with regular cross entropy loss
 - Helps with initialization (pre-training)

Pre-training effects



Variational Autoencoders (VAE)

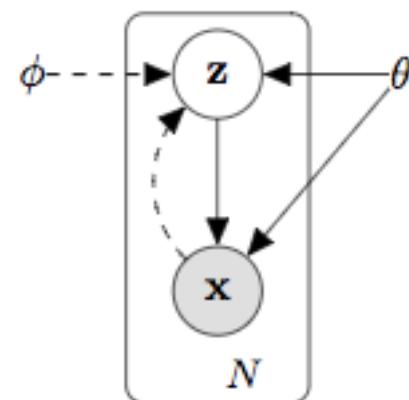
- Assume X is generated from a latent variable Z
 - Ex: a picture of a cat is generated from class label cat
- VAE wants to find this latent variable Z in the coding layer
- Our code can be generated from
 - $q_\phi(z|x)$ (our neural network)
 - which tries to mimick the true $p_\theta(z|x)$ that we don't know.

VAE reward (not loss)

$$E[\log P(X|z)] - D_{KL}[Q(z|X)\|P(z)]$$

Decoder

Encoder True Code



VAE reward explained

- We want to maxmimize

$$E[\log P(X|z)] - D_{KL}[Q(z|X)\|P(z)]$$

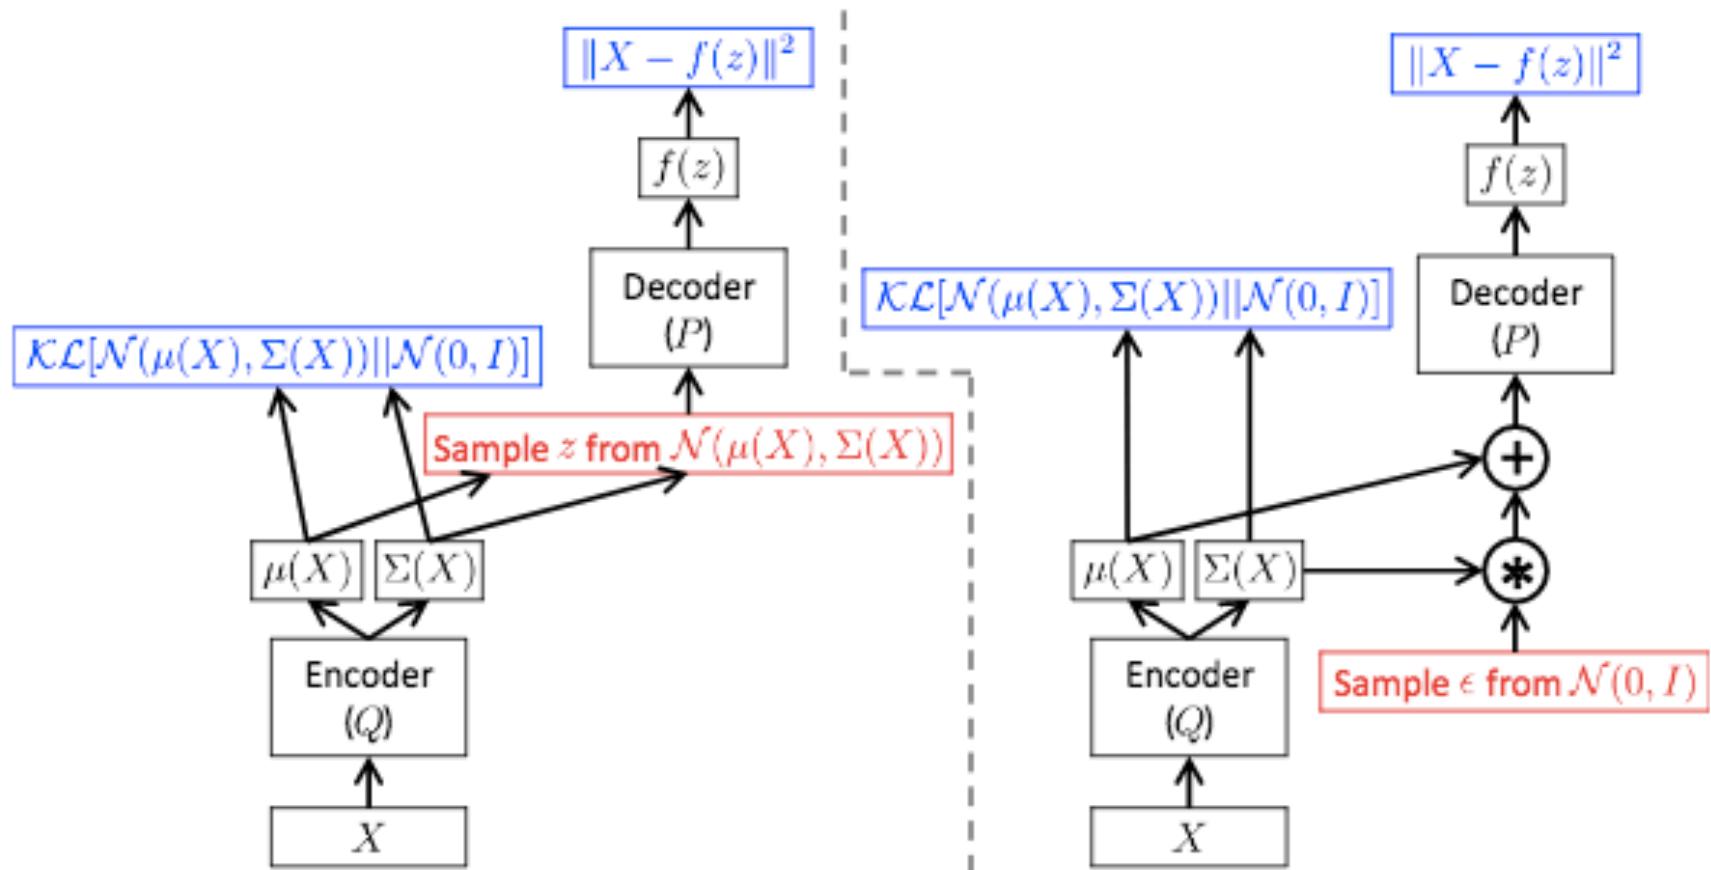
- D_{KL} is the KL-divergence
 - Distance between distributions
- The first term means we need to generates X in the best way possible from z. This is actually maximum likelihood estimation.
- What about $P(z)$?
 - This is what we get to pick, and the VAE will try to match it.
 - Simplest choice is $N(0,1)$
- For more info, read <https://arxiv.org/abs/1606.05908>

$$D_{KL}(P\|Q) = - \sum_i P(i) \log \frac{Q(i)}{P(i)},$$

And <https://wiseodd.github.io/techblog/2016/12/10/variational-autoencoder/>

VAE Training

$$E[\log P(X|z)] - D_{KL}[Q(z|X)\|P(z)]$$



Generative Adversarial Networks

- Let's recall what generative models are

Generative Models

- Naïve Bayes, Bayes classifiers are generative models
- Learn the model for each class y given input features x
 $p(x|y)$. (The likelihood probability)
 - Note we also model $p(y)$ and essentially $p(x,y)$
- To do classification we want to solve for the best y given input feature x (the posterior)

$$y^* = \operatorname{argmax}_y P(y|x)$$

- We can use Bayes' rule

$$y^* = \operatorname{argmax}_y \frac{P(x|y)P(y)}{P(x)}$$

Discriminative models

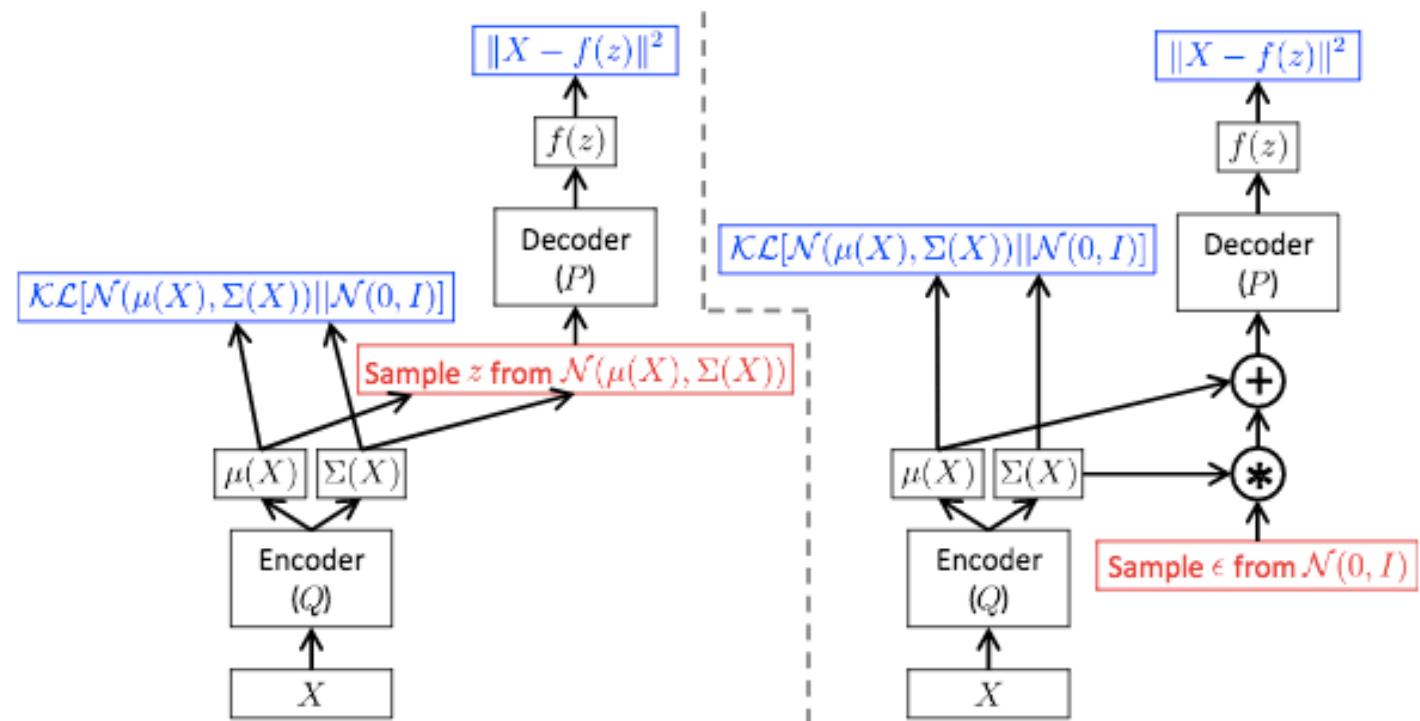
- Discriminative models model $P(y|x)$ directly

$$y^* = \operatorname{argmax}_y P(y|x)$$

- $P(y|x)$ is called the posterior probability
- Generally, $P(y|x)$ can be any function $h(y,x)$ that gives a score for each class
 - Logistic regression
 - SVM
 - Neural networks

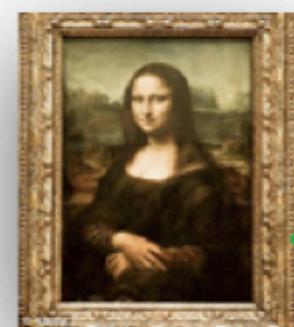
Generating via sampling

- With $P(X|Y)$ we can create a new X by being given Y



Generative Adversarial Networks (GAN)

- Consider a money counterfeiter
 - He wants to make fake money that looks real
 - There's a police that tries to differentiate fake and real money.
- The counterfeiter is the **adversary** and is generating fake inputs. – Generator network
- The police is try to discriminate between fake and real inputs. – Discriminator network



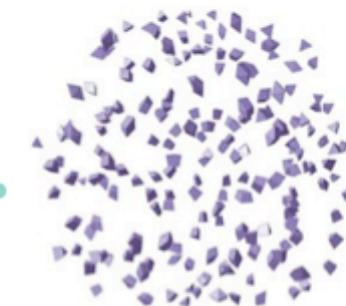
R: Real Data



D: Detective



G: Generator (Forger)



I: Input for Generator

GAN in a nutshell

$D(\cdot)$ – discriminator, outputs probability of real input

$G(\cdot)$ – generator, generates new data from random noise

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))].$$

Samples from data (real), The discriminator wants to give a high probability

Samples from z (noise), and generate $G(z)$. The discriminator wants to give a low probability

- Minmax: minimize the loss in the worst case scenario.
- At equilibrium, the discriminator will output 0.5 for real and fake data.

GAN training

- Create two networks, D and G
 - Discriminator training
 - Generates mini-batch of fake things + real things
 - GSD with loss

$$\frac{1}{m} \sum_{i=1}^m \left[\log D(\mathbf{x}^{(i)}) + \log (1 - D(G(\mathbf{z}^{(i)}))) \right]$$

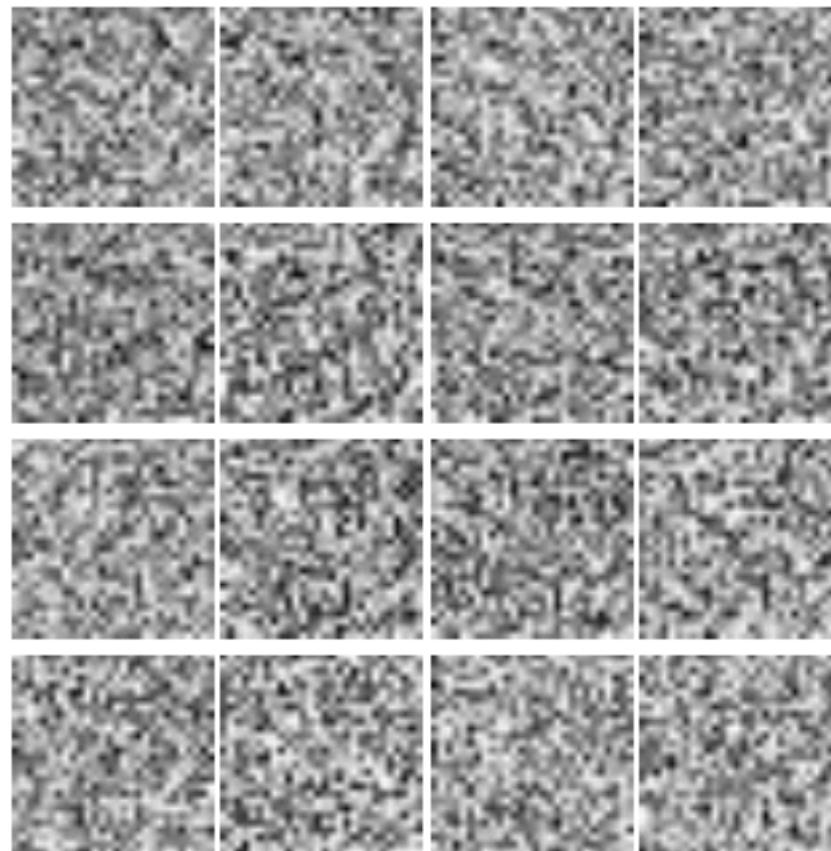
- Generator training
 - Generate fake things
 - GSD with loss

$$\frac{1}{m} \sum_{i=1}^m \log (1 - D(G(\mathbf{z}^{(i)})))$$

- Alternate between the two

GAN example

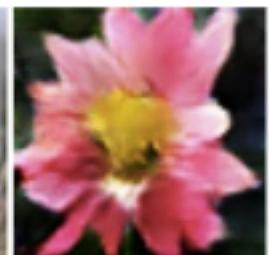
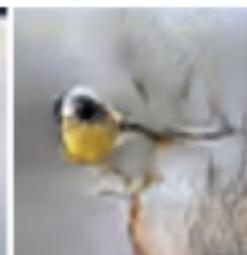
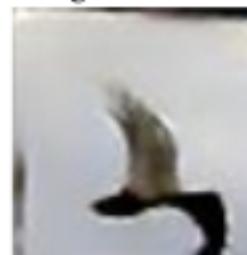
Generator output starts from random noise and gets better as we train.



Other uses

- StackGAN: text to photo

(a) StackGAN
Stage-I
64x64
images



(b) StackGAN
Stage-II
256x256
images



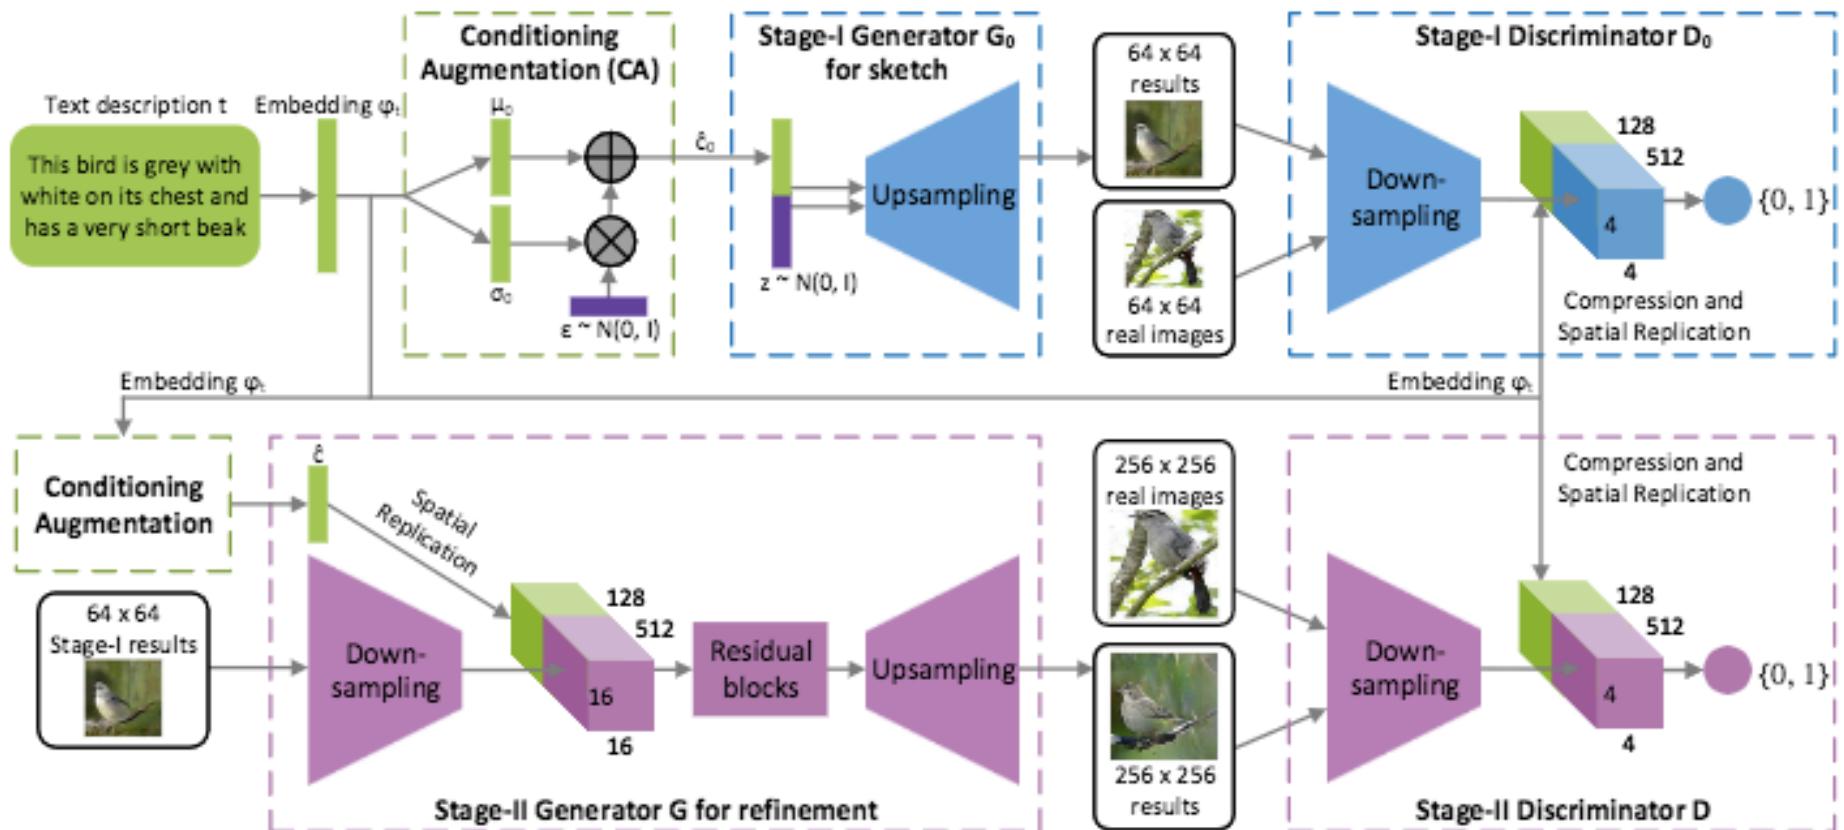
(c) Vanilla GAN
256x256
images



This bird is white with some black on its head and wings, and has a long orange beak

This bird has a yellow belly and tarsus, grey back, wings, and brown throat, nape with a black face

This flower has overlapping pink pointed petals surrounding a ring of short yellow filaments

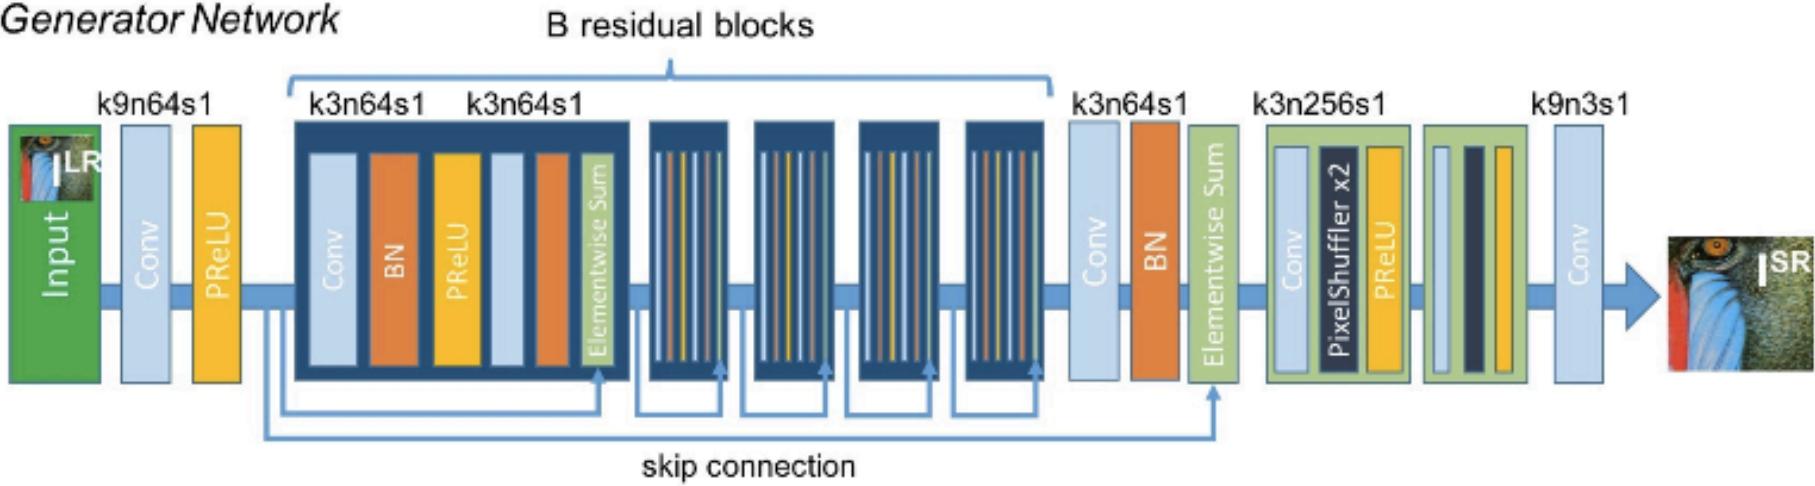


SRGAN

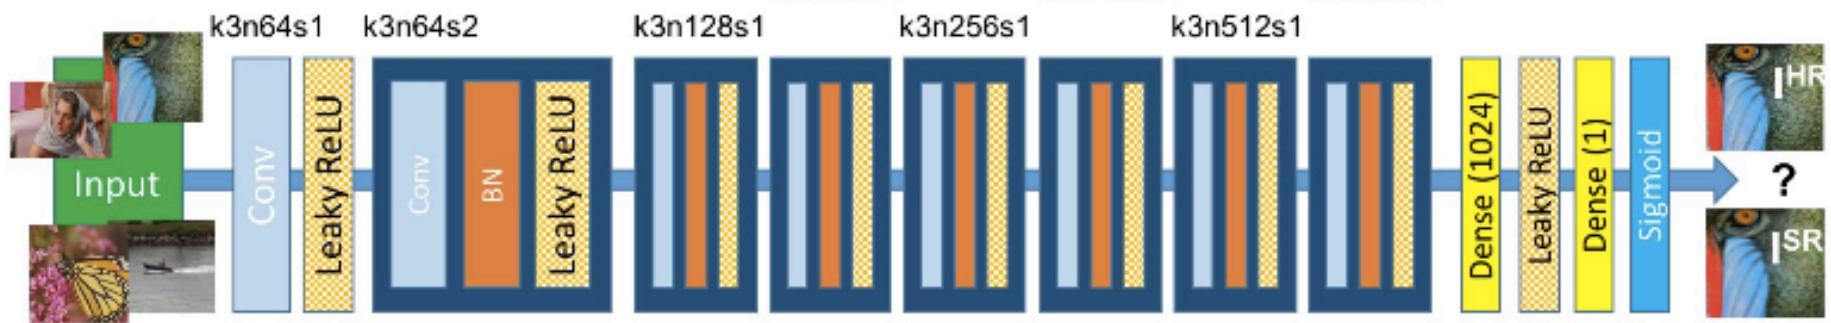
- <https://arxiv.org/pdf/1609.04802.pdf>



Generator Network



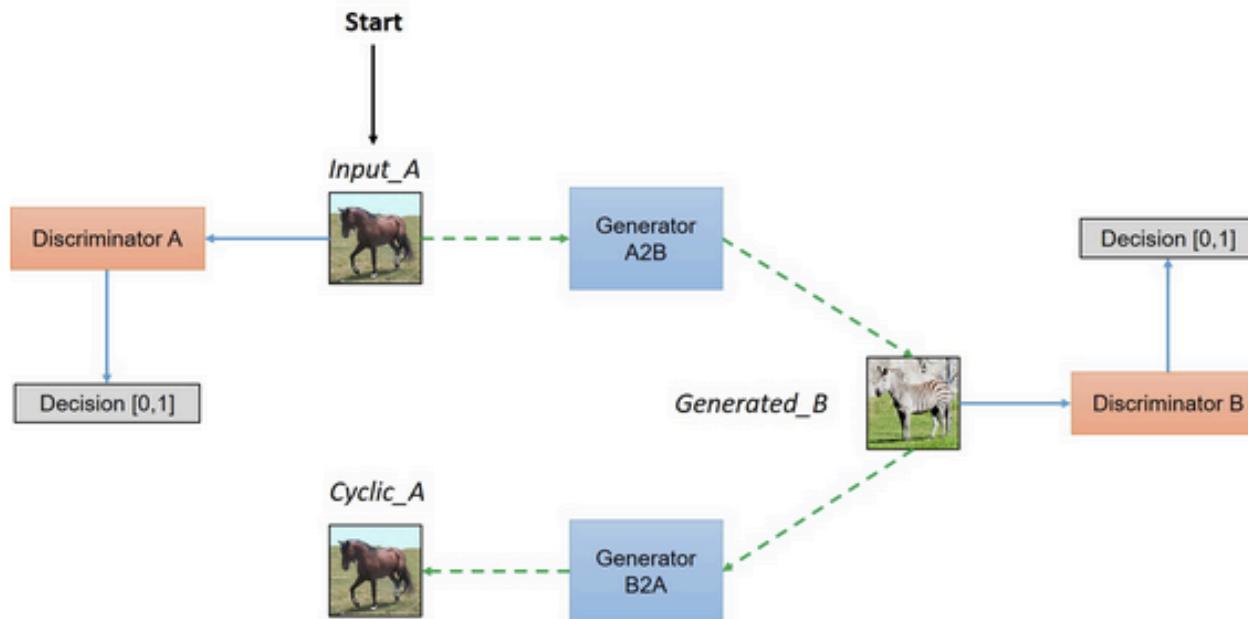
Discriminator Network



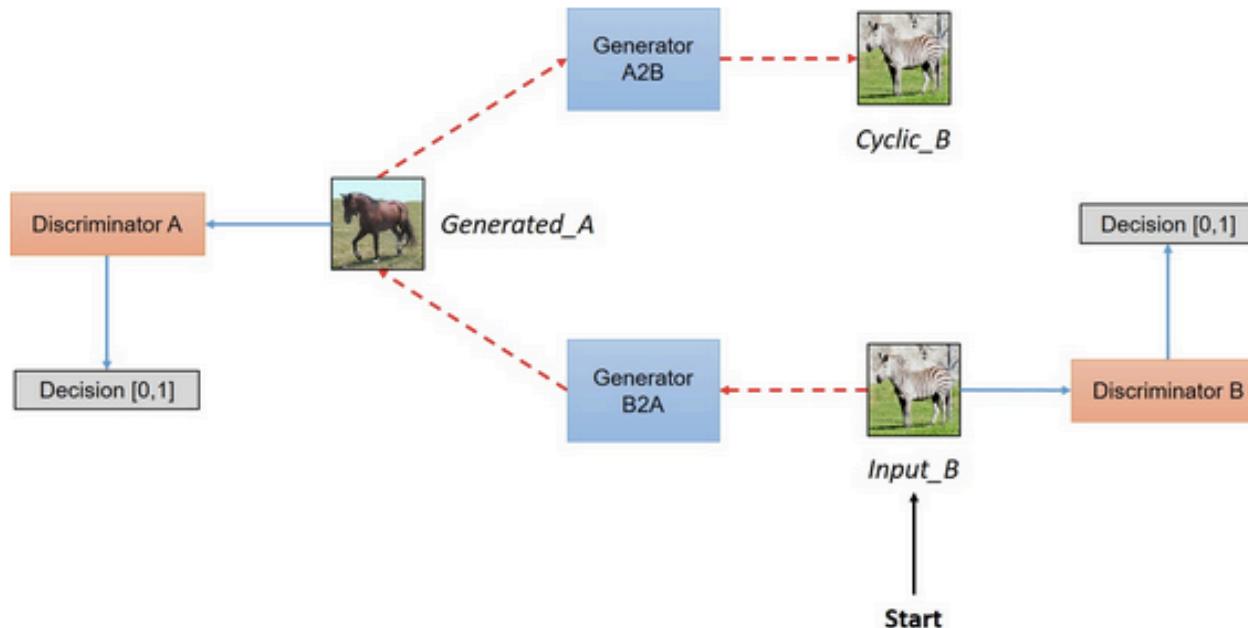
Style transfer with cycleGAN



<https://hardikbansal.github.io/CycleGANBlog/>



Cycle consistency



Note, you don't
need a paired
dataset

Machine translation with no parallel text

- MT usually requires parallel text

This is a cat

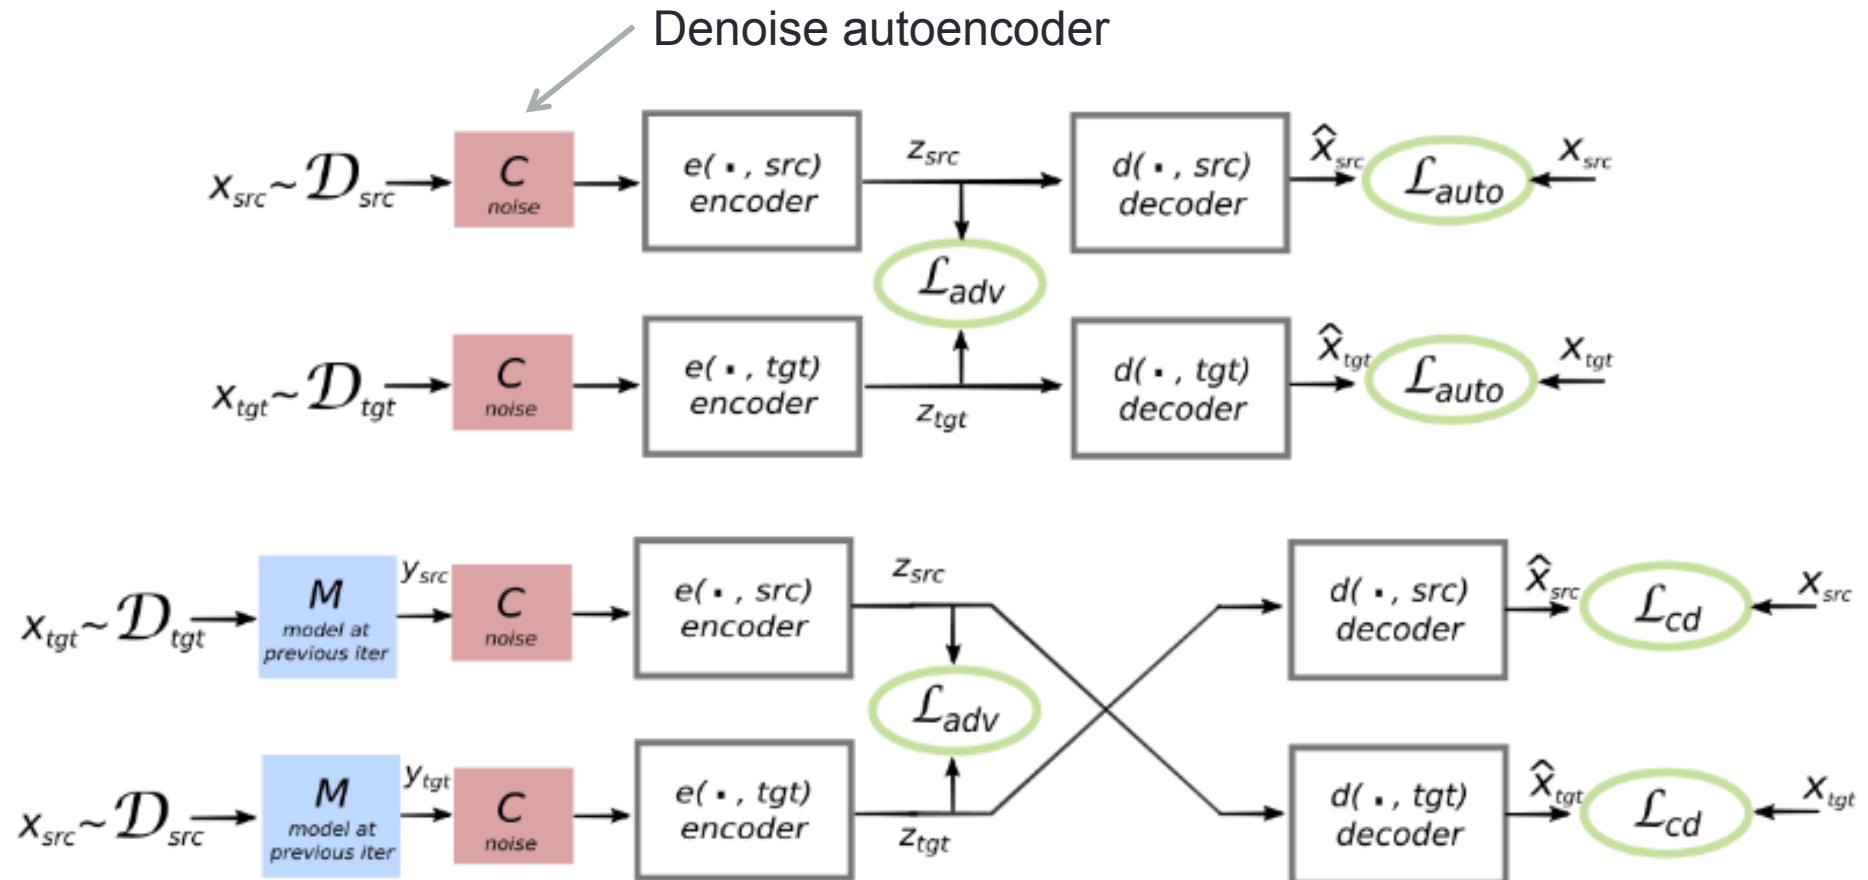
นี่คือแมว

- Most of the time we don't have parallel text. Just text.
- Can we still do MT?
 - Use GANs + Autoencoders

<https://arxiv.org/abs/1711.00043>

Use GAN Loss (\mathcal{L}_{adv}) to enforce that source and target language pairs share the same distributions.

Then enforce the translation distribution matches.



Results

	Multi30k-Task1				WMT			
	en-fr	fr-en	de-en	en-de	en-fr	fr-en	de-en	en-de
Supervised	56.83	50.77	38.38	35.16	27.97	26.13	25.61	21.33
word-by-word	8.54	16.77	15.72	5.39	6.28	10.09	10.77	7.06
word reordering	-	-	-	-	6.68	11.69	10.84	6.70
oracle word reordering	11.62	24.88	18.27	6.79	10.12	20.64	19.42	11.57
Our model: 1st iteration	27.48	28.07	23.69	19.32	12.10	11.79	11.10	8.86
Our model: 2nd iteration	31.72	30.49	24.73	21.16	14.42	13.49	13.25	9.75
Our model: 3rd iteration	32.76	32.07	26.26	22.74	15.05	14.31	13.33	9.64

Summary

- Various ways of clustering
 - K-means, GMM, LDA
- Encoders
 - Denoising, VAE
 - Generate nice representations of the input
- GAN
 - A neural network that we can sample (generate) data
 - Discriminator vs Generator

Reminder

- Mid report due tomorrow
- Meeting Thursday
- Next class: practical issues in machine learning application deployments, e.g. more various random stuff